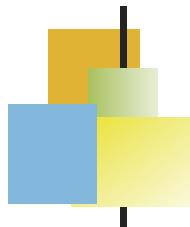


EECS 484: Database Management Systems

Instructor: H. V. Jagadish

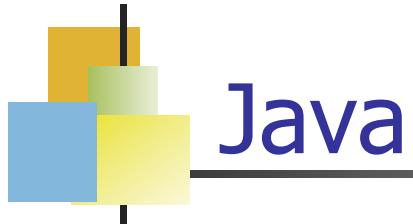
Email: jag@umich.edu

Office Hours: MW 10.30-11, BBB 4601

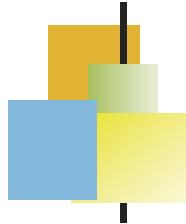


Course Outline – EECS 484

- GOAL: Basic introduction to database management systems.
- Two perspectives:
 - External (*Database user*)
 - Data models, ER model, relational model, SQL, database design ...
 - Java/JDBC Project: Common platform for building database applications
 - Internal (*Database implementer*)
 - File organizations, access methods, sorting, concurrency control, recovery, ...
 - Minirel Project: Build components of a Relational Database System
- Textbook “Database Management Systems”, by Raghu Ramakrishnan & Johannes Gehrke. 3rd ed.

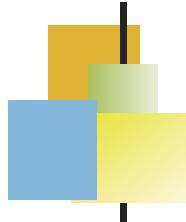


- Databases are most often accessed via SQL.
- But SQL is usually embedded in, and called from, a traditional programming language.
- Java is the most common choice, and so the one we have chosen for this course.
- Discussion this week is a tutorial on: Intro to Java for C++ programmers



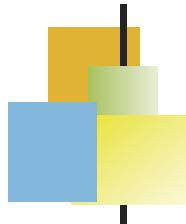
Groups

- Total of four group projects.
- Group of size 2
- Same partner for all projects
- Start looking for partners now!
- Register your group next week by following the link we provide.
 - If you are not in a group by next Friday, you will be assigned a partner at random.



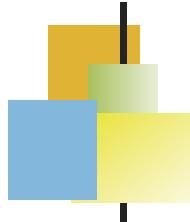
Project Grading

- Mostly autograder, some human.
- Limited number of submissions, even for autograded portion.
 - Make sure to test extensively.
- Same score for both partners.



Discussion Sections

- Not optional!
- Project covered in the discussion section.
- Exams may have project-related questions.
- 5 sections on Fridays – identical content, but different instructor.



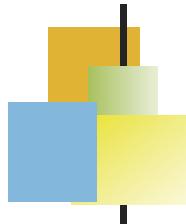
Academic Honesty

CoE Honor Code for all students.

Specifics on course website.

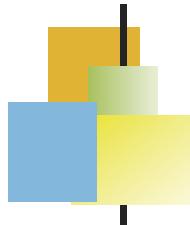
Can discuss, but cannot copy.

Questions? – Ask me first!



Course Policies

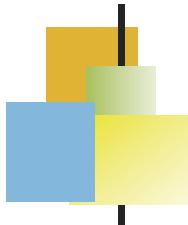
- Projects
 - Due by 11.55 p.m. on due date
 - 2 free late days total for all projects.
 - 1% course grade for each late day (or part thereof) used beyond the free day.
 - Up to 4 late days allowed per project.
- Assignments
 - Due by 11.55 p.m. on due date
 - No late submissions accepted.
 - Worst assignment dropped from total.



Course Grading

First Exam	25%
Second Exam	25%
4 of 5 written homework assignments [each worth 2.5%]	10%
Four projects [each worth 10%]	40%

No make up exams



Karma Points

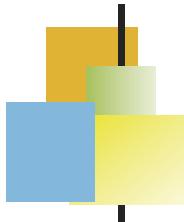
You can earn plus/minus 2% karma points.

You earn positive points by:

- Helpful posts on piazza
- Good questions in class
- Etc.

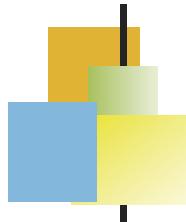
You earn negative points by:

- Posting duplicate questions on piazza without checking first
- Disrupting class
- Etc.



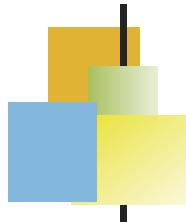
Exams

- Two exams:
- One midterm (Oct 23) and one final (Dec 15).
- Non-cumulative.
- No alternate exams, no make-up exams.
- Closed book but a one page, one-side, handwritten cheat sheet is permitted.



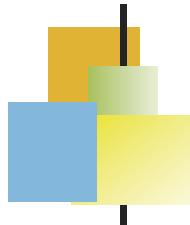
This week

- There is discussion.
 - Java basics for a C++ programmer.
 - (You will need this for project 2).
 - No need to go if you know Java.
- No office hours.
 - Regular office hours start next week.



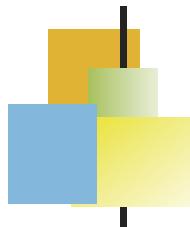
Lectures

- Notes are posted on canvas the night before.
- Sometimes updated after lecture.
 - To fix errors
 - To add clarifications
- Video recordings will be posted on canvas, usually a day or two after.



What Is a DBMS?

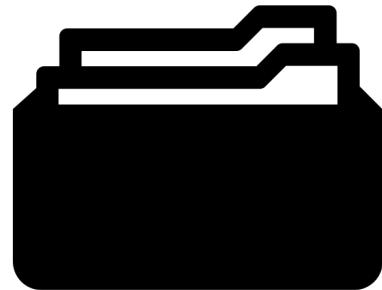
- DBMS = Database Management System
- Database: Large, integrated collection of data.
- Models some real-world *enterprise*
 - Entities (e.g., students, courses)
 - Relationships (e.g., Lisa Simpson is taking EECS 484)
- **DBMS:** a software package designed to store and manage databases

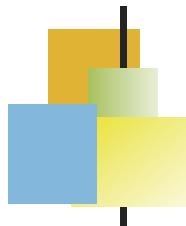


Old-time Solution: Sorted Student Folders

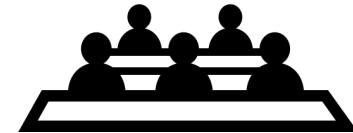


- Advantages?
- Disadvantages?

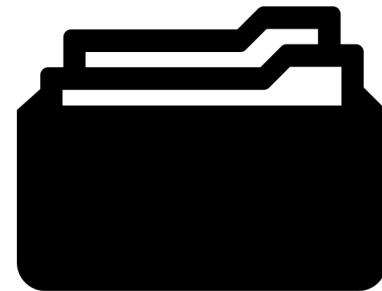


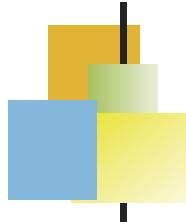


Old-time Solution: Sorted Student Folders

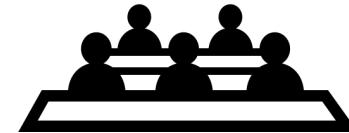


- Advantages?
 - cheap
- Disadvantages?
 - Large physical footprint
 - No sharing
 - No ad-hoc queries

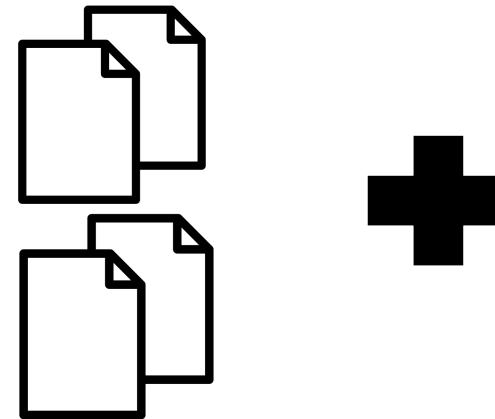


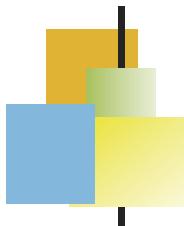


Other Solution: Flat Files



- Access?
 - using programs in C, Java, Python etc.
- Layout for the student records?





Other Solution: Flat Files



- Access?
 - using programs in C, Java, etc.
- Layout for the student records?

CSV:

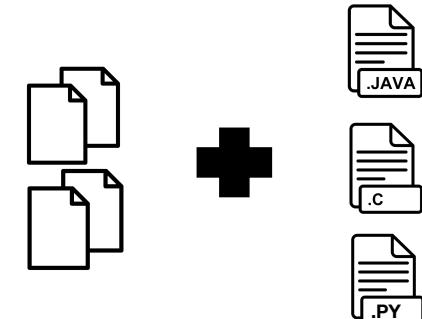
Brown, Lisa, lbrown, db, A, os, B

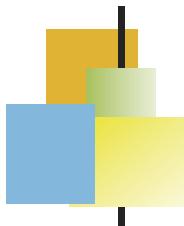
Smith, Bart, bsmith

Tompson, Mary, mtom, vis, B+, db, A-

...

...

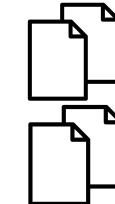




Other Solution: Flat Files



- Access?
 - using programs in C, Java, etc.
- Layout for the student records?



Multiple files: 

Brown, Lisa, lbrown

Smith, Bart, bsmith

Tompson, Mary, mtom

...

...



Ibrown, db, A

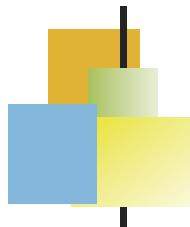
Ibrown, os, B

mtom, vis, B+

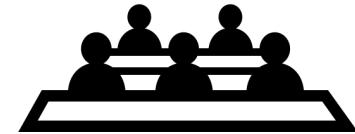
mtom, db, A-

...

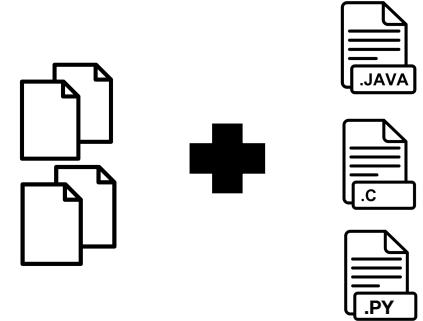
...

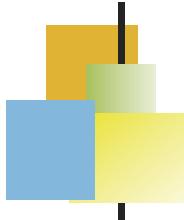


Other Solution: Flat Files

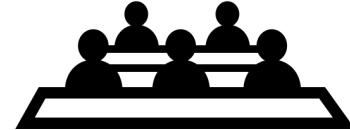


- Problems?

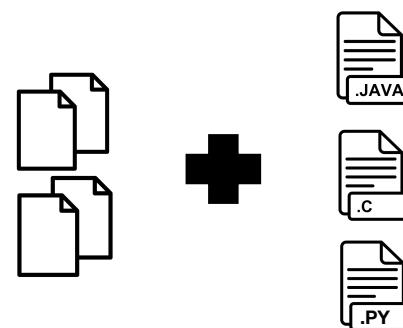


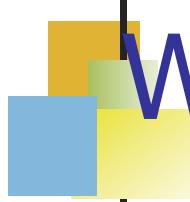


Other Solution: Flat Files



- Problems?
 - Inconvenient access to data
 - requires programming experience and knowledge of file layout
 - Data redundancy
 - Integrity problems
 - Atomicity problems (concurrent access issues)
 - Security problems

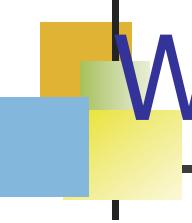




Why use a DBMS?

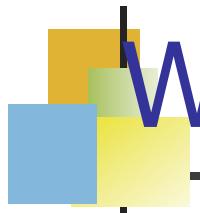


- It solves ALL these problems!
 - Data independence
 - Apps need a view of the data, not info about internal representation and storage
 - Efficient storage and access
 - Centralized data administration
 - Data integrity and security
 - Concurrent access, recovery from crashes
 - Reduced application dev time



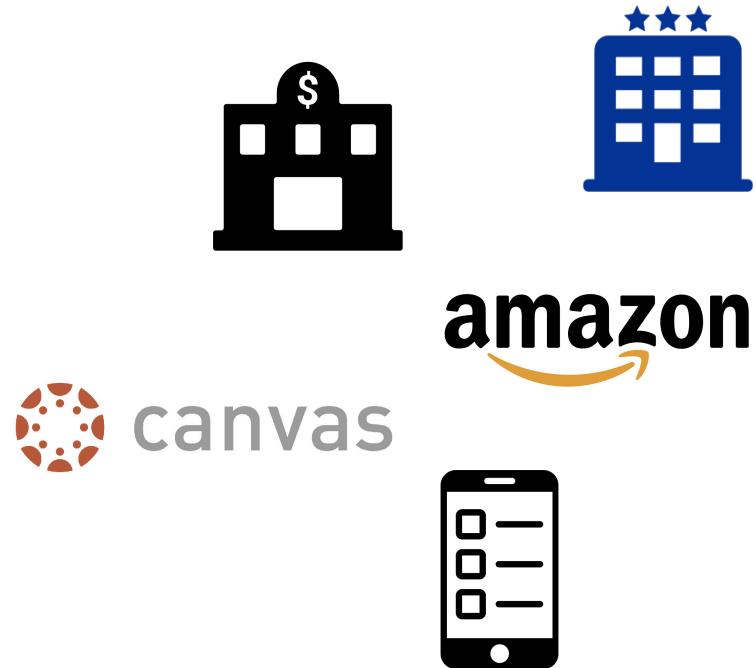
Who uses a DBMS?





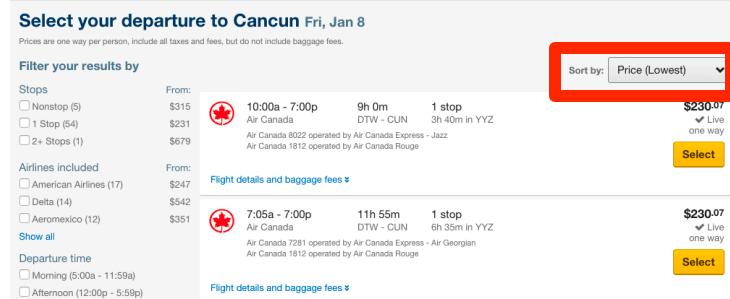
Who uses a DBMS?

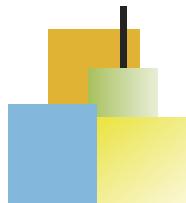
- Everyone!
 - Your bank
 - Your university
 - Your coffee shop
 - Your favorite hotel
 - Your favorite website
 - Your phone
 - Your government
- How many databases have you used so far today?



Why Study Databases?

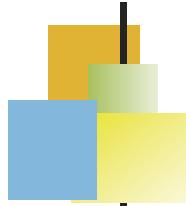
- Data is useless without the tools to extract information (queries)
 - “Optimal” pricing of an airline ticket
- Datasets increasing in diversity and volume
 - Websites, digital libraries, interactive video, human genome project, mobile applications
- Databases touch most of CS
 - OS, languages, theory, AI, multimedia, logic, ...





Data Models

- **Data model**: a collection of concepts for describing data.
- **Schema**: a description of a particular collection of data, using a given data model.
- **Relational model**: the most widely-used model today.
 - Data model: Database is a collection of **relations**
A relation is a table with rows and columns.
 - Every relation has a schema, which describes the columns (also called the fields or attributes).
- **Entity-Relationship (ER) model**: A “semantic” data model, i.e. a higher-level more user-intuitive model
 - A (relational) DBMS only understands the relational model
→ Must translate an ER schema to a relational schema



Relational and Other Data Models

- **DBMS using the relational DM** ('70s-'80s)

- IBM DB2
- Informix
- Oracle
- Sybase
- Microsoft Access
- Tandem
- Teradata
- ...

- **Other data models**

- ✧ Hierarchical (mid '60s-'70s)
 - IBM IMS
- ✧ Network ('70s)
 - IDMS, IDS
- ✧ Object-oriented (~'90s)
 - ObjectStore
- ✧ Object-relational (relational model + object DB concepts)
 - Oracle
- ✧ ...



Relational (Data) Model

- The most widely-used model today
- **Data model** = a collection of concepts for describing data
 - A collection of **relations**
 - **Relation** = set of records – think of it as a table with rows and columns

Students

sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Courses

cid	cname	cred.
E-484	EECS484	4
E-584	EECS584	3

Enrolled

sid	cid	grade
41	E-484	A-
13	E-584	A+



Relational (Data) Model

- **Schema** = a description of data in terms of a data model
 - Every relation has a schema
 - Specifies the **name** of the **relation**, the **name** and **type** of the **columns** (or fields or attributes)
 - Each row also called a **tuple** or a record

Students(sid:**string**, name:**string**, login:**string**, age:**integer**)

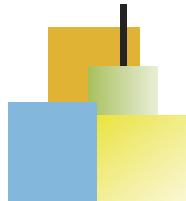
Courses(cid:**string**, cname:**string**, credits:**integer**)

Enrolled(sid:**string**, cid:**string**, grade:**string**)

Students			
sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

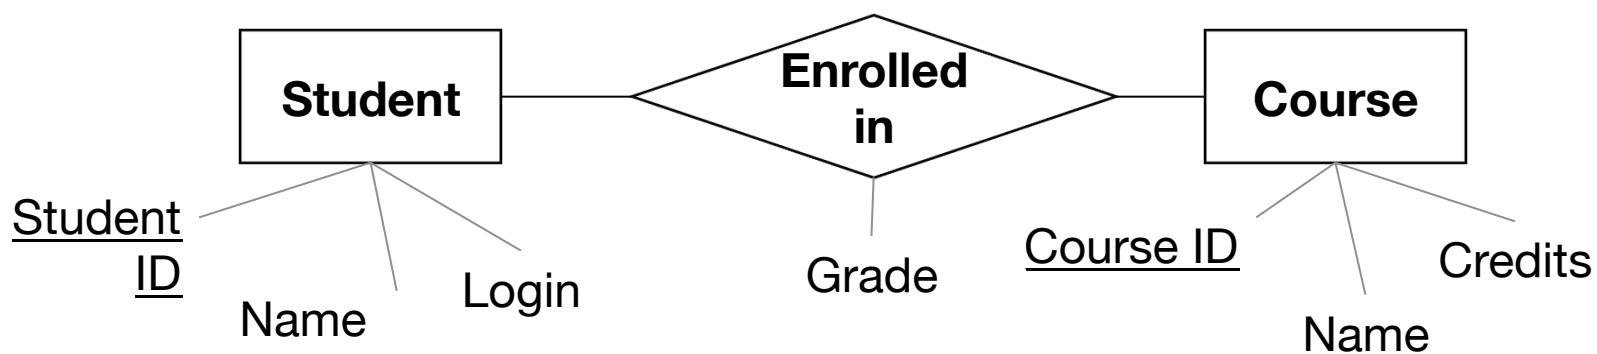
Courses		
cid	cname	cred.
E-484	EECS484	4
E-584	EECS584	3

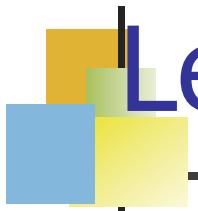
Enrolled		
sid	cid	grade
41	E-484	A-
13	E-584	A+



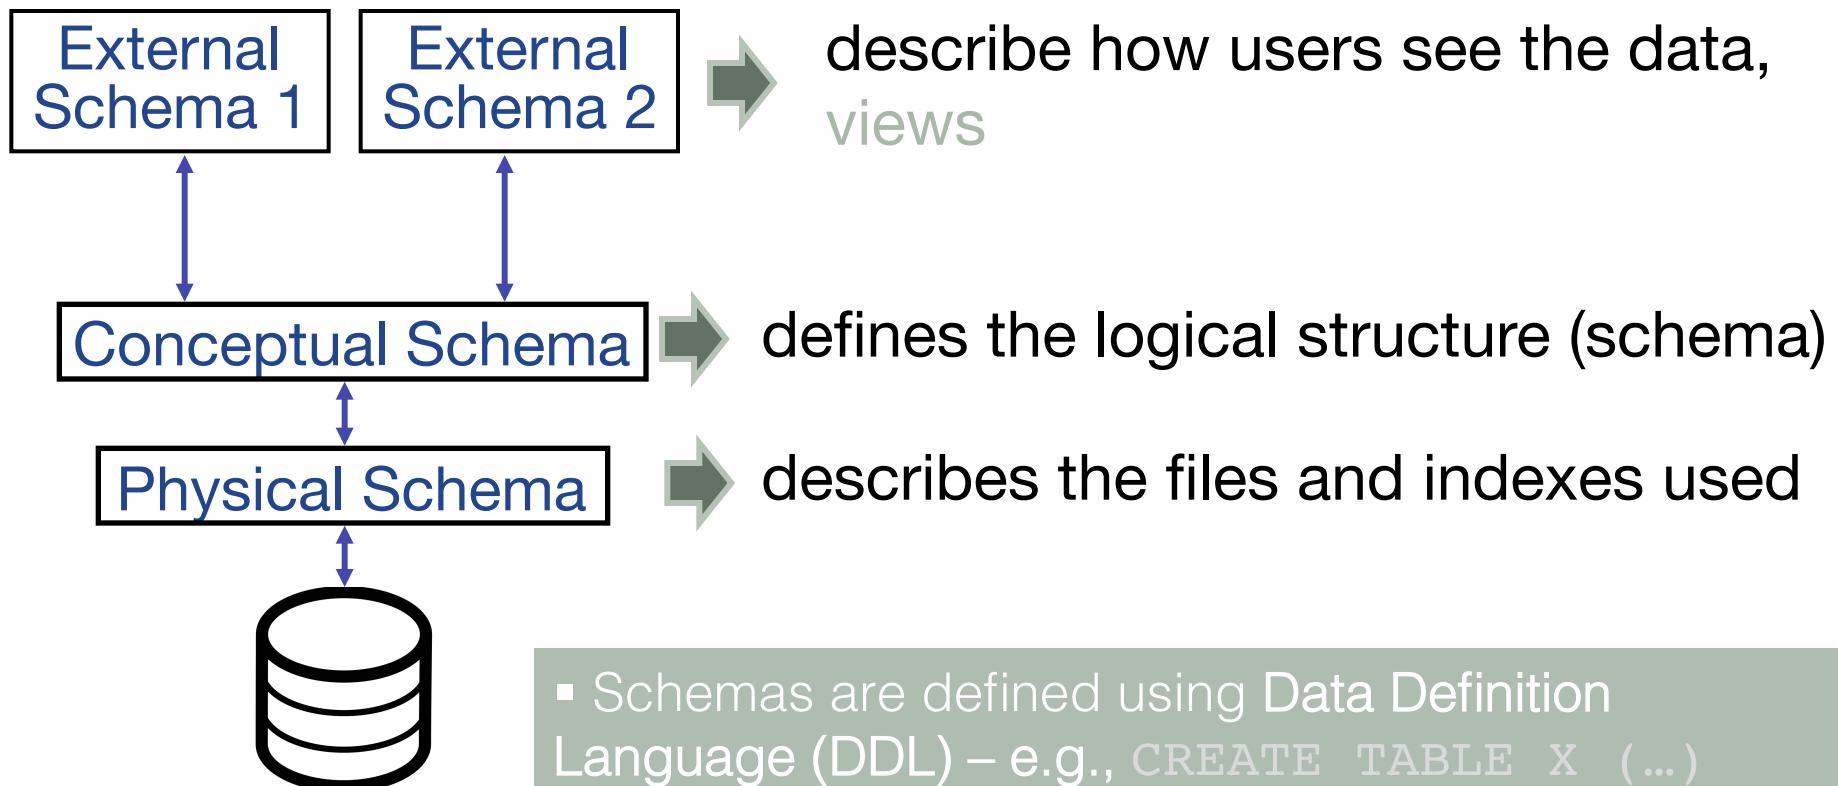
Entity-Relationship (ER) Model

- A “semantic” data model
 - a higher-level, more user-intuitive model
 - A (relational) DBMS understands the relational model
 - Must translate an ER schema to a relational schema
- Entity-Relationship diagram:
 - **Entities:** Student, Course
 - **Relationship:** Enrolled_in



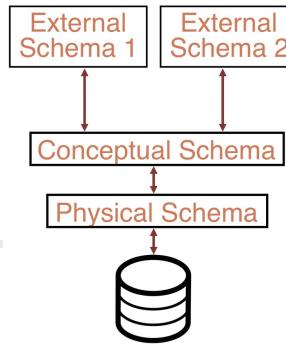


Levels of Abstraction



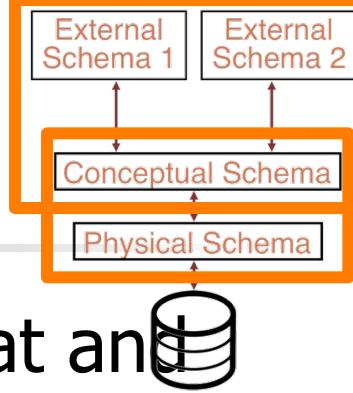
- Schemas are defined using Data Definition Language (DDL) – e.g., CREATE TABLE X (...)
- Data is modified/queried using Data Manipulation Language (DML) – e.g., SELECT FROM X WHERE ...

Example

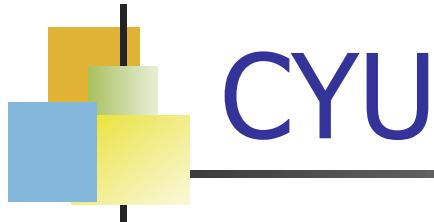


- Conceptual schema (1):
 - `Students(sid:string, name:string, login:string, age:integer)`
 - `Courses(cid:string, cname:string, credits:integer)`
 - `Enrolled(sid:string, cid:string, grade:string)`
- Physical schema (1):
 - Relations stored as unordered files.
 - Index on first column of `Courses`.
- External Schema (≥ 1):
 - View: `Course_info(cid:string, enrollment:integer)`
 - View: `Class_rank(sid:string, gpa:real, rank:integer)`

Data Independence

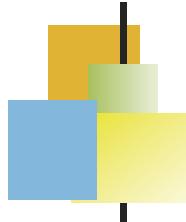


- Applications insulated from data format and storage details
- Logical data independence: Protection from changes in *logical* structure of data
 - External / Conceptual schemas
- Physical data independence: Protection from changes in *physical* structure of data
 - Conceptual / Physical schemas



- Which of these are more suitable for storing in a DBMS rather than files in an OS?
 - (a) Grades for students at the university
 - (b) Source code for a program
 - (c) Contents of a textbook

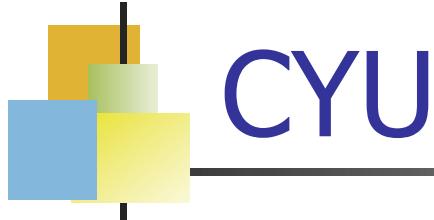




CYU

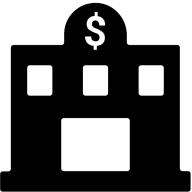
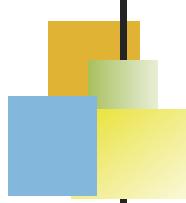
- Let's say UM provides you access to a relational table that gives just your grades in various courses. Does that relation represent:
 - An external schema?
 - A conceptual schema?
 - A physical schema?



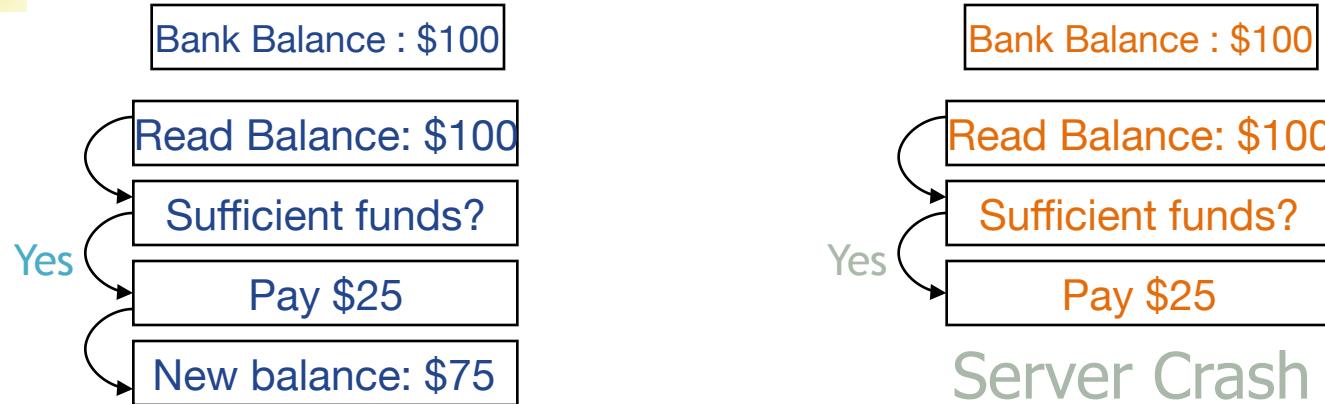


- The relational table with student grade information is very large and stored on multiple servers for performance. Does the storage scheme represent:
 - a) An external schema?
 - b) A conceptual schema?
 - c) A physical schema?

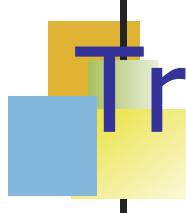




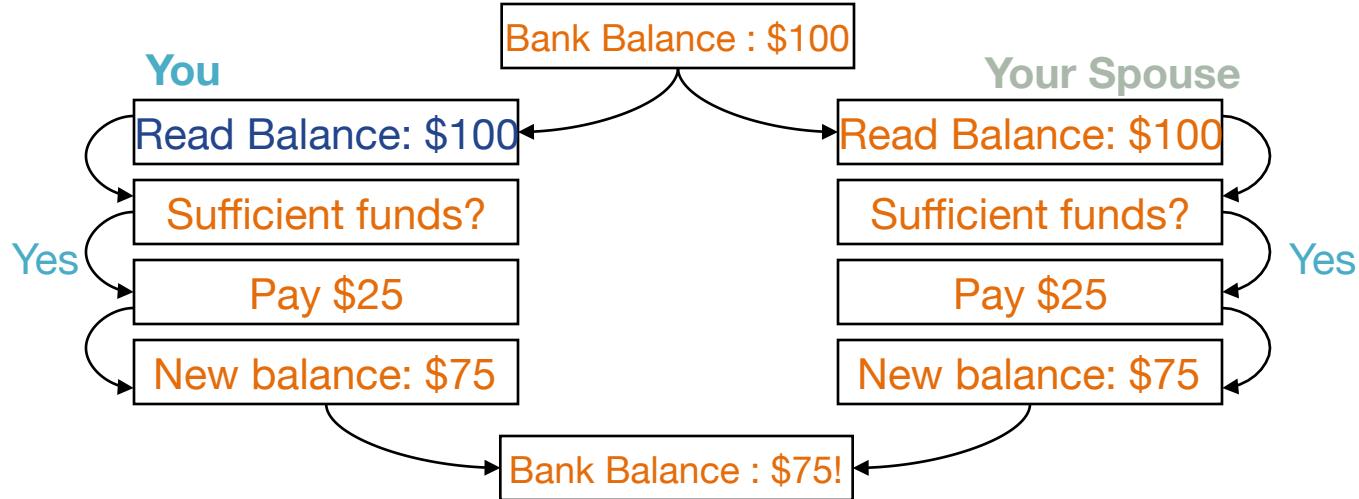
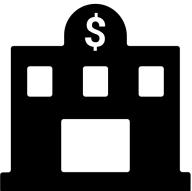
Transactions (I)



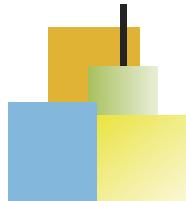
- **Transaction:** any one execution of a user program
in a DBMS
- Inconsistency caused by incomplete operations
- DBMS ensures atomic operations!
 - i.e., all or nothing!
 - Automatic recovery from crashes!



Transactions (II)



- Inconsistency caused by interleaving actions of different user programs
- DBMS provides the illusion of a “single-user” system
 - Key concept: **Transaction**, an atomic sequence of R/W
 - Concurrency control, transaction management

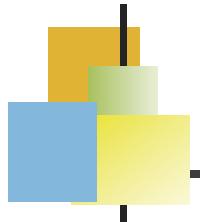


Lots of People use DBMS ...

- DBMS vendors
- DB application programmers
 - E.g. smart webmasters
- *Database administrator (DBA)*
 - Designs logical /physical schemas
 - Handles security and authorization
 - Data availability, crash recovery
 - Database tuning as needs evolve



Must understand how a DBMS works!



Summary

- DBMS used to maintain, query large datasets.
- Benefits include recovery from system crashes, concurrent access, quick application development, data integrity and security.
- Levels of abstraction give data independence.
- DBAs hold responsible jobs and are **well-paid!**
- DBMS R&D is one of the most exciting areas in CS.