

# **Logical Database Design: Mapping ER to Relational**

Chapter 3, Section 3.5

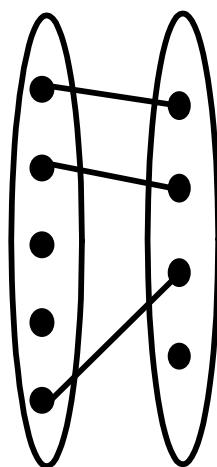
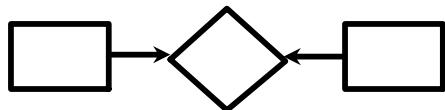
# ER Model vs. Relational Model

- ER Model used for conceptual design
- Relational Model implemented by modern DBMS
- Important Step: Translate ER diagram to Relational schema

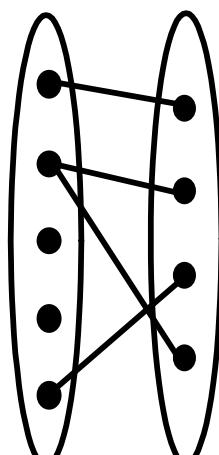
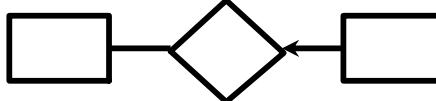
# Recall ER Constructs

- Basic Constructs
  - Entity Sets
  - Relationship Sets
  - Attributes (of entities and relationships)
- Additional Constructs
  - ISA Hierarchies
  - Weak Entities
  - Aggregation
- Integrity Constraints
  - Key constraints
  - Participation constraints
  - Overlap / Covering constraints for ISA hierarchies

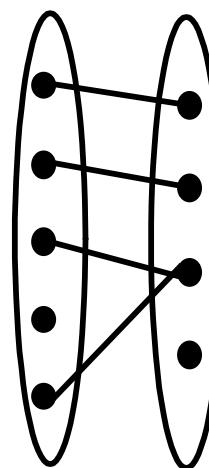
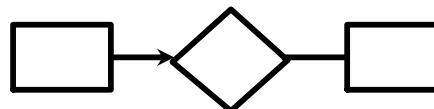
# Review: ER relationship types



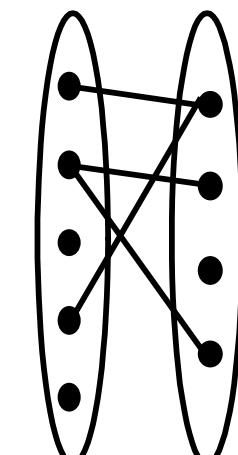
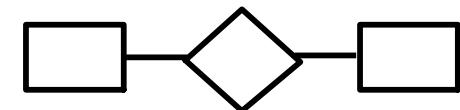
1-to-1



1-to Many  
1-to-N



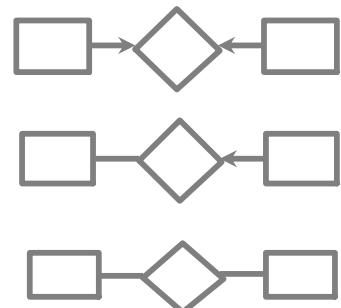
Many-to-1  
N-to-1



Many-to Many  
M-to-N

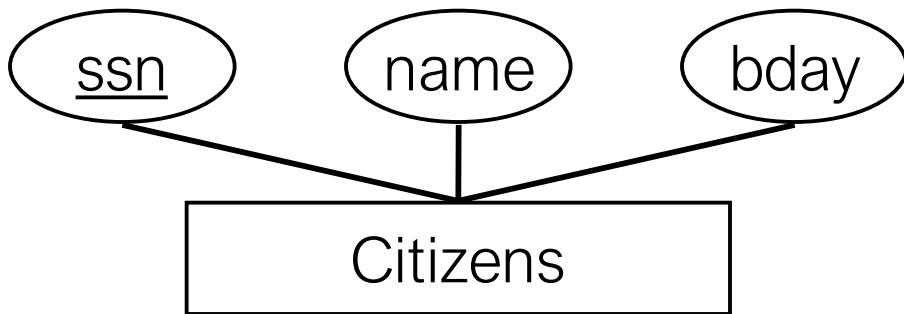
# ER to Tables: Basics

- Strong entities:
  - key = primary key
- Binary relationships: keys come from the participating entities
  - 1-to-1: either key (other = candidate key)
  - 1-to-N: the key of the ‘many’ (N) part
  - M-to-N: both keys





# Entity Sets to Tables

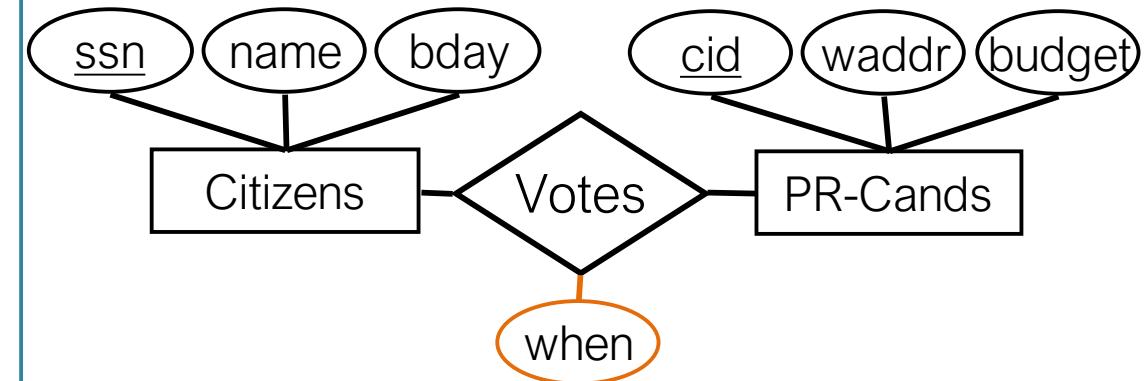


```
CREATE TABLE Citizens  
  (ssn      CHAR(11),  
   name     CHAR(20),  
   bday     DATE,  
   PRIMARY KEY (ssn))
```

Can ssn have a null value?



# Relationship Sets to Tables

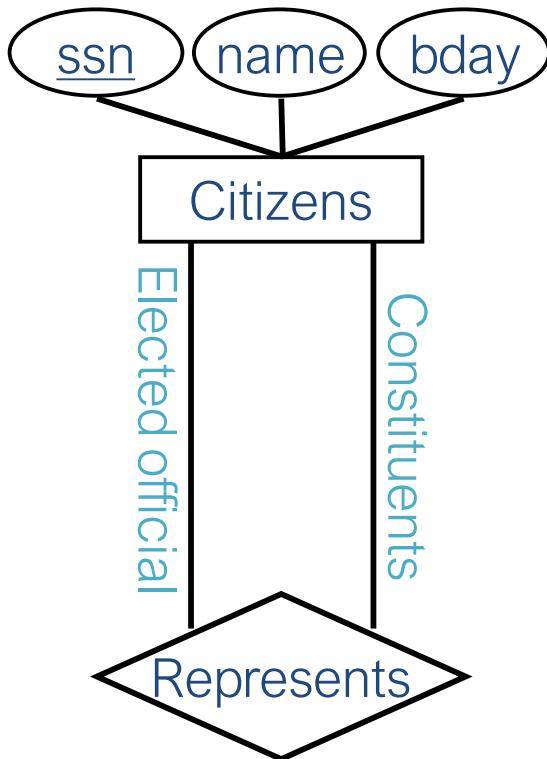


```
CREATE TABLE Votes (
    ssn      CHAR(11),
    cid      INTEGER,
    when     DATE,
PRIMARY KEY (ssn, cid),
FOREIGN KEY (ssn) REFERENCES Citizens,
FOREIGN KEY (cid) REFERENCES PR-Cands)
```

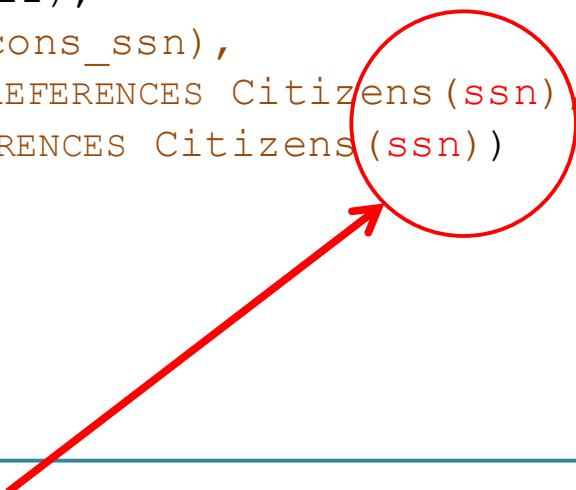
- Foreign key: keys from participating entity sets
- Descriptive attributes

Generalizes to n-ary relationships  
(we will see example later)

# Relationship Sets to Tables



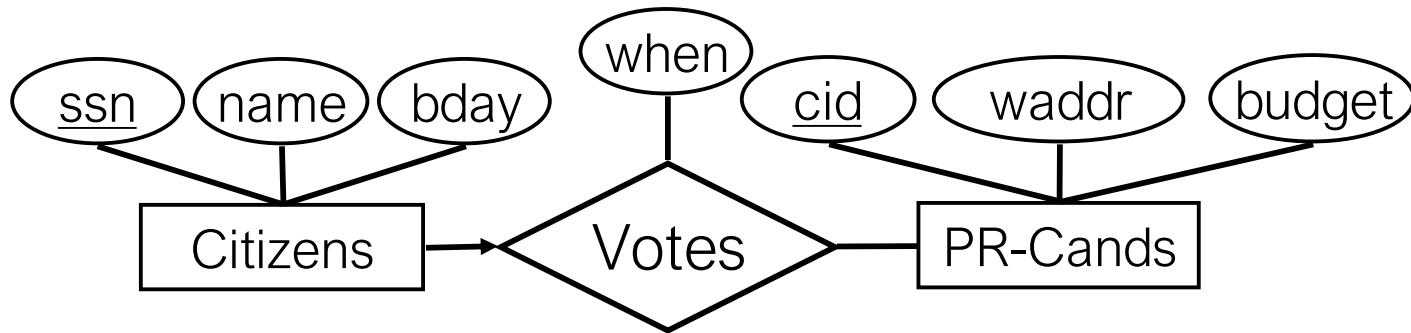
```
CREATE TABLE Represents (
    elected_ssn  CHAR(11),
    cons_ssn      CHAR(11),
PRIMARY KEY (elected_ssn, cons_ssn),
FOREIGN KEY (elected_ssn) REFERENCES Citizens(ssn),
FOREIGN KEY (cons_ssn) REFERENCES Citizens(ssn))
```



Note that you need to specify the column that you are referring to in the `Citizens` table, as `Citizens` does not have “`elected_ssn`” nor “`cons_ssn`”.



# Key Constraints

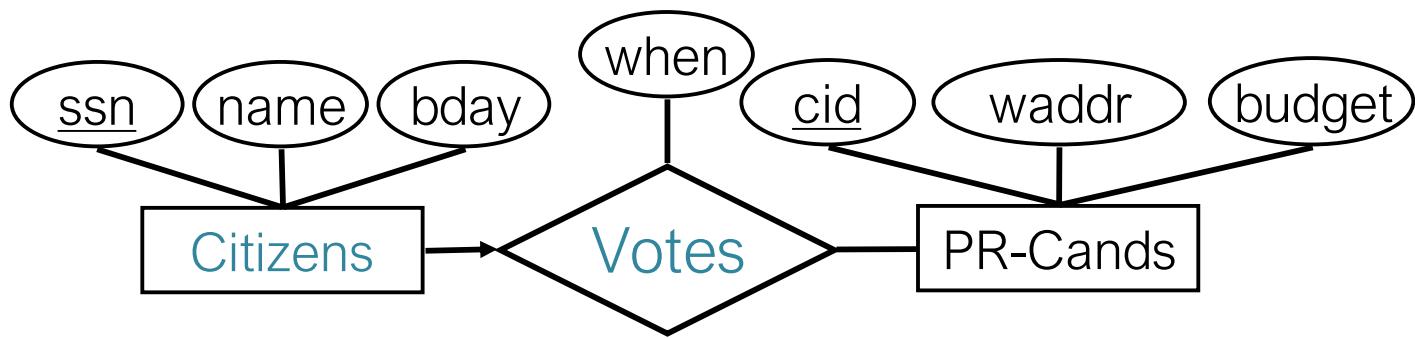


- Approach 1: **Three Tables** (Citizens, Votes, PR-Cands)

```
CREATE TABLE Votes
(  ssn      CHAR(11),
   cid      INTEGER,
   when    DATE,
   PRIMARY KEY (ssn),
   FOREIGN KEY (ssn) REFERENCES Citizens,
   FOREIGN KEY (cid) REFERENCES PR-Cands)
```



# Key Constraints



*Each citizen can only vote once, so OK to fold 'Votes' relationship into 'Citizens' entity*

- Approach 2: Two Tables (Citizen\_Votes, PR-Cands)

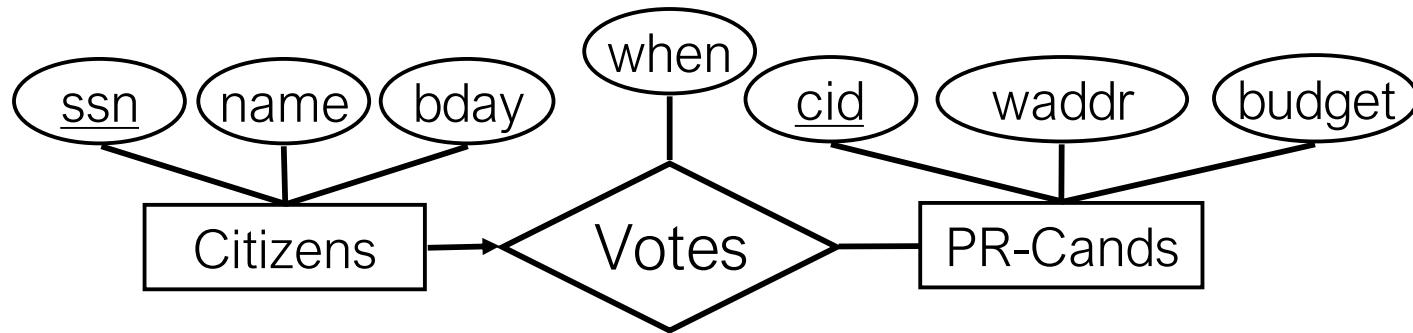
```
CREATE TABLE Citizen_Votes (
    ssn      CHAR(11),
    name     CHAR(20),
    bday     DATE,
    when    DATE,
    cid      INTEGER,
PRIMARY KEY  (ssn),
FOREIGN KEY (cid) REFERENCES PR-Cands)
```

Q: Can cid be null?  
Q: What if many citizens don't vote?  
Q: Which approach is better?





# Key Constraints

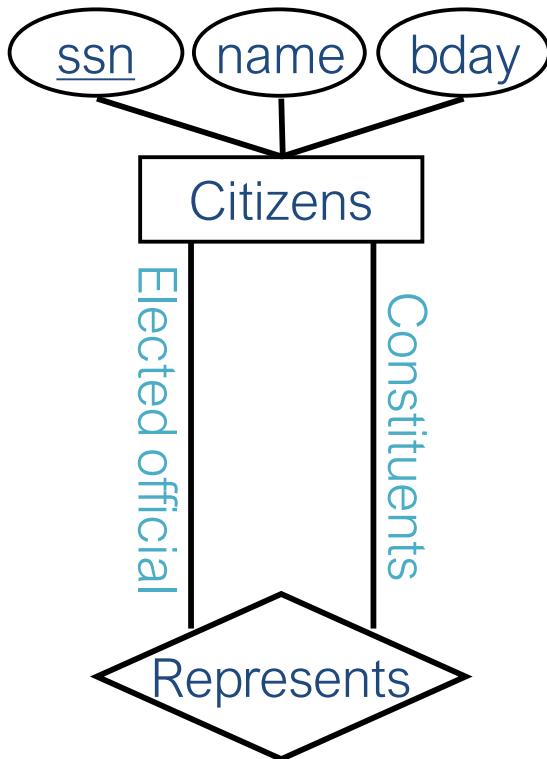


- What about **one** table? 
- No! This is bad design.  
e.g., For each citizen that votes for Trump, we have to store Trump's information (cid, waddr, budget) => **REDUNDANCY!**

# Primary Keys and ICs

- In (most) DBMSs the primary key (or any of its parts, if it is composite) cannot get NULL value.
  - In this class, we will stick with this rule (despite the fact that there are exceptions)
- Although you can choose how to handle UPDATE/DELETE actions when there are references (via foreign keys), some options may be in conflict with the primary key (PK) constraints
  - The system will not allow changes that violate the PK constraints.

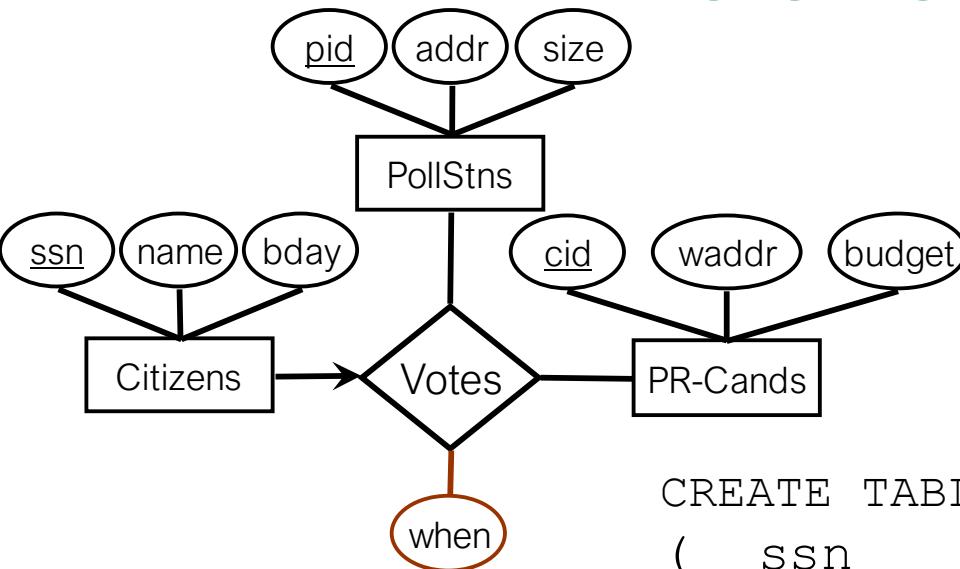
# Relationship Sets to Tables



```
CREATE TABLE Represents (
    elected_ssn  CHAR(11),
    cons_ssn      CHAR(11),
PRIMARY KEY (elected_ssn, cons_ssn),
FOREIGN KEY (elected_ssn) REFERENCES Citizens(ssn),
FOREIGN KEY (cons_ssn) REFERENCES Citizens(ssn))
```

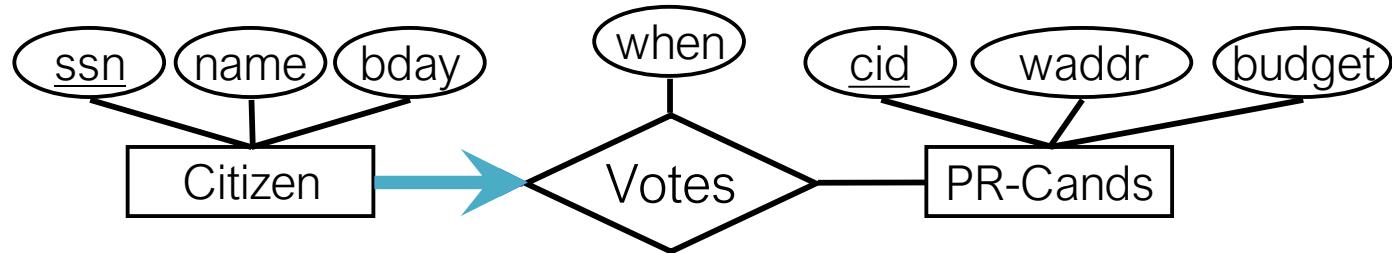
Note that you need to specify the column that you are referring to in the `Citizens` table, as `Citizens` does not have “`elected_ssn`” nor “`cons_ssn`”.

# Key Constraints: N-ary relationships



```
CREATE TABLE Citizens_Votes
(   ssn        CHAR(11),
    name       CHAR(20),
    bday       DATE,
    when      DATE,
    pid        INTEGER,
    cid        INTEGER,
    PRIMARY KEY  (ssn),
    FOREIGN KEY (cid) REFERENCES PR-Cands,
    FOREIGN KEY (pid) REFERENCES PollStns)
```

# Participation Constraints



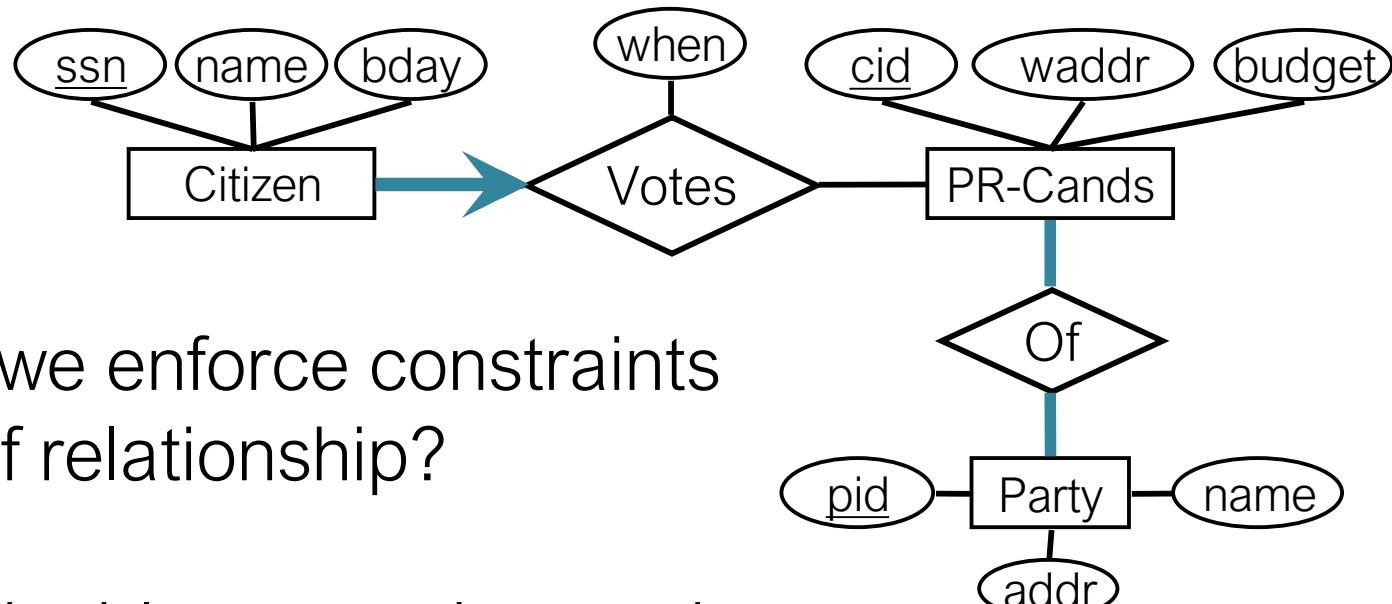
## Using Approach 2

```
CREATE TABLE Citizen_Votes (
    ssn      CHAR(11),
    name     CHAR(20),
    bday     DATE,
    when    DATE,
    cid      INTEGER NOT NULL,
    PRIMARY KEY (ssn),
    FOREIGN KEY (cid) REFERENCES PR_Cands
        ON DELETE NO ACTION);
```

Can we enforce the participation constraint using Approach 1 (three tables) ?



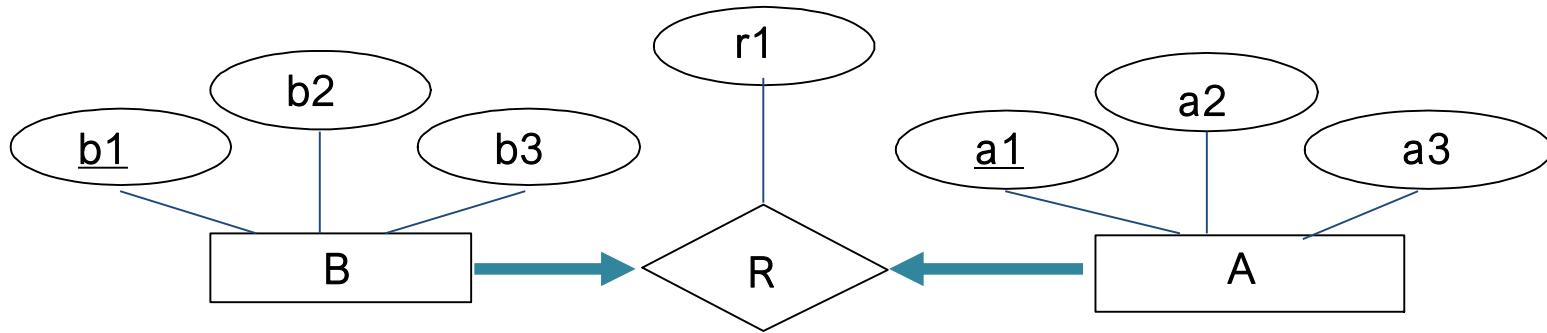
# Participation Constraints



Can we enforce constraints  
on Of relationship?

Need table constraints and  
assertions (later).

# Mapping Participation Constraints (1-to-1)

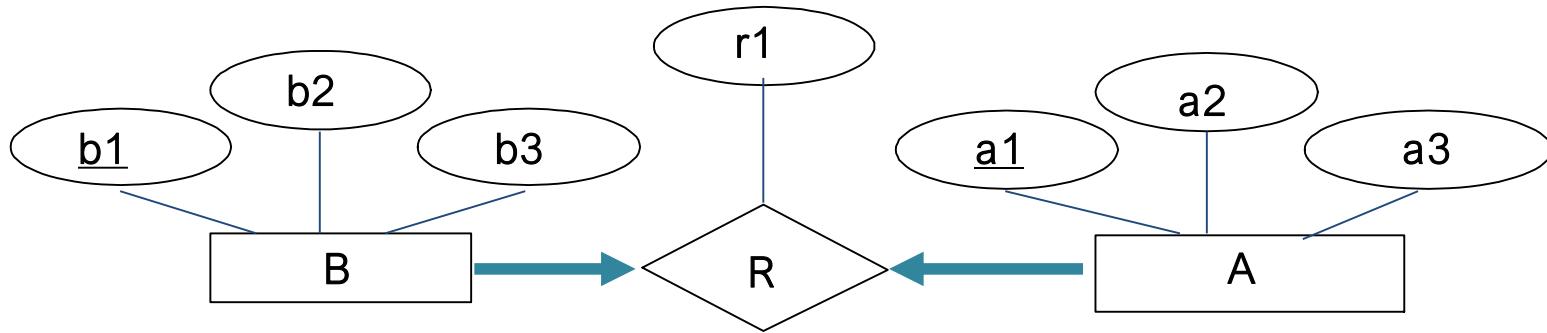


```
CREATE TABLE RAB (
    r1 Integer,
    a1 Integer,
    a2 Integer,
    a3 Integer,
    b1 Integer,
    b2 Integer,
    b3 Integer ...)
```

Key constraints?

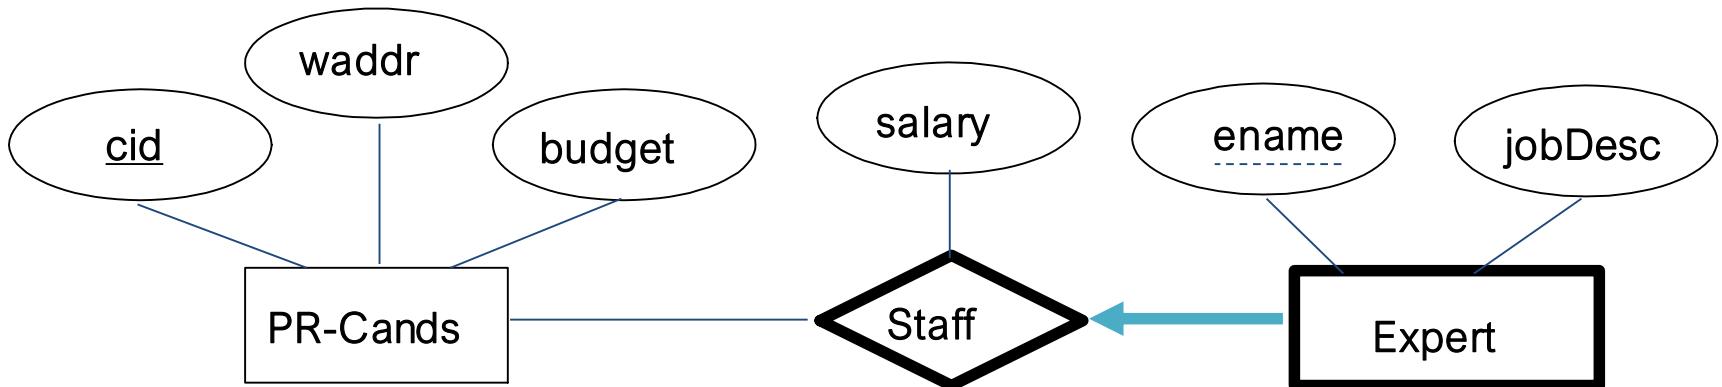


# Mapping Participation Constraints (1-to-1)



```
CREATE TABLE RAB (
    r1 Integer,
    a1 Integer,
    a2 Integer,
    a3 Integer,
    b1 Integer NOT NULL,
    b2 Integer,
    b3 Integer,
    UNIQUE (b1), PRIMARY KEY (a1))
```

# Weak Entities

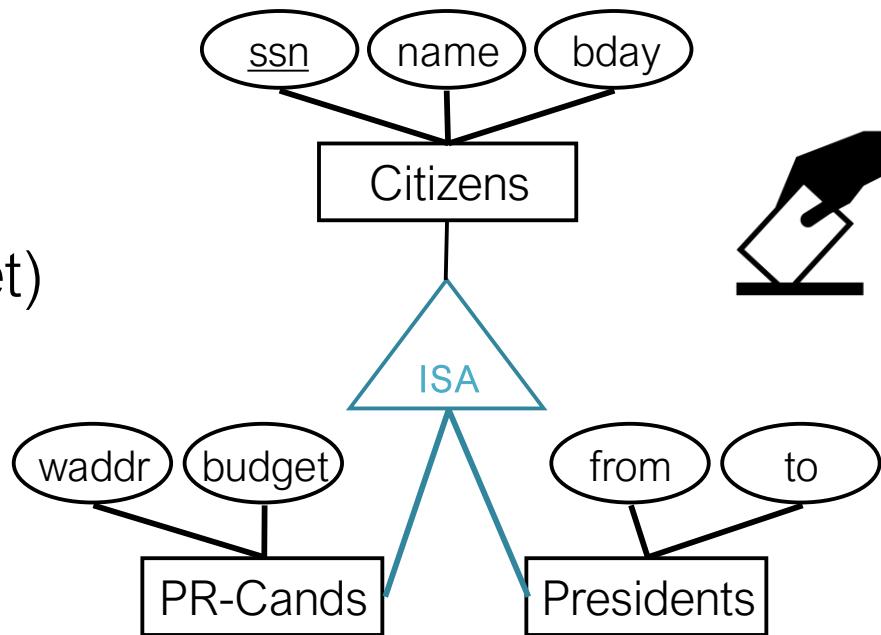


- Approach 2: Combine weak entity and owning relationship into one relation
  - Delete all weak entities when an owner entity is deleted.

```
CREATE TABLE Expert_Staff (
    ename      CHAR(20),
    jobDesc   CHAR(40),
    salary     REAL,
    cid       INTEGER,
    PRIMARY KEY (ename, cid),
    FOREIGN KEY (cid) REFERENCES PR-Cands
    ON DELETE CASCADE)
```

# ISA Hierarchies: General Approach

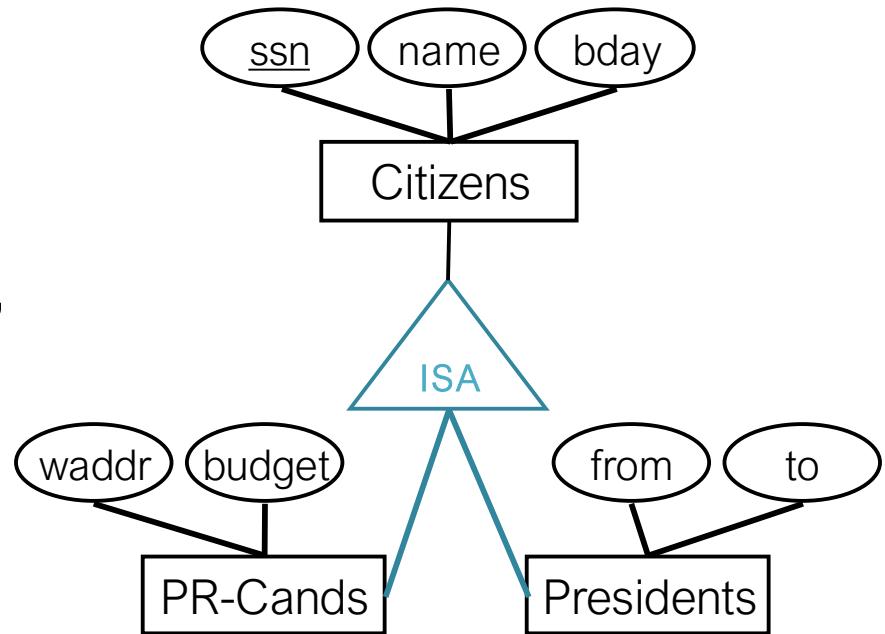
- Three relations:
  - Citizens (ssn, name, bday)
  - PR-Cands (ssn, waddr, budget)
  - Presidents (ssn, from, to)
- Queries:
  - Involving all citizens => Easy
  - Involving just PR-Cands => need to join PR-cands with Citizens to get some attributes



# ISA Hierarchies: Alternative



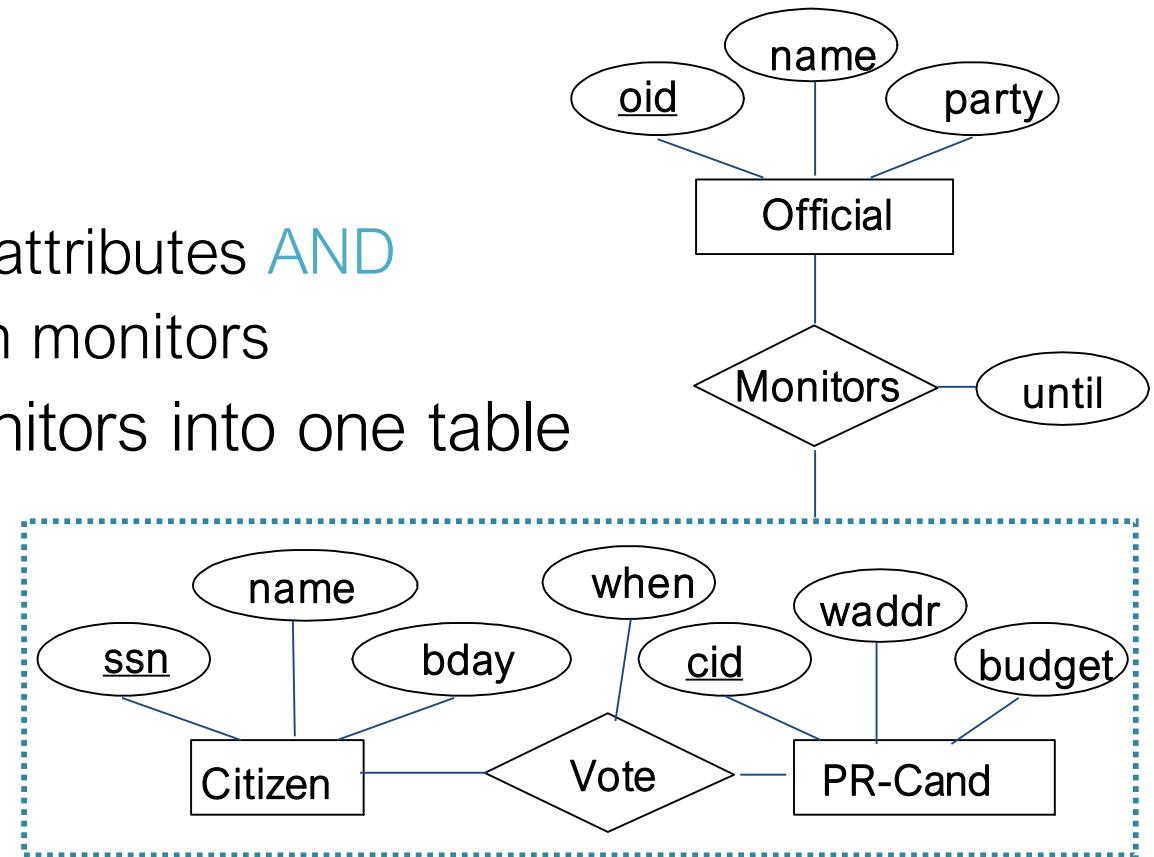
- Two relations:
  - PR-Cands (ssn, name, bday, waddr, budget)
  - Presidents (ssn, name, bday, from, to)
- Problems:
  - What if citizen is both?
    - Redundancy
  - What if citizen is neither?
    - Use General approach



# Aggregation



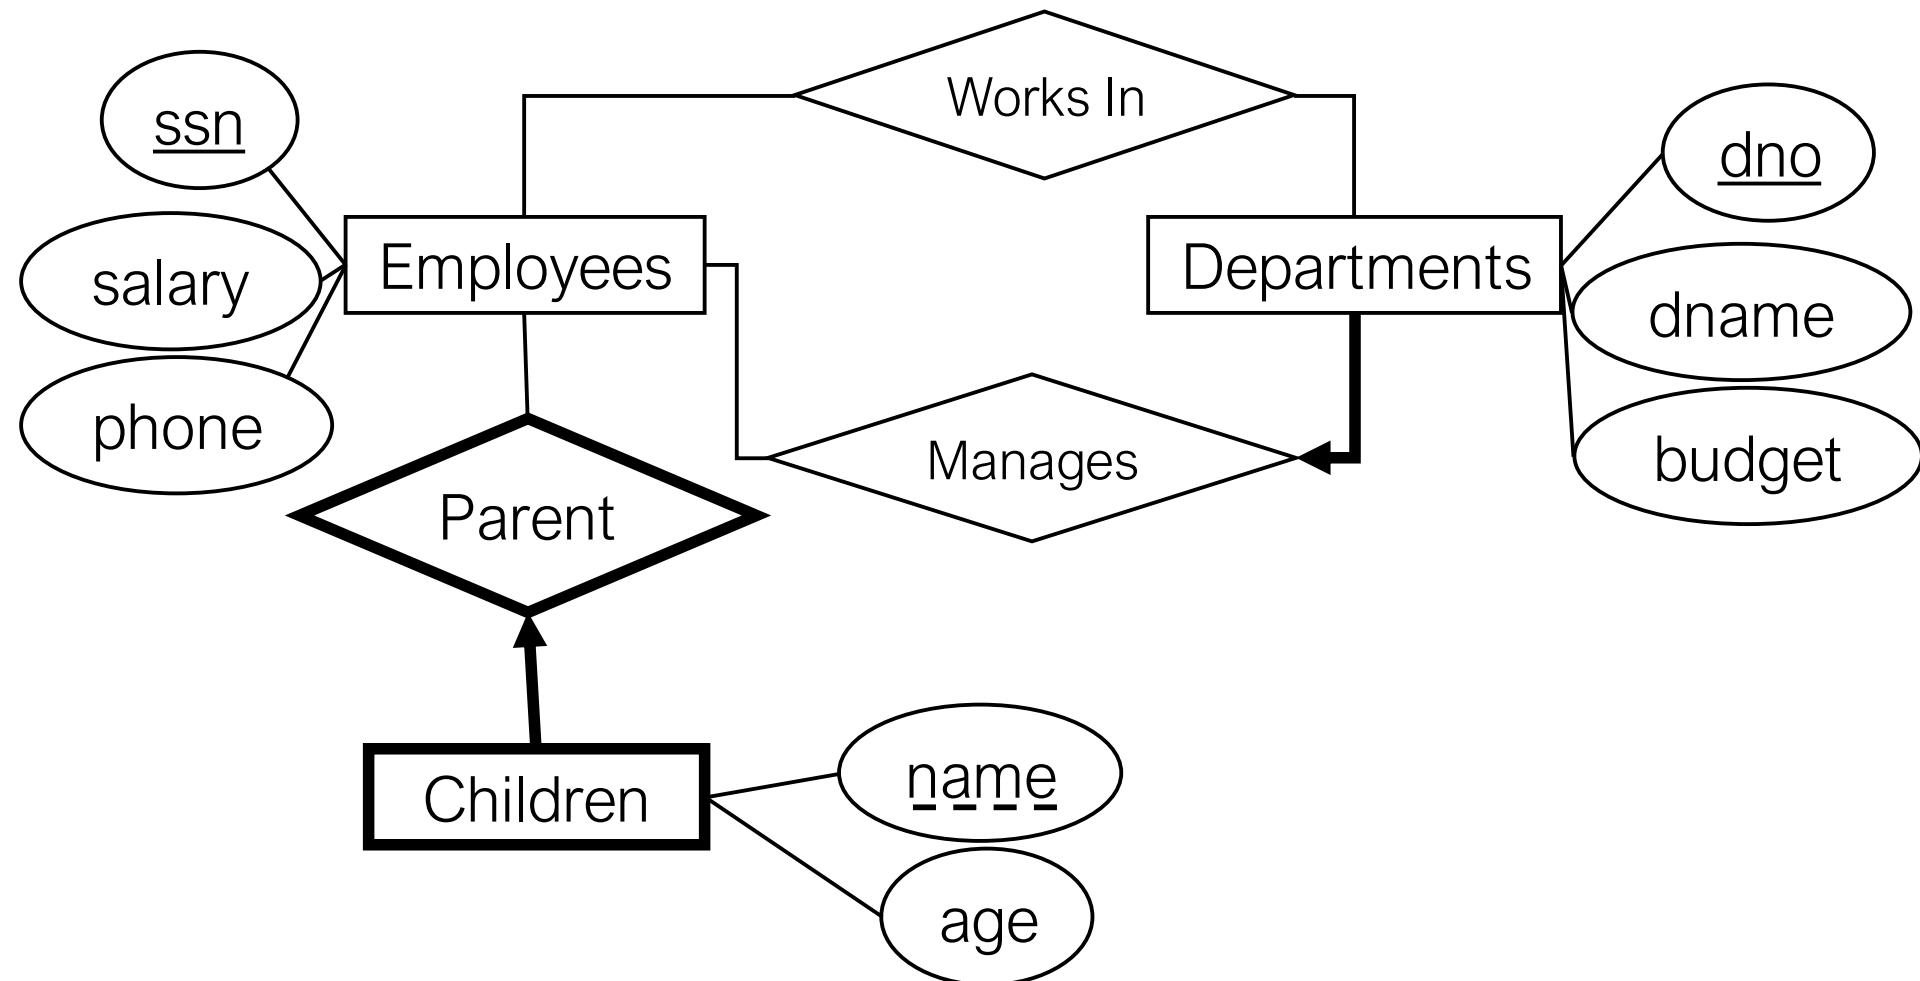
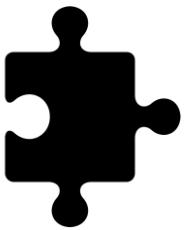
- Keep keys of **all** participating entity sets; decide primary
- Keys for Monitors
  - oid
  - (ssn, cid)
- Q: What if Vote:
  - has no descriptive attributes **AND**
  - total participation in monitors
- A: Fold Vote and Monitors into one table



# Exercise – Part 1

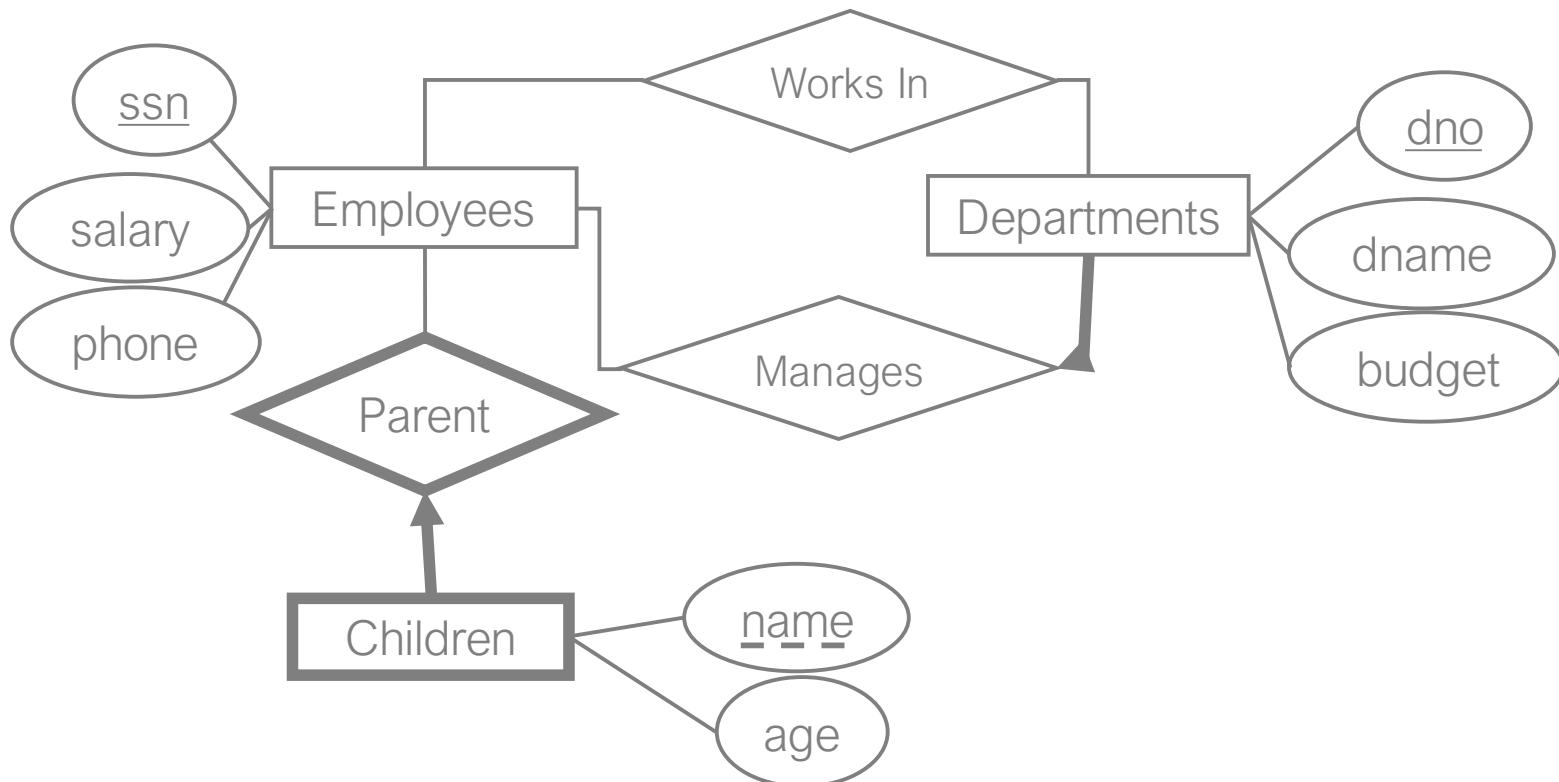
- A company database needs to store information about
  - **employees** (identified by ssn, with salary and phone attributes),
  - **departments** (identified by dno, with dname and budget attributes), and
  - **children of employees** (with name and age attributes).
- Employees work in (zero or more) departments
- Each department is managed by exactly one employee
- A child must be identified uniquely by name when the parent (who is an employee; assume only one parent works for the company) is known.
- We are not interested in information about a child once the parent leaves the company.
- **Draw an ER diagram that captures this information**

# ER Diagram (one solution)

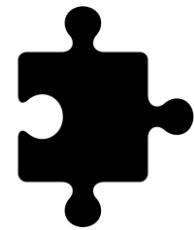


# Exercise – Part 2

- Write SQL statements to create the corresponding relations, and to capture as many of the constraints as possible.



# SQL DDL – One solution



```
CREATE TABLE Employees (  
    ssn INTEGER,  
    salary REAL,  
    phone CHAR(10),  
    PRIMARY KEY(ssn))
```

```
CREATE TABLE Works (  
    ssn INTEGER,  
    dno INTEGER,  
    PRIMARY KEY (ssn, dno),  
    FOREIGN KEY (ssn)  
        REFERENCES employees,  
    FOREIGN KEY (dno)  
        REFERENCES departments)
```

```
CREATE TABLE Departments (  
    dno INTEGER,  
    dname CHAR(20),  
    budget real,  
    manager INTEGER NOT NULL,  
    PRIMARY KEY (dno),  
    FOREIGN KEY (manager)  
        REFERENCES employees)  
    REFERENCES employees)
```

```
CREATE TABLE Children (  
    name CHAR(20),  
    age REAL,  
    parent INTEGER NOT NULL,  
    PRIMARY KEY(name, parent),  
    FOREIGN KEY(parent)  
        REFERENCES employees  
        ON DELETE CASCADE)
```

# Integrity Constraints

- Describe conditions that must be satisfied by every legal instance
- Types of integrity constraints
  - Domain constraints
  - Primary key constraints
  - Foreign key constraints
  - **General constraints**



# Table Constraints (5.7 in book)

- More general than key constraints
- Can use a query to express constraint
  - Constraints checked each time table updated
  - CHECK constraint always true for empty relation

```
CREATE TABLE    Sailors
(   sid        INTEGER,
    sname      CHAR(10),
    rating     INTEGER,
    age        REAL,
    PRIMARY KEY  (sid),
    CHECK    ( rating >= 1
              AND rating <= 10)
) ;
```

# Try it out in sqlplus or sqlite

```
CREATE TABLE Sailors
( sid      INTEGER,
  sname    CHAR(10),
  rating   INTEGER,
  age      REAL,
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10));
INSERT INTO Sailors VALUES (1, 's1', 11, 25);
```

Do you get a constraint violation error?

# Try it out in sqlplus or sqlite

```
CREATE TABLE Sailors
( sid      INTEGER,
  sname    CHAR(10),
  rating   INTEGER,
  age      REAL,
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10));
INSERT INTO Sailors VALUES (1, 's1', 11, 25);
```

Do you get a constraint violation error?

> Error: CHECK constraint failed: Sailors



# More general CHECK

- Interlake boats cannot be reserved.
- Note: these are not supported in Oracle or SQLite

```
CREATE TABLE Reserves
  ( sname  CHAR(10) ,
    bid      INTEGER,
    day      DATE,
    PRIMARY KEY (bid,day),
    CONSTRAINT noInterlakeRes
    CHECK ('Interlake' NOT IN
           ( SELECT B.bname
             FROM Boats B
            WHERE B.bid=bid)))
```



# Constraints Over Multiple Relations

- For general constraint over multiple tables, use an assertion
- Number of boats plus number of sailors is < 100

```
CREATE ASSERTION smallClub
CHECK
( (SELECT COUNT (S.sid) FROM Sailors S) +
  (SELECT COUNT (B.bid) FROM Boats B) < 100)
```

# Practical Considerations

- CHECK with subqueries and ASSERTIONS
  - Part of SQL standard (since 1992?)
  - But, they are not supported in many major databases
  - **Main concern:** Performance issues; CHECK or ASSERTION constraints over multiple are very slow
- Instead: **Triggers**
  - Most major database systems require them
  - Triggers are procedural

# Active Databases & Triggers (5.8 in the book)

**Trigger:** Procedure that starts automatically if specified changes occur to the DBMS

- Three parts:
  - Event (activates the trigger)
  - Condition (test that is run when the trigger is activated)
  - Action (what happens if the trigger runs)
    - Before and After Triggers
- Trigger Execution
  - Row-level Triggers: Once per modified row
  - Statement-level Triggers: Once per SQL statement



# Oracle Trigger Example

- First trigger executed before the activating statement, second executes after the activating statement.
- In combination with:
  - FOR EACH ROW - execute once per modified record
  - (default) - execute once per activating statement
- Activating statements:
  - INSERT
  - DELETE
  - UPDATE

```
CREATE TRIGGER init_count
BEFORE INSERT ON Student          /* Event */
DECLARE
    count INTEGER;
BEGIN
    /* Action */
    count := 0
END;

CREATE TRIGGER incr_count
AFTER INSERT ON Student           /* Event */
FOR EACH ROW
WHEN (new.age >= 18)            /* Condition */
*/
BEGIN
    /* Action */
    count := count + 1;
END;
```

# Recall CASCADE constraints

```
CREATE TABLE Athlete
(aid INTEGER PRIMARY KEY,
name CHAR(30),
country CHAR(20),
sport CHAR(20));
```

```
CREATE TABLE Olympics
(oid INTEGER PRIMARY KEY,
year INTEGER,
city CHAR(20));
```

```
CREATE TABLE Compete
(aid INTEGER,
oid INTEGER,
PRIMARY KEY (aid, oid),
FOREIGN KEY (aid) REFERENCES Athlete
    ON DELETE CASCADE
FOREIGN KEY (oid) REFERENCES Olympics
);
```



# CASCADE Using Triggers

```
CREATE TABLE Compete
  (aid INTEGER,
   oid INTEGER,
   PRIMARY KEY (aid, oid),
   FOREIGN KEY (aid) REFERENCES Athlete
   FOREIGN KEY (oid) REFERENCES Olympics);
```

```
CREATE OR REPLACE TRIGGER cascade_on_delete
AFTER DELETE ON Athlete
FOR EACH ROW
BEGIN
  DELETE FROM Compete
  WHERE Compete.aid = :OLD.aid;
END;
```





# Trying out triggers

- Try out the file `athlete_trigger_cascade.sql` in sqlplus
- Also try it out in sqlite to see if it supports triggers the same way
- Try removing a row from athlete. Does it cascade to Compete via the trigger?
- Try dropping the trigger:
  - `DROP TRIGGER triggername;`
- What happens now on deleting a row from Athlete?

# Triggers: Pitfalls and Pain

- Triggers can be recursive!
  - Chain of triggers can be hard to predict, which makes triggers difficult to understand and debug
- Errors with “mutating” table
  - A table that is currently being modified by an UPDATE, DELETE, or INSERT statement, or a table that might be updated by the effects of a DELETE CASCADE constraint
  - The session that issued the triggering statement cannot query or modify a mutating table
  - Used to prevent a trigger from seeing **inconsistent data**