

SQL Queries

Chapter 5

[https://lagunita.stanford.edu/courses/
DB/SQL/SelfPaced/about](https://lagunita.stanford.edu/courses/DB/SQL/SelfPaced/about)

Structured Query Language

- Create a Table
 - Add new records
 - Retrieve records
 - Update records
 - Delete records
-
- Create a View
 - Update a View

Create (Delete/Modify)
Insert
Select
Update
Delete

DDL

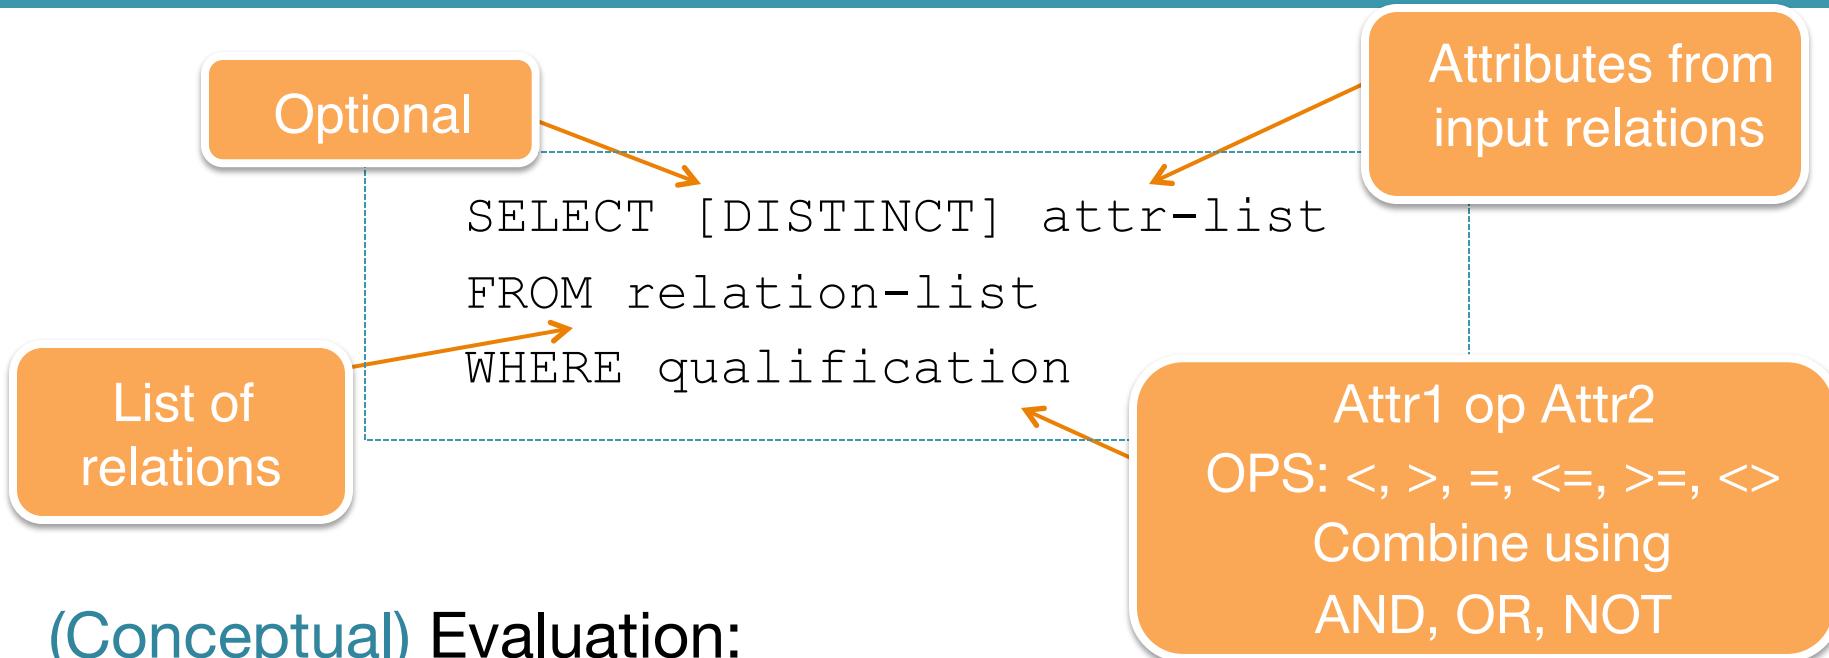
Create
Update



SQL Query Language

- Implements relational algebra...
Select, Project, Join, Set operators “Relationally complete”
- And so much more...
 - Correlated subqueries
 - Ordering of results
 - Aggregate queries (e.g., SUM, MAX, AVG)
 - Three-valued logic for NULL values
 - Etc.

Basic SQL Query



(Conceptual) Evaluation:

1. Take cross-product of relation-list
2. Select rows satisfying qualification
3. Project columns in attr-list
(eliminate duplicates only if DISTINCT)

Optimizer chooses efficient plan!

Example of Basic Query



- Schema:

Sailors (*sid*, *sname*, *rating*, *age*)

Boats (*bid*, *bname*, *color*)

Reserves (*sid*, *bid*, *rday*)

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

- Find the names of sailors who have reserved boat #103

```
SELECT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid AND R.bid = 103;
```

$$\pi_{sname}((\sigma_{bid=103} \text{Reserves}) \bowtie \text{Sailors})$$


Eliminating Duplicates



```
SELECT DISTINCT sname  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid;
```



Another Example

- Schema:

Sailors (**sid**, sname, rating, age)

Boats (**bid**, bname, color)

Reserves (**sid**, **bid**, rday)

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

- Find the colors of boats reserved by a sailor named Rusty

```
SELECT B.color
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND
      S.sname = 'Rusty';
```

Note on Range Variables



Needed when same relation appears twice in FROM clause

```
SELECT S1.sname, S2.sname  
FROM Sailors S1, Sailors S2  
WHERE S1.age > S2.age;
```



What does this
Query compute?



Good style to always use range variables anyway...



Another Example

- Find pairs of sailors where the first one has half the rating of the second one:

```
SELECT S1.sname AS name1, S2.sname AS name2  
FROM Sailors S1, Sailors S2  
WHERE 2*S1.rating = S2.rating;
```



Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red



Incrementing the result

- Compute increments for the ratings of persons who have sailed two different boats on the same day:

```
SELECT S.sname, S.rating+1 AS rating  
FROM Sailors S, Reserves R1, Reserves R2  
WHERE S.sid=R1.sid AND S.sid=R2.sid AND  
R1.day=R2.day AND R1.bid<>R2.bid;
```



Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

sid	bld	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

bld	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Showing JOINS Explicitly



Find the names of sailors who have reserved boat #103.
All the following are equivalent:

$$\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$$

Cross-product syntax:

```
SELECT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid AND  
R.bid = 103;
```

Join syntax:

```
SELECT S.sname  
FROM Sailors S JOIN Reserves  
R ON S.sid = R.sid  
WHERE R.bid = 103;
```

Natural join syntax (joins on common attributes):

```
SELECT S.sname  
FROM Sailors S NATURAL JOIN Reserves R  
WHERE R.bid = 103;
```



INNER Joins

The join we just saw is also called an **INNER JOIN**
(We will see outer joins shortly)

Join syntax:

```
SELECT S.sname  
FROM Sailors S JOIN Reserves R ON S.sid = R.sid  
WHERE R.bid = 103;
```

Eqvt. Inner join syntax:

```
SELECT S.sname  
FROM Sailors S INNER JOIN Reserves R ON  
    S.sid = R.sid  
WHERE R.bid = 103;
```

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid = R.sid AND
R.bid = 103;
```

Example of Basic Query:

Reserves

sid	bid	rday
22	101	10/10
58	103	11/12

Sailors

sid	sname	rating	age
22	Dustin	7	45
58	Rusty	10	35
31	Lubber	8	55

Reserves x Sailors

sid	bid	rday	sid	sname	rating	age
22	101	10/10	22	Dustin	7	45
22	101	10/10	58	Rusty	10	35
22	101	10/10	31	Lubber	8	55
58	103	11/12	22	Dustin	7	45
58	103	11/12	58	Rusty	10	35
58	103	11/12	31	Lubber	8	55

ORDER BY Clause



Helps sort the result for presentation

Attribute(s) in ORDER BY clause (must be) in SELECT list

Find the names
and ages of all sailors,
in increasing order of age

```
SELECT S.sname, S.age  
FROM Sailors S  
ORDER BY S.age [ASC]
```

Find the names
and ages of all sailors,
in decreasing order of age

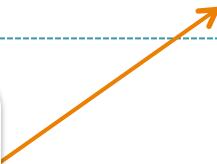
```
SELECT S.sname, S.age  
FROM Sailors S  
ORDER BY S.age DESC
```



ORDER BY Clause

```
SELECT S.sname, S.age, S.rating  
FROM Sailors S  
ORDER BY S.age ASC, S.rating DESC
```

What does this query compute?



Find the names, ages, and rankings of all sailors.

Sort the result in increasing order of age.

If there is a tie, sort those tuples in decreasing order of rating.

Set Operators

- UNION (eliminates duplicates)
- UNION ALL (keeps duplicates)
- INTERSECT
- EXCEPT or MINUS (set difference)



Union Example

Find names of sailors who have reserved a red or a green boat.

Try without UNION



Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

sid	bld	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

bld	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red



Union Example

Find names of sailors who have reserved a red or a green boat.

```
SELECT DISTINCT S.sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
AND (B.color = 'red' OR B.color = 'green');
```



```
SELECT S.sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
UNION
SELECT S.sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid and R.bid = B.bid AND B.color = 'green';
```



Intersect

Find names of sailors who have reserved a red and a green boat.

```
SELECT S.sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid  
AND (B.color = 'red' AND B.color = 'green');
```

What is wrong with the above query?



Intersect



Find names of sailors who have reserved a red and a green boat.

```
SELECT S.sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid  
AND (B.color = 'red' AND B.color = 'green');
```

What is wrong with the above query?





Intersect

Find names of sailors who have reserved a red and a green boat.

```
SELECT S.sname  
FROM Sailors S, Reserves R1, Boats B1,  
      Reserves R2, Boats B2  
WHERE S.sid = R1.sid AND R1.bid = B1.bid  
      AND S.sid = R2.sid AND R2.bid = B2.bid  
      AND B1.color = 'red' AND B2.color = 'green'
```

```
SELECT S.sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'  
INTERSECT  
SELECT S.sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid and R.bid = B.bid AND B.color = 'green';
```

Bug?



Set Difference Example

Find tuples in A that are not in B

```
SELECT * FROM A  
MINUS  
SELECT * FROM B;
```



Set Difference Example

Find sids of sailors who have reserved red, but not green boats.

```
SELECT S.sid  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid  
      AND B.color = 'red'
```

EXCEPT

```
SELECT S.sid  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid  
      AND B.color = 'green'
```

Simpler?



Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Set Difference Example



Find sids of sailors who have reserved red, but not green boats.

```
SELECT R.sid  
FROM Reserves R, Boats B  
WHERE R.bid = B.bid AND B.color = 'red'
```

EXCEPT

```
SELECT R.sid  
FROM Reserves R, Boats B  
WHERE R.bid = B.bid AND B.color = 'green'
```

Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

sid	bld	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

More Set Comparison Operators

- Set comparisons:
 - $attr \text{ IN } R$: true if R contains attr
 - $\text{EXISTS } R$: true if R is not an empty relation
 - $\text{UNIQUE } R$: true if no duplicates in R
 - You can use NOT with these, e.g. NOT EXISTS
- Also available ANY or ALL: (op is $<$, \leq , $>$, \geq , $=$, \neq)
 - $attr > \text{ANY } R$: some element of R satisfies the condition that attr > that element
 - $attr < \text{ALL } R$: all elements of R satisfy the condition that attr < element



Nested Queries

Query with another query embedded inside

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid IN  
(SELECT R.sid  
FROM Reserves R  
WHERE R.bid=103)
```

What does this query compute?

Conceptual evaluation:

For each row of Sailors, evaluate the subquery over reserves.

To find sailors who have not reserved 103, use NOT IN.

Over-Use of Nesting



- Common error by novice SQL programmers
- Query optimizers not as good at optimizing queries across nesting boundaries
- Try hard first to write non-nested

```
SELECT DISTINCT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid AND R.bid = 103;
```



Example

- Q1: What does this query compute?



```
SELECT S.sid  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                      FROM Sailors S2  
                      WHERE S2.name = 'John' );
```

- Q2: Rewrite the query without using a nested query



Example

- Q1: What does this query compute?



```
SELECT S.sid  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                      FROM Sailors S2  
                      WHERE S2.name = 'John' );
```

- Q2: Rewrite the query without using a nested query

```
SELECT DISTINCT S.sid  
FROM Sailors S, Sailors S2  
WHERE S.rating > S2.rating AND S2.name = 'John' ;
```



Example

Find sailors (all their info) whose rating is greater than that of **all** sailors called Horatio:



```
SELECT  *
FROM   Sailors S
WHERE  S.rating > ALL  (SELECT  S2.rating
                         FROM   Sailors S2
                         WHERE  S2.sname='Horatio')
```

Aggregate Operators



```
SELECT COUNT (*)  
FROM Sailors S
```

```
SELECT COUNT  
  (DISTINCT S.name)  
FROM Sailors S
```

```
SELECT AVG(S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating = (SELECT MAX(S2.rating)  
                  FROM Sailors S2)
```

COUNT (*)
COUNT ([DISTINCT] A)
SUM ([DISTINCT] A)
AVG ([DISTINCT] A)
MAX (A) *Can use Distinct*
MIN (A) *Can use Distinct*

single column*

```
SELECT AVG(DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10
```

Sailors

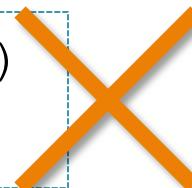
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Aggregate Query - Example



Find name and age of the oldest sailor(s)

```
SELECT S.sname, MAX (S.age)  
FROM Sailors S
```



```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age = (SELECT MAX (S2.age)  
               FROM Sailors S2)
```

How many tuples
in the result?

**Does not work in
sqlite3**

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age >= ALL (SELECT S2.age  
                     FROM Sailors S2)
```

GROUP BY



Conceptual evaluation

- Partition data into groups according to some criterion
- Evaluate the aggregate for each group

Example:

For each rating level, find the age of the youngest sailor

```
SELECT MIN (S.age), S.rating  
FROM Sailors S  
GROUP BY S.rating
```

How many tuples
in the result?



GROUP BY and HAVING

```
SELECT [DISTINCT] target-list  
FROM relation-list  
WHERE qualification  
GROUP BY grouping-list  
HAVING group-qualification
```

Target-list contains:

- 1. Attribute names**
(subset of grouping-list)
- 2. Aggregate operations**
e.g. min(age)

Conceptual Evaluation:

1. Eliminate tuples that don't satisfy qualification
2. Partition remaining data into groups
3. Eliminate groups according to group-qualification
4. Evaluate aggregate operation(s) for each group

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors



```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) >= 2
```

Sailors

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
71	Zorba	10	16.0
64	Horatio	7	35.0
29	Brutus	1	33.0
58	Rusty	10	35.0

rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

rating	age
7	35.0

Answer
relation

NULL Values in SQL

- NULL represents ‘unknown’ or ‘inapplicable’
- Query evaluation complications
 - Q: Is $(\text{rating} > 10)$ true when rating is NULL?
 - A: Condition evaluates to ‘unknown’ (not T or F)
- What about AND, OR connectives?
 - Need 3-valued logic
- WHERE clause eliminates rows that don’t evaluate to true

p	q	p AND q	p OR q
T	T	T	T
T	F	F	T
T	U	U	T
F	T	F	T
F	F	F	F
F	U	F	U
U	T	U	T
U	F	F	U
U	U	U	U



NULL Values Example

Sailors

sid	sname	rating	age
22	Dustin	7	45
58	Rusty	10	NULL
31	Lubber	8	55

```
SELECT sname  
FROM sailors  
WHERE age > 45 OR age <= 45
```

What does this query return?



NULL Values in Aggregates



NULL values are generally ignored when computing aggregates

Sailors

sid	sname	rating	age
22	Dustin	7	45
58	Rusty	10	NULL
31	Lubber	8	55

```
SELECT AVG(age)  
FROM sailors
```

Returns 50!



Inner Joins

Find the names of sailors who have reserved boat #103.
All the following are equivalent:

$$\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$$

Cross-product syntax:

```
SELECT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid AND  
R.bid = 103;
```

Join syntax:

```
SELECT S.sname  
FROM Sailors S JOIN Reserves  
R ON S.sid = R.sid  
WHERE R.bid = 103;
```

Natural join syntax (joins on common attributes):

```
SELECT S.sname  
FROM Sailors S NATURAL JOIN Reserves R  
WHERE R.bid = 103;
```



Outer Joins

Sailors

sid	sname	rating	age
22	dustin	7	45.0
58	rusty	10	35.0

Reserves

sid	bid	day
22	101	10/10/99

```
SELECT S.sid, R.bid  
FROM Sailors S NATURAL LEFT [OUTER]  
JOIN Reserves R
```

Result

sid	bid
22	101
58	null

Similarly:

- Right Outer Join
- Full Outer Join

Note: OUTER is default, when using LEFT, RIGHT, or FULL

More Outer Joins

```
SELECT S.sid, R.bid  
FROM Sailors S RIGHT[OUTER] JOIN  
Reserves R ON S.sid=R.sid;
```

Sailors

sid	sname	rating	age
58	Rusty	10	35
31	Lubber	8	55

Result

sid	bid
null	101
58	103

Reserves

sid	bid	rday
22	101	10/10
58	103	11/12

JOIN Syntax with Multiple Tables

Sailors

sid	sname	rating	age
22	dustin	7	45.0
58	rusty	10	35.0

Reserves

sid	bid	day
22	101	10/10/99

```
SELECT S.sname, B.bname  
FROM Sailors S JOIN Reserves R ON (S.sid = R.sid)  
                  JOIN Boats B ON (R.bid = B.bid)  
WHERE S.name = 'dustin';
```

Similarly:

- RIGHT [OUTER] JOIN ON...
- FULL [OUTER] JOIN ON...
- LEFT [OUTER] JOIN on...
- NATURAL JOINS (outer and inner)

For each red boat, find the number of reservations for this boat*



```
SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

```
SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
GROUP BY B.bid  
HAVING B.color = 'red'
```

Would this work?
Note: one color per bid





Subtle Errors

Find the sid of sailors who have reserved exactly one boat

```
SELECT S1.sid FROM Sailors S1  
MINUS  
SELECT R1.sid  
FROM Reserves R1, Boats B1, Reserves R2, Boats B2  
WHERE R1.sid=R2.sid AND R1.bid=B1.bid  
      AND R2.bid=B2.bid AND R1.bid <> R2.bid;
```

There is a subtle error in the above



Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red



Error Fixed

Find the sid of sailors who have reserved exactly one boat

```
SELECT R3.sid FROM Reserves R3  
MINUS  
SELECT R1.sid  
FROM Reserves R1, Boats B1, Reserves R2, Boats B2  
WHERE R1.sid=R2.sid AND R1.bid=B1.bid  
      AND R2.bid=B2.bid AND R1.bid <> R2.bid;
```

Error Fixed: Another Solution



Find the sid of sailors who have reserved exactly one boat

```
SELECT S.sid  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid = R.sid AND B.bid = R.bid  
GROUP BY S.sid  
HAVING COUNT(*) = 1;
```



Intersect on Non-Key

Find the names of sailors who have reserved a red and a green boat

```
SELECT S.sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'  
INTERSECT  
SELECT S.sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid and R.bid = B.bid AND B.color ='green'
```

What is wrong with the above query?

Fix it: (1) by using a view
and (2) without a view.



Error Fixed



Find the names of sailors who have reserved a red and a green boat

```
CREATE VIEW RedGreenSailors AS
SELECT S.sid
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
INTERSECT
SELECT S.sid
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid and R.bid = B.bid AND B.color = 'green';

SELECT S.sname
FROM Sailors S, RedGreenSailors R
WHERE S.sid = R.sid;
DROP VIEW RedGreenSailors;
```

Error Fixed: Another Solution



Find the names of sailors who have reserved a red and a green boat

Get rid of the VIEW

```
SELECT S.sname  
FROM Sailors S,  
(SELECT S.sid  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'  
INTERSECT  
SELECT S.sid  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid and R.bid = B.bid AND B.color = 'green')  
RedGreenSailors  
WHERE S.sid = RedGreenSailors.sid;
```

Another sol:
See Q8 in
Ch. 5, p.150

Find the age of the youngest sailor with age > 18, for each rating with at least 2 sailors (of any age)



```
SELECT S.rating, MIN (S.age) AS MINAGE  
FROM Sailors S  
WHERE S.age > 18  
GROUP BY S.rating  
HAVING 1 < (SELECT COUNT (*) FROM Sailors S2  
              WHERE S2.rating=S.rating)
```

- Subquery in the HAVING clause
- Compare this with the query where we considered only ratings with 2 sailors over 18!

Find the age of the youngest sailor with age > 18, for each rating with at least 2 sailors (of any age)



```
SELECT S.rating, MIN (S.age) AS MINAGE  
FROM Sailors S  
WHERE S.age > 18  
GROUP BY S.rating  
HAVING 1 < (SELECT COUNT (*) FROM Sailors S2  
              WHERE S2.rating=S.rating)
```

- Compare this with the query where we considered only ratings with 2 sailors over 18!

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
WHERE S.age >= 18  
GROUP BY S.rating  
HAVING COUNT (*) >= 2
```



Find ratings for which the average age is the minimum of the average age over all ratings*

Aggregate operations
cannot be nested!

WRONG: 

```
SELECT S.rating
FROM Sailors S
WHERE AVG(S.age) =
  (SELECT MIN (AVG (S2.age))
   FROM Sailors S2
   GROUP BY S2.rating)
```

Correct solution (in SQL/92 – may fail in practice):

```
SELECT T.rating, T.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) T
WHERE T.avgage = (SELECT MIN (T.avgage) FROM T);
```

If above does not work, one solution is to define T as a view



Solution Using Views

```
CREATE VIEW AVG_AGE_BY_RATING AS  
SELECT S.rating, AVG(S.age) AS avgage  
FROM Sailors S  
GROUP BY S.rating;  
  
SELECT T.rating, T.avgage  
FROM AVG_AGE_BY_RATING T  
WHERE T.avgage= (SELECT MIN(A.avgage)  
                  FROM AVG_AGE_BY_RATING A);
```