

Database Application Programming


Chapter 6 (JDBC Section)

Databases “In the Wild”

- So far, we've talked about the DBMS as a standalone system
 - Access interactively by writing SQL queries (e.g., using SQL*Plus)
- In practice, DBMS is often part of a larger software infrastructure
 - Multi-tiered system architecture
 - Access database from another program


Database “Ecosystem” (1)

“Client-Server Architecture”




Many Clients

“Client-Server Architecture”




Database “Ecosystem” (2)

“3-Tier Architecture” (Common to add more tiers, too)



Embedding SQL in Application




SQL Integration with PL

- Ugly Problem:
 - Database supports SQL queries
 - Application written in programming language (e.g. Java, C++)
 - What is the interface between the two?
- A Common Solution:
 - “Embed” SQL in host language
 - Provide an API for processing query results
- Also object-relational mapping tools
 - LINQ to SQL, Ruby on Rails...

JDBC (“Java-Database Connectivity”)

- Connect to a database using a JDBC driver
- Send queries over the JDBC connection
- Receive results into a Java ResultSet



Try things out with sqlite / Oracle

- Download JDBCTet.java and the latest jdbc driver for sqlite from <https://bitbucket.org/xerial/sqlite-jdbc/downloads>
- Alternative: Oracle driver posted with Project 2.
- Compile: javac JDBCTet.java
- Run: java –cp .:*sqlite-jdbc-3.7.2.jar* JDBCTet

Opening/Closing Connections

- Get Connection object:
 - Oracle requires passwords
 - Sqlite3 is file-based and does not
 - Connection conn =
- Always close connections before quitting the program
 - conn.close();
- A cool trick in recent Java/JDBC to auto-close:

http://docs.oracle.com/javase/7/docs/technotes/guides/jdbc/jdbc_41.html

JDBC Example



```
Connection conn;  
// Obtain a connection to DB, store in conn  
// (Requires JDBC driver; See sample code in Project 2)  
  
String q = "SELECT name FROM Students WHERE GPA > 3.5";  
try {  
    Statement st = conn.createStatement();  
    ResultSet rs = st.executeQuery(q);  
  
    while (rs.next()) {  
        String name = rs.getString("name");  
        System.out.println(name);  
    }  
    rs.close();  
    st.close();  
}  
catch (SQLException e) {System.err.println(e.getMessage());}
```

Cursor retrieves rows from result one at a time

Full Javadoc for `java.sql` available online:
<http://download.oracle.com/javase/6/docs/api/>

Challenges

- DBMS and PL implement different types
 - “Impedance Mismatch”
- Need to match DB types with PL types

SQL Type	Java Type	ResultSet Method
CHAR	String	getString()
VARCHAR	String	getString()
DOUBLE	Double	getDouble()
INTEGER	Integer	getInt()
NUMBER	(depends)	(depends)

Challenges

- What computation to do in the database vs. the application program?
- Rules of Thumb:
 - Avoid fetching more data than necessary
 - “Push” data processing to the DBMS when possible

JDBC Example




What does the following code snippet do?

```
String query = "SELECT NAME, AGE FROM CITIZENS";
try {
    double sum = 0;
    double count = 0;
    Statement st = conn.createStatement();
    ResultSet rs = st.executeQuery(query);
    while (rs.next()) {
        String name = rs.getString("NAME");
        sum += rs.getDouble("AGE");
        count++;
    }
    System.out.println(sum/count);
    // be sure to close rs and st as well here
    // and in exception handlers. Omitted for brevity.
}
catch (SQLException e) {System.err.println(e.getMessage());}
```

What Happens?

Compute the average age




JDBC Example (Revised)

Push the computation “closer” to the data...
Make DBMS do processing it does well.

```
String query = "SELECT AVG(AGE) FROM CITIZENS";
try {
    Statement st = conn.createStatement();
    ResultSet rs = st.executeQuery(query);
    while (rs.next()) {
        Double avg = rs.getDouble(1);
        System.out.println(avg);
    }
    // close omitted for brevity
}
catch (SQLException e) {System.err.println(e.getMessage());}
```

What Happens Now?

Compute the average age



Statement vs. PreparedStatement

- Statement (Base class)
 - Arbitrary SQL query
 - DBMS parses and optimizes each query
- PreparedStatement
 - Parameterized SQL query
 - DBMS “pre-compiles” the query (parses, optimizes, and stores query execution plan)
 - Can be used multiple times
 - Amortizes optimization cost across multiple uses

Statement vs. PreparedStatement

```
public List<String> getNames (int age) {  
    String q = "SELECT name FROM Students WHERE age = " + age;  
    Statement st = conn.createStatement();  
    ResultSet rs = st.executeQuery(q);  
    ...  
}
```

```
public List<String> getNames (int age) {  
    String q = "SELECT name FROM Students WHERE age = ?";  
    PreparedStatement ps = conn.prepareStatement(q);  
    ps.setDouble(1, age);  
    ResultSet rs = ps.executeQuery();  
    ...  
}
```



Which is more efficient if getNames() called many times?


PreparedStatement

- PreparedStatement also does some type checking on parameters and removal of escape characters
- Helpful for preventing some security problems (e.g. SQL Injection)

Security Issues / SQL Injection

- Common vulnerability in database applications
 - SQL Injection is simple, yet surprisingly common
 - Can be prevented with defensive coding

Web Front-End



SQL Injection – Example

JDBC

```
Authenticate (String n, String p) {  
...  
String query =  
"SELECT grade FROM students WHERE name = " + n + ",  
AND password = '" + p + "'";  
...  
}
```



Input:

name = bart; password = mypwd

SQL:

```
SELECT grade FROM students WHERE name = 'bart' AND  
password = 'mypwd'
```

SQL Injection – Example

JDBC

```
Authenticate (String n, String p) {  
...  
String query =  
"SELECT grade FROM students WHERE name =  
'" + n + "' AND password = '" + p + "'";  
...  
}
```



Input:

name = lisa; password = n' OR 'x'='x


SQL:

```
SELECT grade FROM students WHERE name = 'lisa' AND  
password = 'n' OR 'x' = 'x'
```

SQL Injection – Example

JDBC

```
Authenticate (String n, String p) {  
...  
String query =  
"SELECT grade FROM students WHERE name =  
'" + n + "' AND password = '" + p + "'";  
...  
}
```



Input:

name= **foo;** password= **n'; UPDATE students SET grade= 'A**

SQL:

```
SELECT grade FROM students WHERE name = 'foo' AND  
password = 'n'; UPDATE students SET grade = 'A'
```

SQL Injection



What is a simple way to prevent this kind of attack?

Validate input!

Use a prepared statement for user input parameters!

HI, THIS IS
YOUR SON'S SCHOOL.
WE'RE HAVING SOME
COMPUTER TROUBLE.



OH, DEAR - DID HE
BREAK SOMETHING?
IN A WAY -)



DID YOU REALLY
NAME YOUR SON
Robert'); DROP
TABLE Students;-- ?




OH, YES. LITTLE
BOBBY TABLES,
WE CALL HIM.

WELL, WE'VE LOST THIS
YEAR'S STUDENT RECORDS.
I HOPE YOU'RE HAPPY.



AND I HOPE
YOU'VE LEARNED
TO SANITIZE YOUR
DATABASE INPUTS.

Stored Procedures and UDFs



Database applications you write (e.g., using JDBC) are usually located outside the DBMS.

Exceptions:

- Stored Procedures
- User-Defined Functions (UDFs)
- Stored and executed within the DBMS

Summary

- DBMS is often part of a larger software infrastructure
 - E.g., Database-backed web applications
- Integrating SQL with application code is a messy problem
 - Efficiency: Push computation “close” to data
 - Security: Validate input
 - Testing/Debugging