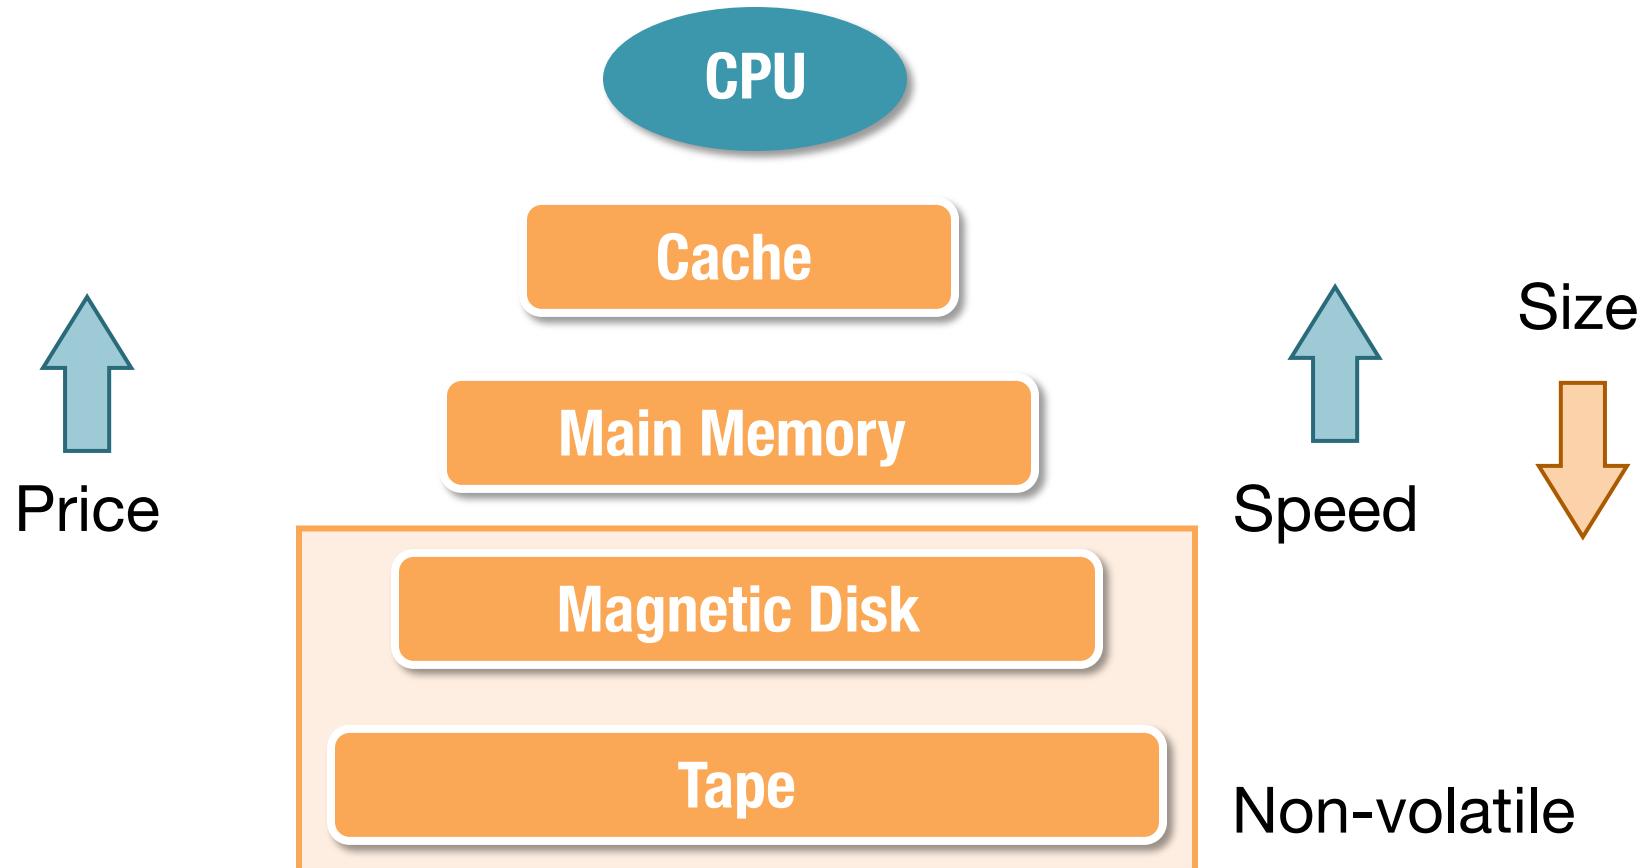


Storage and Indexing

Chapter 8

The Memory Hierarchy



Performance of Microprocessors and Memory
improving faster than disks and tapes

The Memory Hierarchy

Reading from **disk** is **many times slower** than from **memory**

- Memory access is as fast as Chuck Norris,
- whereas disk access is slower than the line at the DMV



The DMV in a nutshell.



VS

Why Not Store Everything in Main Memory?

- Too expensive: RAM costs 100-1000x Disk per GB
- Main memory is volatile: Want data to persist between runs
- Typical storage hierarchy:
 - Main memory (RAM) for currently used data
 - Disk for the main database (secondary storage)
 - Non-volatile storage
 - Tapes for archiving older versions of data (tertiary storage)
 - Sequential access devices

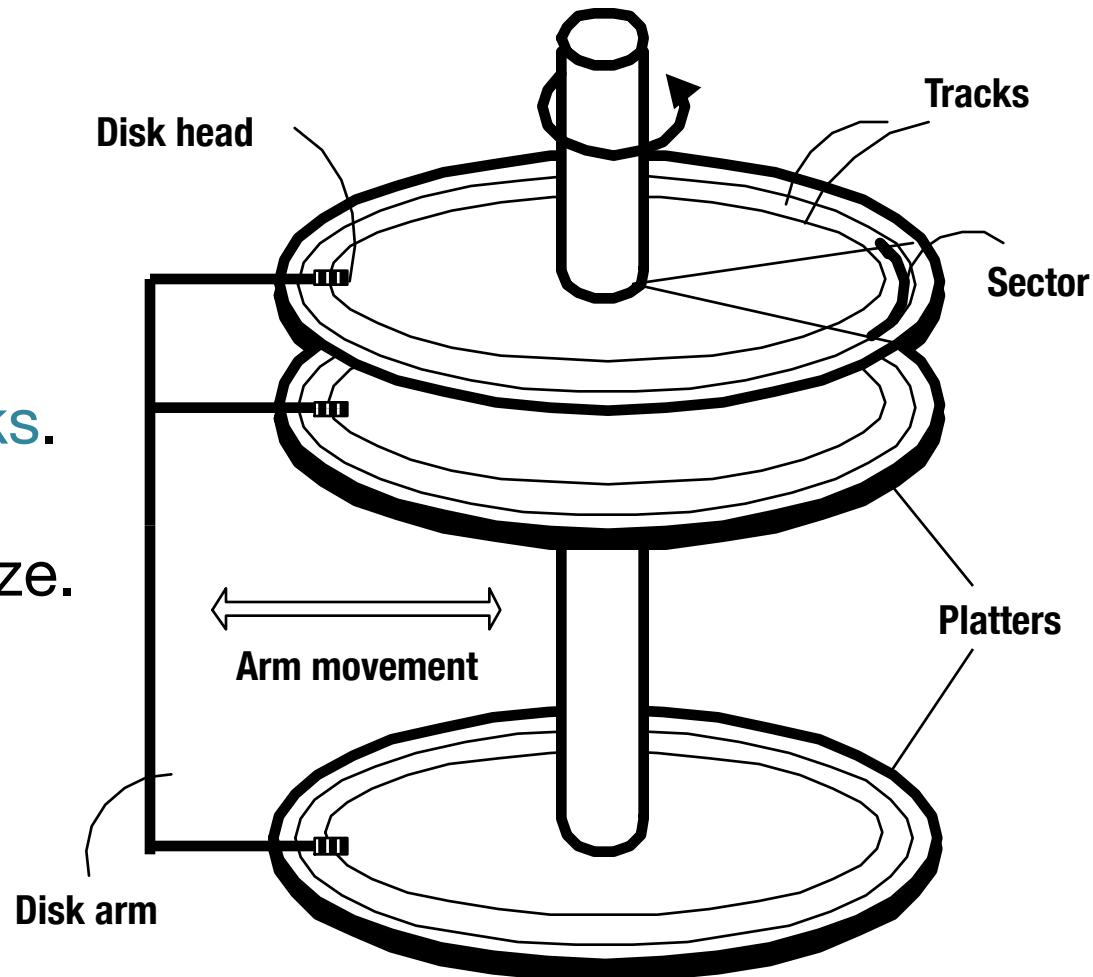


Disks

Set of tracks with same diameter called a **cylinder**.

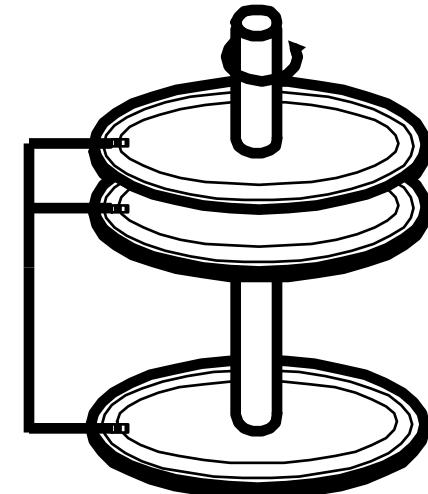
Data stored in **blocks**.
Size of block is a multiple of sector size.

Only one disk head reads or writes at a time.



Performance Implications

- Data must be in memory for DMBS to use.
- Unit of transfer is a block. Whole block must be transferred. Reading or writing a disk block is called an **I/O**.
- Disk geometry affects access time (hard drives)
 - **Seek time:** time to move disk head to appropriate track
 - **Rotational delay:** time waiting for block to move under disk head
 - **Transfer time:** time to read or write block once head is positioned



Arranging Blocks on Disk

Access time = seek time + rotational delay + transfer time

- **GOAL:** Minimize seek time and rotational delay
- ‘**Next**’ block concept:
 - blocks on same track, followed by
 - blocks on same cylinder, followed by
 - blocks on adjacent cylinder
- Arranging blocks so they are read and written **sequentially** is important to reducing time spent doing disk I/O

Comparison

- Data must be brought into memory to be read/written
- Typically: Memory size \ll Disk size
- Why?
 - Cost: memory (100x), SSD(10x), HD (1x)
 - Performance:
 - Seq I/O: memory (50x) vs. SSD (5x) vs. HD(1x)
 - Random I/O: memory (50x) vs. SSD (2-5x) vs. HD(0.01x)

Hard drives much cheaper, but also slower. Arrange in a hierarchy to get the best speed at the lowest cost.

Disk Properties

- Comparison with main memory
 - non-volatile (disks) vs. volatile (main memory)
 - Both random-access
 - Unit of access: one or more blocks (disks) vs. a byte or word
 - Speed: very slow (disk) vs. fast (memory)
- Another kind of drive: solid state
 - Cost/MB: 10x over hard drives currently.
 - Can provide lower latencies and higher throughput than hard drives
 - Access times independent of block placement
- Hybrid non-volatile storage:
 - Use a combination of solid state drives and hard drives

Disk Performance

- Traditional disks
 - Sequential I/O faster than random I/O
- Solid-state drives becoming popular
 - Same speed: sequential or random I/O (Layout not relevant)
 - 10x higher cost, but lower latencies
- Most databases today still use traditional disks, but could change over time
- Hybrid drives that use solid state drives as a cache also becoming common

Pages and Records

- Databases store data on disks
- Unit of information for disks is a **page**
 - **Physical abstraction**
 - Page size is typically 4KB to 16KB
- How are tables stored in pages?
 - A table is stored as a *file of records*
 - **Records** are **logical** units
 - Records stored in pages
 - The term “**file of records**” here simply means a **named set of pages**, not necessarily an OS file
- Typically, **one page has multiple records**
- One table (file) will typically require multiple pages

Check Your Understanding

- Why do databases need disks? Why not just use main memory?
- Which has the lowest per GB cost?
 - Main memory
 - Solid state drives
 - Magnetic drives
- Of the above, for best speed, non-volatile storage, and highest capacity, which of the above would you use?



Operations on File - Example

- Employees (Name, Age, Salary)
- Operations
 - **Scan:** Fetch all employees from disk
 - **Equality Search:** age = 21
 - **Range Selection:** age \geq 18 AND age $<$ 65
 - Insert a record
 - Delete a record



How do we organize the file
to accomplish these tasks efficiently?

File Organization

- Each record in a file has a unique Record ID (RID), which is sufficient to locate the record on the disk
 - e.g. an RID may be (page#, offset, length)
- Some methods of arranging **file of records** on disk
 - **Heap (random order) files:** No particular order defined for records
 - **Sorted Files:** Records sorted based on one or more attributes
 - **Indexes:** Organize records using trees or hashing

Indexes

- A **data structure** that organizes data records on disk to optimize certain operations.
- Speed up selections on the **search key** field of the index (denoted **k**)
 - Any **subset of the fields of a relation** can be the search key for an index on the relation
 - **Search key** is **not** the same as **key** (minimal set of fields that uniquely identify a record in a relation)
- An index file contains a collection of **data entries k^*** for each **search key value k**
 - k^* should allow us to get to the record contents

Data Entries k^*

- A data entry k^* must give us a way to get to the data for k :
- Alternative 1: Data entry k^* is an actual record (with search key k)
 - Index File == File of Records
- Alternative 2: Data entry k^* is (k, rid) pair, where rid refers to a record with search key k
 - Actual data records stored in a different file
- Alternative 3: Data entry k^* is $(k, \text{rid-list})$ pair, where rid-list refers to list of records with search key k

How many indexes can use Alternative 1?



Choosing Among Alternatives

- Only one index can use alternative 1
 - Otherwise, you would get redundancy
 - Other indexes can be Alternative 2 or 3
- Alternative 3 is more compact than alternative 2, but leads to variable-length index entries

Check Your Understanding



Can search key k be a composite value,
e.g. pair of values or attributes?

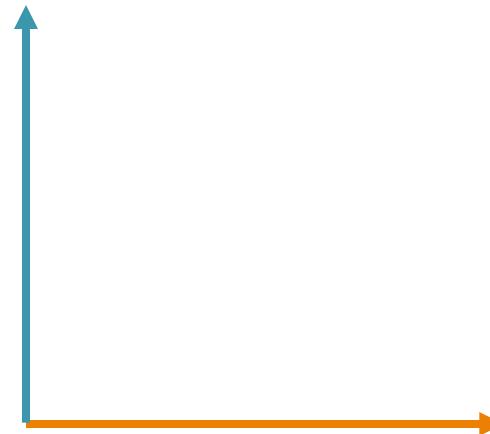
Can a file (relation) have more than one index?

Can a file (relation) have more than one index
using data entry alternative 1?

Index Design Space

Organization Structure for k^*

- Hash-based
 - (+) Equality search
- Tree-based
 - (+) Range, equality search
 - B+Tree (dynamic)
 - ISAM (static)



Data Entry (k^*) Contents

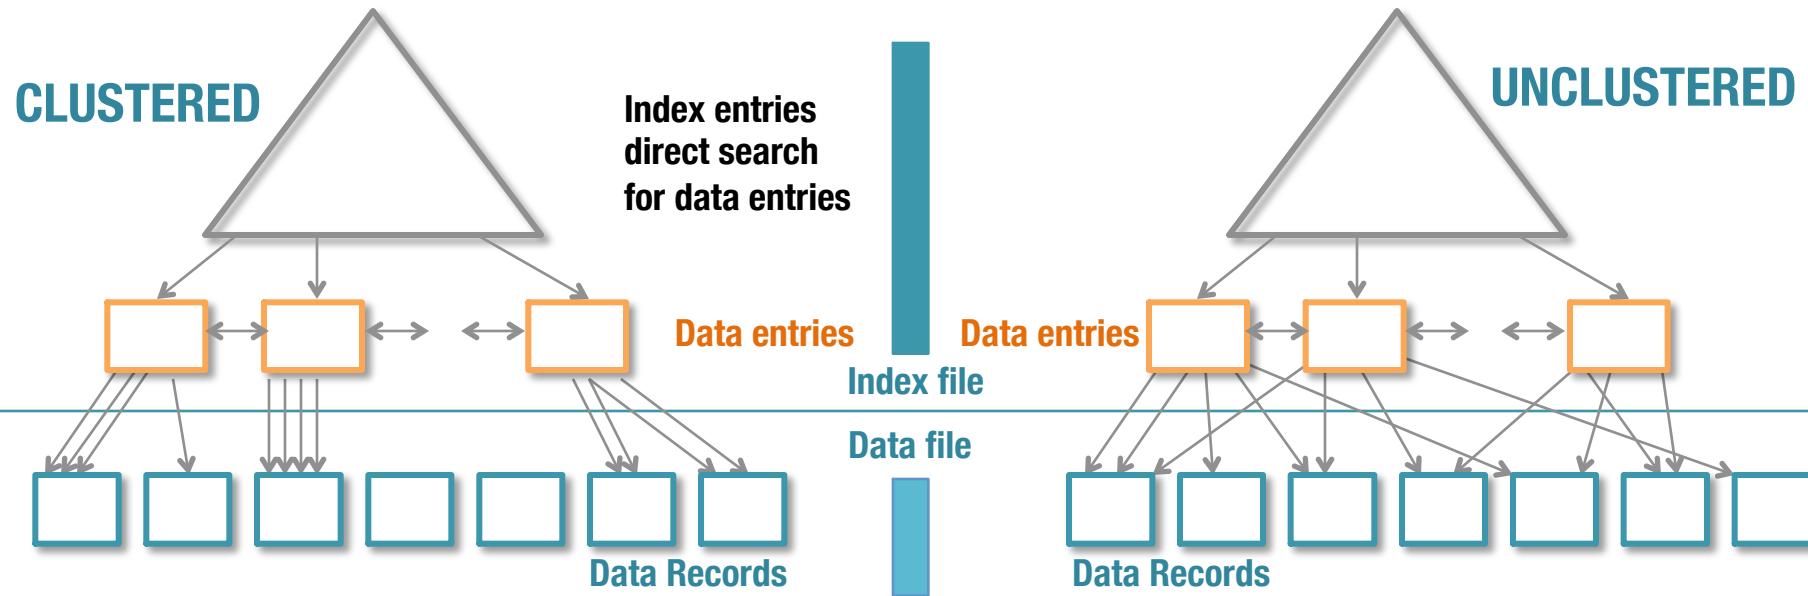
1. Actual Data record
 $\text{index} = \text{file}$
2. $\langle k, \text{rid} \rangle$
actual records in a diff file
3. $\langle k, \text{list of rids} \rangle$

Indexing Terminology



- **Primary vs. secondary index:** If search key contains primary key, then called primary index
 - Otherwise, called secondary index
- **Clustered vs. unclustered index:** If order of data records is the same as, or ‘close to’, order of data entries, then called clustered index
 - Alternative 1 implies clustered index. A file that uses Alternative 1 is also called a *clustered file*
 - But not every clustered index uses Alternative 1
- A file can be clustered on at most one search key

Clustered vs. Unclustered Index



- Suppose: data records in heap file, index with Alt. 2
- To build **clustered** index, first sort the Heap file

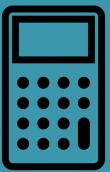


Suppose you have an index on Age,
and you want to do a range selection (e.g., $\text{Age} \geq 18$)
Which index would you prefer? Why?

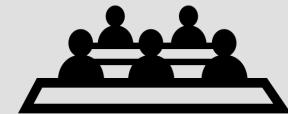
Cost of a Search Query

- Cost of Page I/O >> memory read
- # of pages read more important than # of records read
- Searching a heap (unordered)
 - All pages must be read

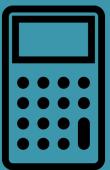
Search on Unclustered Index



- Let's say you want to **find records** of U. of Michigan **students who are 18 years old**.
- Unclustered index on age
- 25% of students are 18 years old
- Each page can hold 10 records
- 100,000 records
- How many pages are you likely to read?
 - Close to 100%, i.e. 10,000 pages?
 - Close to 25%, i.e. 2,500 pages?



Search on Clustered Index



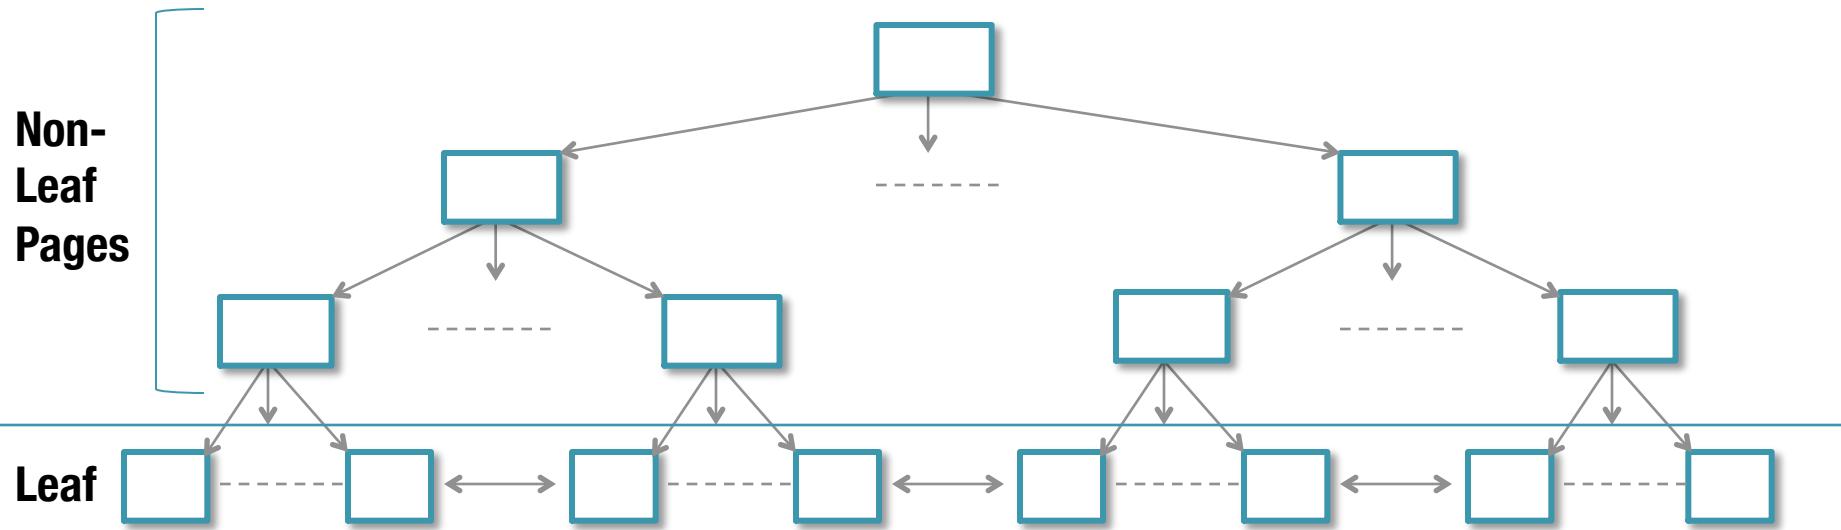
- Let's say you want to **find records** of U. of Michigan **students who are 18 years old**.
- **Clustered index** on age
- 25% of students are 18 years old
- Each page can hold 10 records
- 100,000 records
- What percent of pages are you likely to read, taking advantage of the index?
 - Close to 100%, i.e. 10,000 pages?
 - Close to 25%, i.e. 2,500 pages?



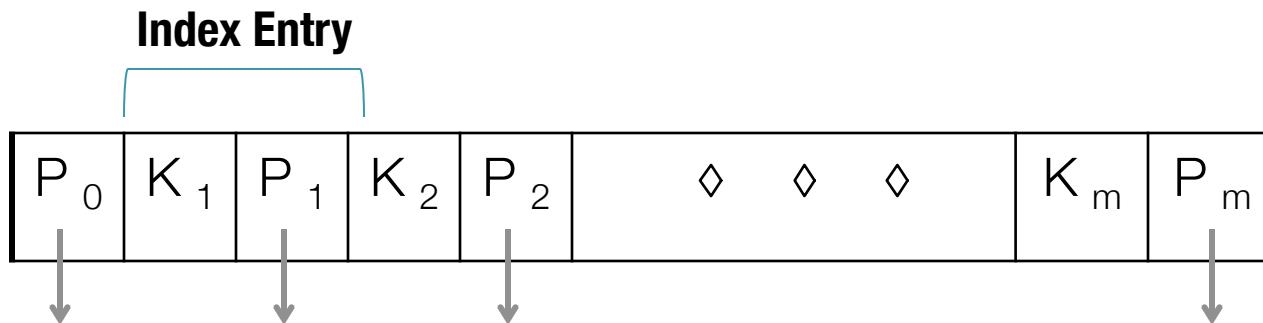
Search on Clustered Index

- In practice, pages are often only around 2/3rd full because deleting records will leave holes in each page over time 
- Thus, the file will occupy around 15,000 pages = $10,000 * 3/2$
- So, in the last problem, number of pages fetched will be approximately:
 - 3,750 pages (25% of 15,000 pages)

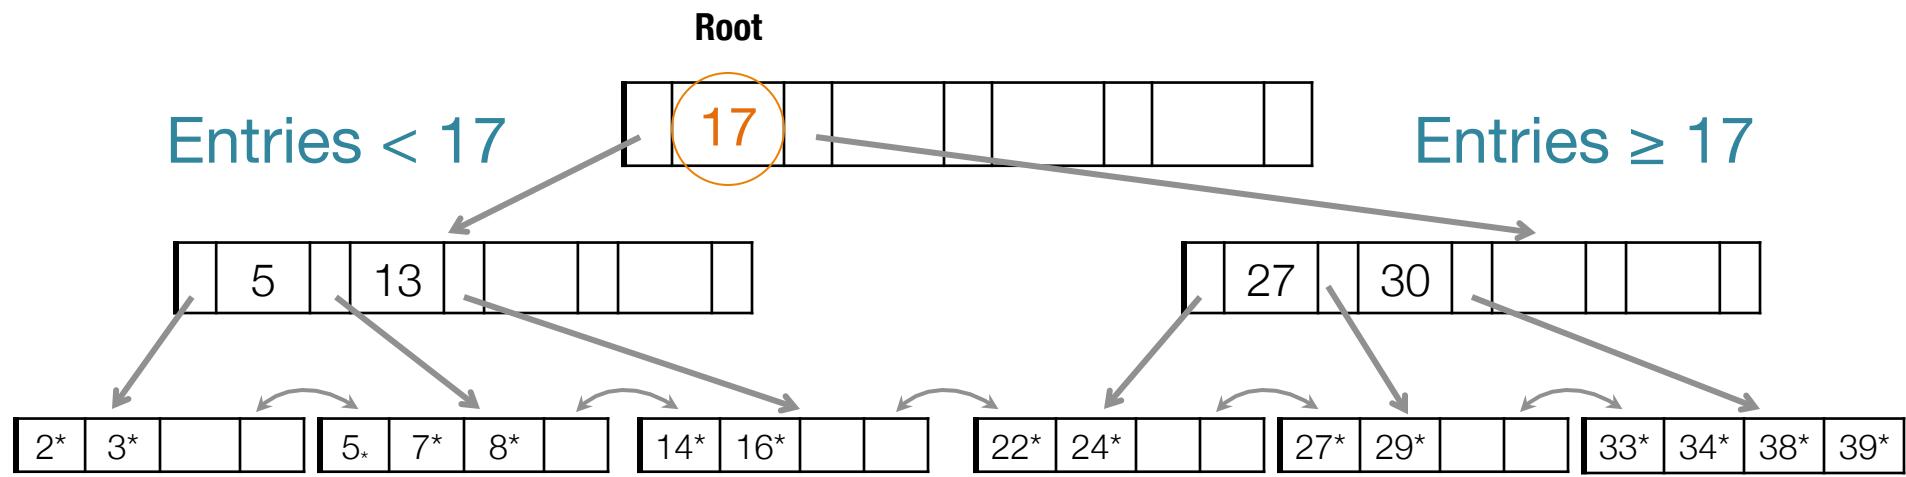
B+ Tree Indexes



- Leaf pages contain **data entries**, and are chained (prev & next)
- Non-leaf pages contain **index entries** and direct searches:



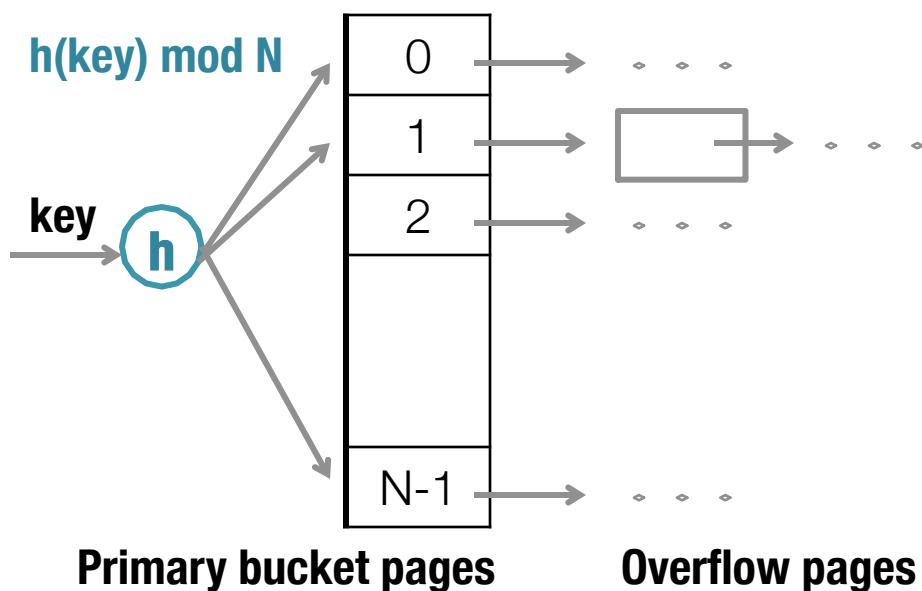
Example B+ Tree



- Find **29***? **28***? All $> 15^*$ and $< 30^*$
- Insert/delete: Find data entry in leaf, then change it
Need to adjust parent sometimes
 - And change sometimes bubbles up the tree

Hash-Based Indexes

- Good for equality selections
 - Index is a collection of **buckets**. Bucket = **primary page** plus zero or more **overflow pages**
 - Hashing function **h**: $h(r)$ = bucket in which record r belongs
h looks at the **search key** fields of r



Buckets contain:

- If Alternative (1) is used → data rec
- Alternative 2 → $\langle \text{key}, \text{rid} \rangle$
- Alternative 3 → $\langle \text{key}, \text{rid-list} \rangle$ pairs

Comparing File Organizations & Indexes Example

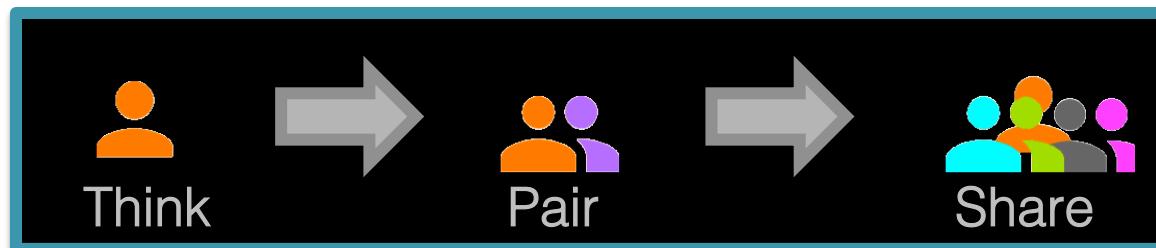


- Employees (Name, Age, Salary)
- 10 records/page and 10,000,000 records
- Assume that people are uniformly distributed between ages 1 to 100
- Operations:
 - Fetch the names of employees:
 - Scan: all employees
 - Equality Search: age = 21
 - Range Selection: age \geq 18 AND age $<$ 65
 - Insert a record
 - Delete a record

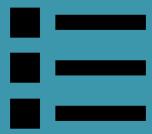
Analysis of I/O Cost

- For simplicity, ignore CPU cost
- Important Factors for I/O Cost:
 - How many pages (approx) read/written?
 - Are I/Os sequential or non-sequential?
- With your neighbor, do this for:
 - HEAP and Clustered index on Age
 - All 5 ops: SCAN, EQUALITY, RANGE, INS, DEL

Assumptions



Heap File



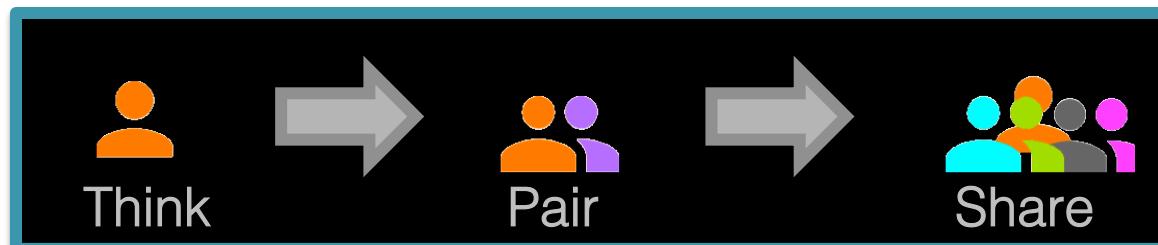
- Scan:
 - Read all pages in files sequentially
- Equality Search (age = 21):
 - Read all pages in files sequentially
 - worst case
- Range Selection on age:
 - Read all pages in file sequentially
- Insert:
 - Read and write last page or write into a new page (2 page I/Os)
- Delete:
 - Searching cost (expensive) + rewrite the page with the record

Clustered Index (on Age)

- Scan:
 - Approx. $1.5 \times$ heap file due to 2/3rd page occupancy (sequential)
- Equality Search (age = 18):
 - Traverse height of tree, read records page (non-sequential)
 - If multiple tuples qualify, read them sequentially
- Range Selection:
 - Traverse height of tree (non-sequential)
 - Scan leaf records (sequential)
- Insert:
 - Traverse height of tree (non-sequential) + write
- Delete:
 - Traverse height of tree (non-sequential) + write

Analysis of I/O Cost

- With your neighbor, do this for:
 - Heap-with-Unclustered-**Tree-Index** (on Age)
 - Heap-with-Unclustered-**Hash-Index** (on Age)
 - All 5 ops: SCAN, EQUALITY, RANGE, INS, DEL



Heap File w/ Unclustered Tree Index (on Age)

- Scan:
- Equality Search:
- Range Selection:

**Midterm
Review
on Thursday**

- Insert:
- Delete:

Heap File w/ Unclustered Hash Index (on Age)

- Scan:
- Equality Search:
- Range Selection:
- Insert:
- Delete:

**Midterm
Review
on Thursday**

Cost estimation

	scan	eq	range	ins	del
Heap	B	B/2	B	2	Search+1
sorted	B	$\log_2 B$	<- +m	Search+B	Search+B
Clust.	1.5B	h	<- +m	Search+1	Search+1
u-tree	$\sim B$	$1+h'$	$<- +m'$	Search+2	Search+2
u-hash	$\sim B$	~ 2	B	Search+2	Search+2

Cost estimation - big-O notation:

	scan	eq	range	ins	del
Heap	B	B	B	2	B
sorted	B	$\log_2 B$	$\log_2 B$	B	B
Clust.	B	$\log_F B$	$\log_F B$	$\log_F B$	$\log_F B$
u-tree	B	$\log_F B$	$\log_F B$	$\log_F B$	$\log_F B$
u-hash	B	1	B	1	1