


Tree-Structured Indexes

Chapter 10

Index Design Space

Organization Structure for k^*

- Hash-based
 - (+) Equality search
- Tree-based
 - (+) Range, equality search
 - B+Tree (dynamic)
 - ISAM (static)



Data Entry (k^*) Contents

1. Actual Data record
 $\text{index} = \text{file}$
2. $\langle k, \text{rid} \rangle$
actual records in a diff file
3. $\langle k, \text{list of rids} \rangle$


Motivation

- Range and equality searches very common
- Can scan heap file, but this is expensive
- **Goal:** Create a dynamic index structure that allows for efficient evaluation, and equality / range queries.




B+ Tree

- Height-balanced (dynamic) tree search structure
 - Supports equality and range-searches efficiently
- Main Characteristics:
 1. Minimum 50% occupancy (except for root), i.e. each node contains $d \leq m \leq 2d$ entries, where d = order of the tree.
 2. Height of tree: the length of any path from the root to a leaf.
 3. Remains balanced after insert/delete operation.




**Index entries
(direct search)**

Index Entries:
Entries in the non-leaf pages:
(search key value, pageid)




Searching in a B+ Tree

- Starting from root, examine index entries in non-leaf nodes, and traverse down the tree until a leaf node is reached
 - Non-leaf nodes can be searched using a binary or a linear search.
- Try it out: Search for 5*, 15*, all data entries $\geq 24^*$




B+-Tree Page Format

Non-leaf Page



Leaf Page



B+ Trees: Height

What is the height of a B+ tree? (formula)



Definitions:

F : fanout (average number of children for non-leaf node)

N : total leaf pages


$$\log_F N$$

B+ Trees in Practice

- Typical order 100. Typical fill-factor 66.5%
 - Maximum fanout: 201
 - Average fanout = 133
- Typical capacity:
 - Height = 1: 133 pages of data entries (leaf pages)
 - Height = 2: 133^2 pages of data entries
 - Height = 3: 133^3 (> 2 million) pages of data entries
 - Height = 4: 133^4 (> 300 million) pages of data entries
- Can often keep top levels of index in buffer pool
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 MBytes


Arithmetic Example

- You are given a file of 10 million records
- Suppose you can store 10 data entries per leaf page
- You build a B+ Tree with order 75, 67% average fill-factor
- What is the avg fanout?
- What is the height of your B+ Tree?



A Note on Order

- Some literature uses: $\text{order} = \text{max } \# \text{ of entries}$
- In this class: $\text{order} = \text{minimum } \# \text{ of entries}$
(book uses this)
- *Order d* concept replaced by physical space criterion in practice (e.g. at least half-full)
 - Index (i.e. non-leaf) pages can typically hold many more entries than leaf pages



B+ Tree Operations

- Search
 - Equality
 - Range
- Insert data entry
- Delete data entry
- Bulk load


B+ Tree: Inserting a Data Entry




- Maintain invariants:
 - Search-tree property
 - All nodes must be at least $\frac{1}{2}$ full (except root node), i.e., has between d and $2d$ entries
 - Root node is allowed to have a single entry
- Strategy:
 - Split nodes when they become full and a node is added:
 - Split an overfull node with $(2d + 1)$ entries into two nodes with d and $(d+1)$ entries, resp.




Inserting 8* into B+ Tree



Inserting 8* into B+ Tree




Overflow




Splitting Overfull Non-Leaf Nodes

Overfull




Splitting Overfull Non-Leaf Nodes

Split




Splitting Overfull Non-Leaf Nodes




Splitting Overfull Non-Leaf Nodes

Minimum occupancy is guaranteed in both leaf and index page splits (but not in the root)




The B⁺-tree After Inserting 8*




- Root was split: height increases by 1
- Could avoid split by **re-distributing entries** with a **sibling**
 - Sibling: immediately to left or right, **and same parent**

Inserting 8* via Entry Re-distribution with Siblings



Inserting 8* via Entry Re-distribution with Siblings



- Re-distributing entries with a *sibling*
 - Improves page occupancy, possibly reduces height
 - Usually not used for non-leaf node splits. Why?
 - Increases I/O, especially if we check both siblings
 - Better if split propagates up the tree (rare)
 - Use only for leaf level entries
 - have to set pointers

B+ Tree Operations


- Search
 - Equality
 - Range
- Insert data entry
- Delete data entry
- Bulk load

B+-Tree: Deleting a Data Entry

1. Find the data entry (will always be at a )
2. Delete it
3. Restore the B+ tree invariant
 - If L is at least half-full, *done!*
 - If L has only $d-1$ entries,
 - Try to [re-distribute](#), borrowing from a *sibling* (adjacent node with *same parent* as L)
 - If re-distribution fails, [merge L](#) and sibling
4. On merge, delete relevant entry in parent


Merge could propagate to root, decreasing height.

Example Tree




- **Task:** Delete 22, 20, 24
- Deleting 22 is easy. Invariant maintained.
- Deleting 20 is harder. Node would become less than half full.

Deleting 22* and 20*




Deleting 22* and 20*




Deleting 20* is done with re-distribution.
Notice how middle key is **copied up**.

... And then Deleting 24*





- Must merge
- In the non-leaf node,
toss the index entry with key value = 27

... And then Deleting 24*




- Must merge
- In the non-leaf node,
toss the index entry with key value = 27

... And then Deleting 24*




Another Deletion Example: Non-leaf Re-distribution

Can re-distribute entry from left child of root to right child.




After Re-distribution

- Rotate through the parent node
- re-distribute index entry with key 20



After Re-distribution

It suffices to re-distribute index entry with key 20;
For illustration 17 is also re-distributed



B+ Tree Deletion

- Try redistribution with **all** siblings first, then merge. Why?
 - Good chance that redistribution is possible (large fanout!)
 - Only need to propagate changes to parent node
 - Files typically grow, not shrink!

B+ Tree Operations


- Search
 - Equality
 - Range
- Insert data entry
- Delete data entry
- Bulk load

Adding Entries to a B+ tree

- Option 1: multiple inserts
 - Slow. Repeated re-organization
 - Does not give sequential storage of leaves
- Option 2: Bulk Loading
 - Fewer I/Os during build
 - Leaves will be stored sequentially (and linked, of course)
 - Can control “fill factor” on pages


Bulk Loading of a B+ Tree

Initialization: Sort all data entries, insert pointer to first (leaf) page in a new (root) page.




Bulk Loading (Contd.)

New index entries for leaf pages always entered into right-most



index page just above leaf level.
When this fills up, it splits. (Split may go up right-most path to the root.)

Much faster than repeated inserts!



Summary

- Tree-structured indexes are ideal for range-searches, also good for equality searches.
- B+ tree is a dynamic height-balanced index structure.
 - Insertions/deletions/search costs $O(\log_F N)$.
 - High fanout (F) means depth rarely more than 3 or 4.
 - Max occupancy = $2d$
 - Min occupancy = order = d (at least half full – exc. root)
 - Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.

Suggested Exercises + Readings

Suggested Exercises: 10.1, 10.5, 10.7

Suggested readings for the lecture on hash-based index:

Read the entire chapter 11 (hash index)