

Hash-Based Indexes

Chapter 11

Index Design Space

Organization Structure for k^*

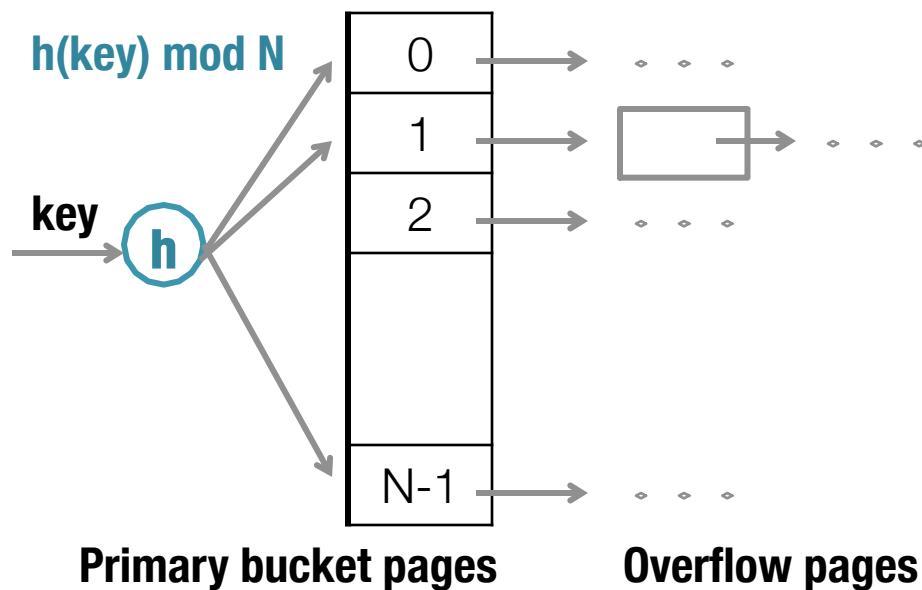
- Tree-based
 - (+) Range, equality search
- **Hash-based**
 - (+) Equality search
 - Static hashing
 - Extendible hashing
 - Linear hashing



1. Actual Data record
 $\text{index} = \text{file}$
2. $\langle k, \text{rid} \rangle$
actual records in a diff file
3. $\langle k, \text{list of rids} \rangle$

Static Hashing

- # primary bucket pages fixed, allocated sequentially, never de-allocated; overflow pages if needed.
- $h(\text{key}) \bmod N$ = bucket to which data entry with **key** belongs. ($N = \# \text{ of buckets}$)



What are
downsides of
static hashing?



Static Hashing (Contd.)

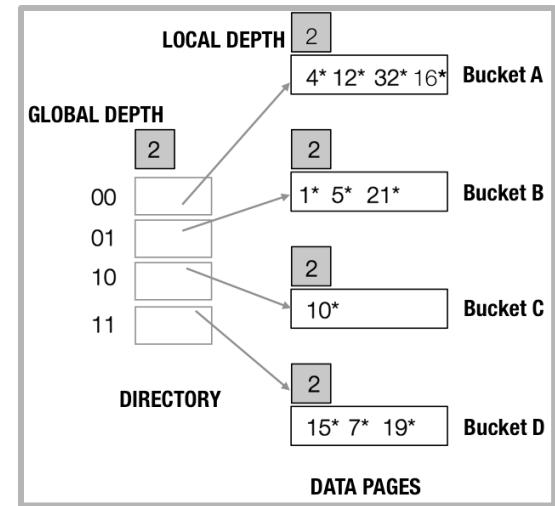
- Buckets contain **data entries**
- Number of buckets (N) is **fixed** ahead of time
- Static structure can be problematic
 - Consider many insertions
 - Long overflow chains can develop (and degrade performance!)
- Might consider periodically doubling N and “rehashing” file
 - Entire file has to be read and written
 - Index unavailable while rehashing
 - Dynamic hashing fixes this problem: **Extendible** and **Linear Hashing**



Why not start w/
a very large N ?

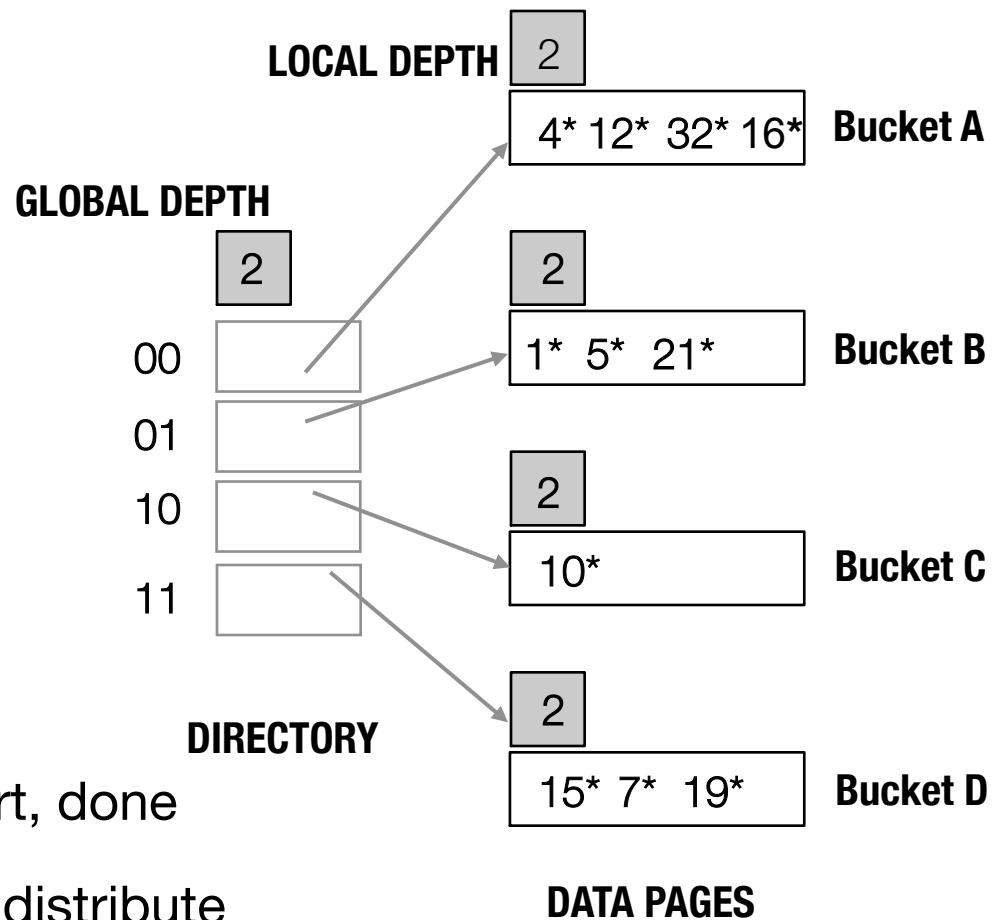
Extendible Hashing

- **Main Idea:** Use a directory of pointers to buckets
- On overflow, double the directory (not # number of buckets)
- Why does this help?
 - Directory much smaller than the file
 - Only one page of data entries is split at a time
 - (Usually*) No overflow pages



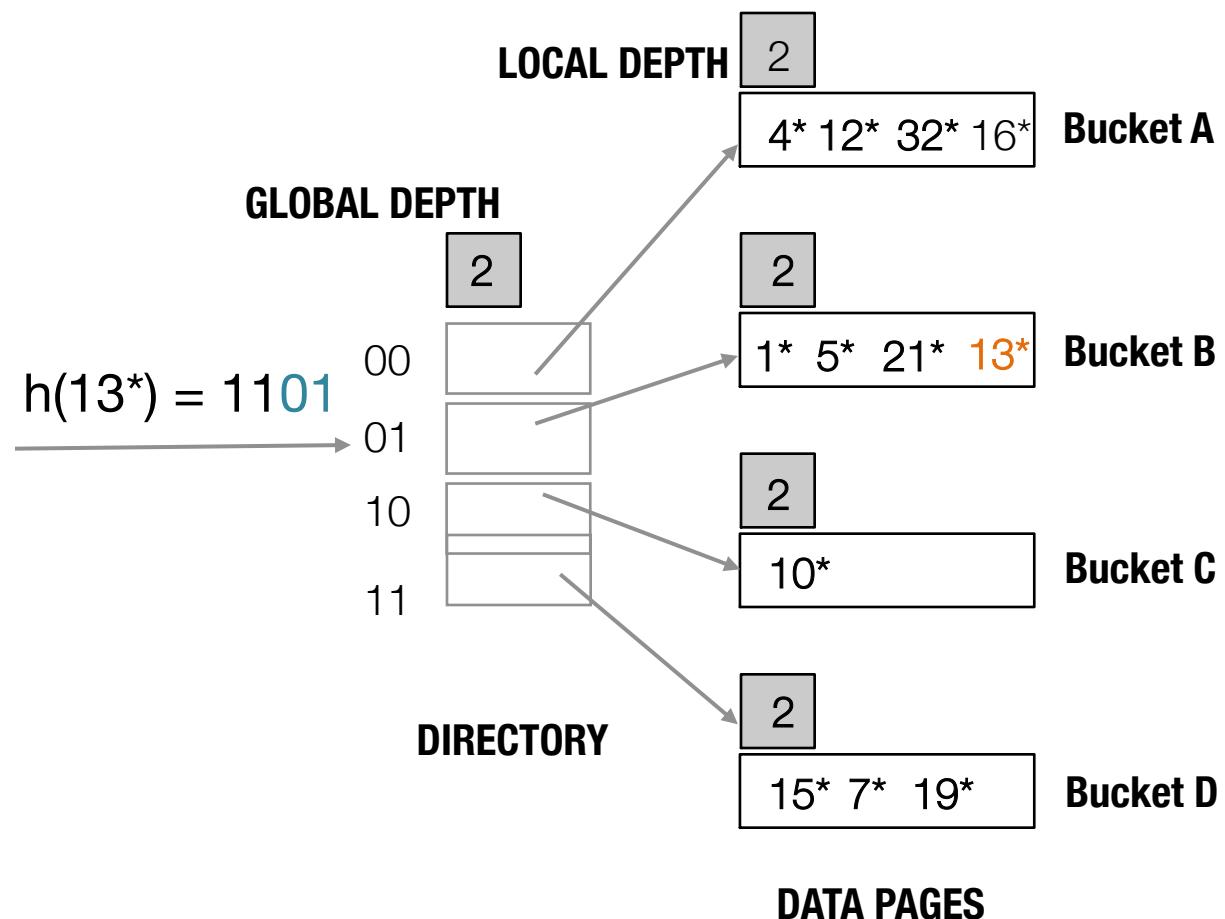
Extendible Hashing: Algorithm

- Directory = an array
- Search for k :
 - Apply hash function $h(k)$
 - Take last **global depth** # bits of $h(k)$
- Insert:
 - If bucket has space, insert, done
 - If bucket is full, **split** it, re-distribute
 - If necessary, double the directory



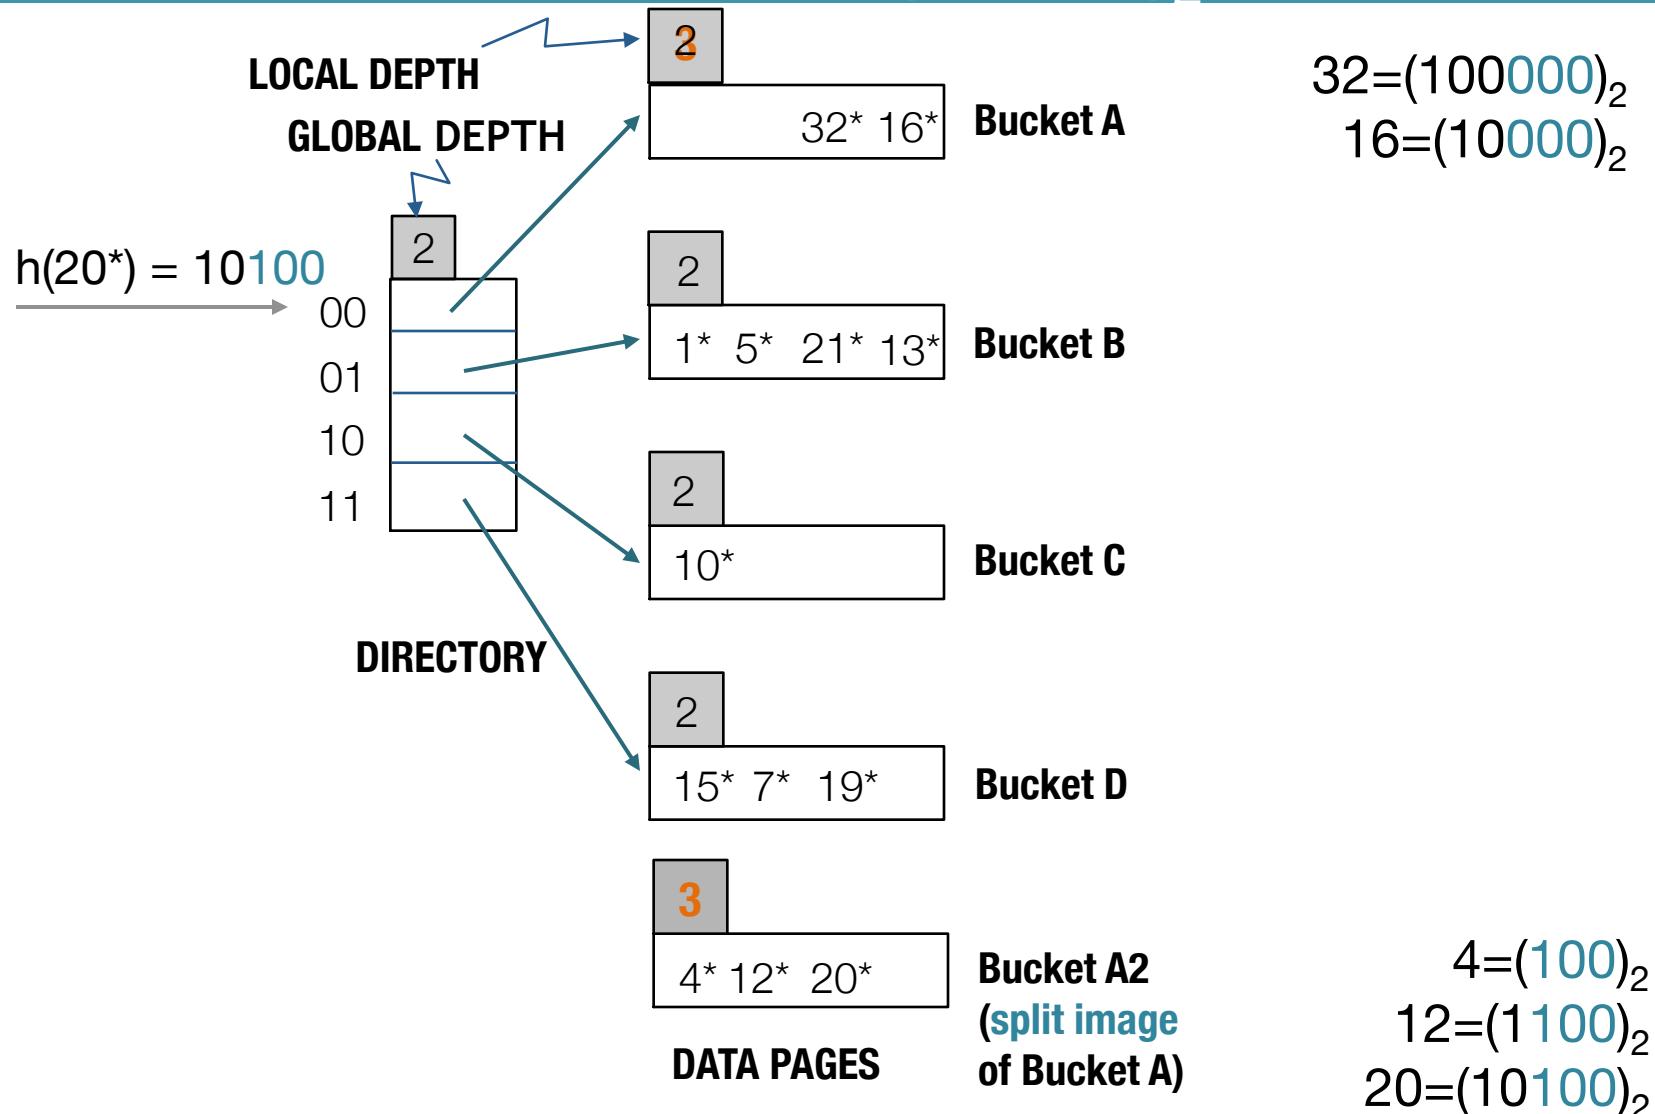
Example: Insertion into a free bucket

- Insert 13^*
- Suppose
 $h(13^*) = 1101$



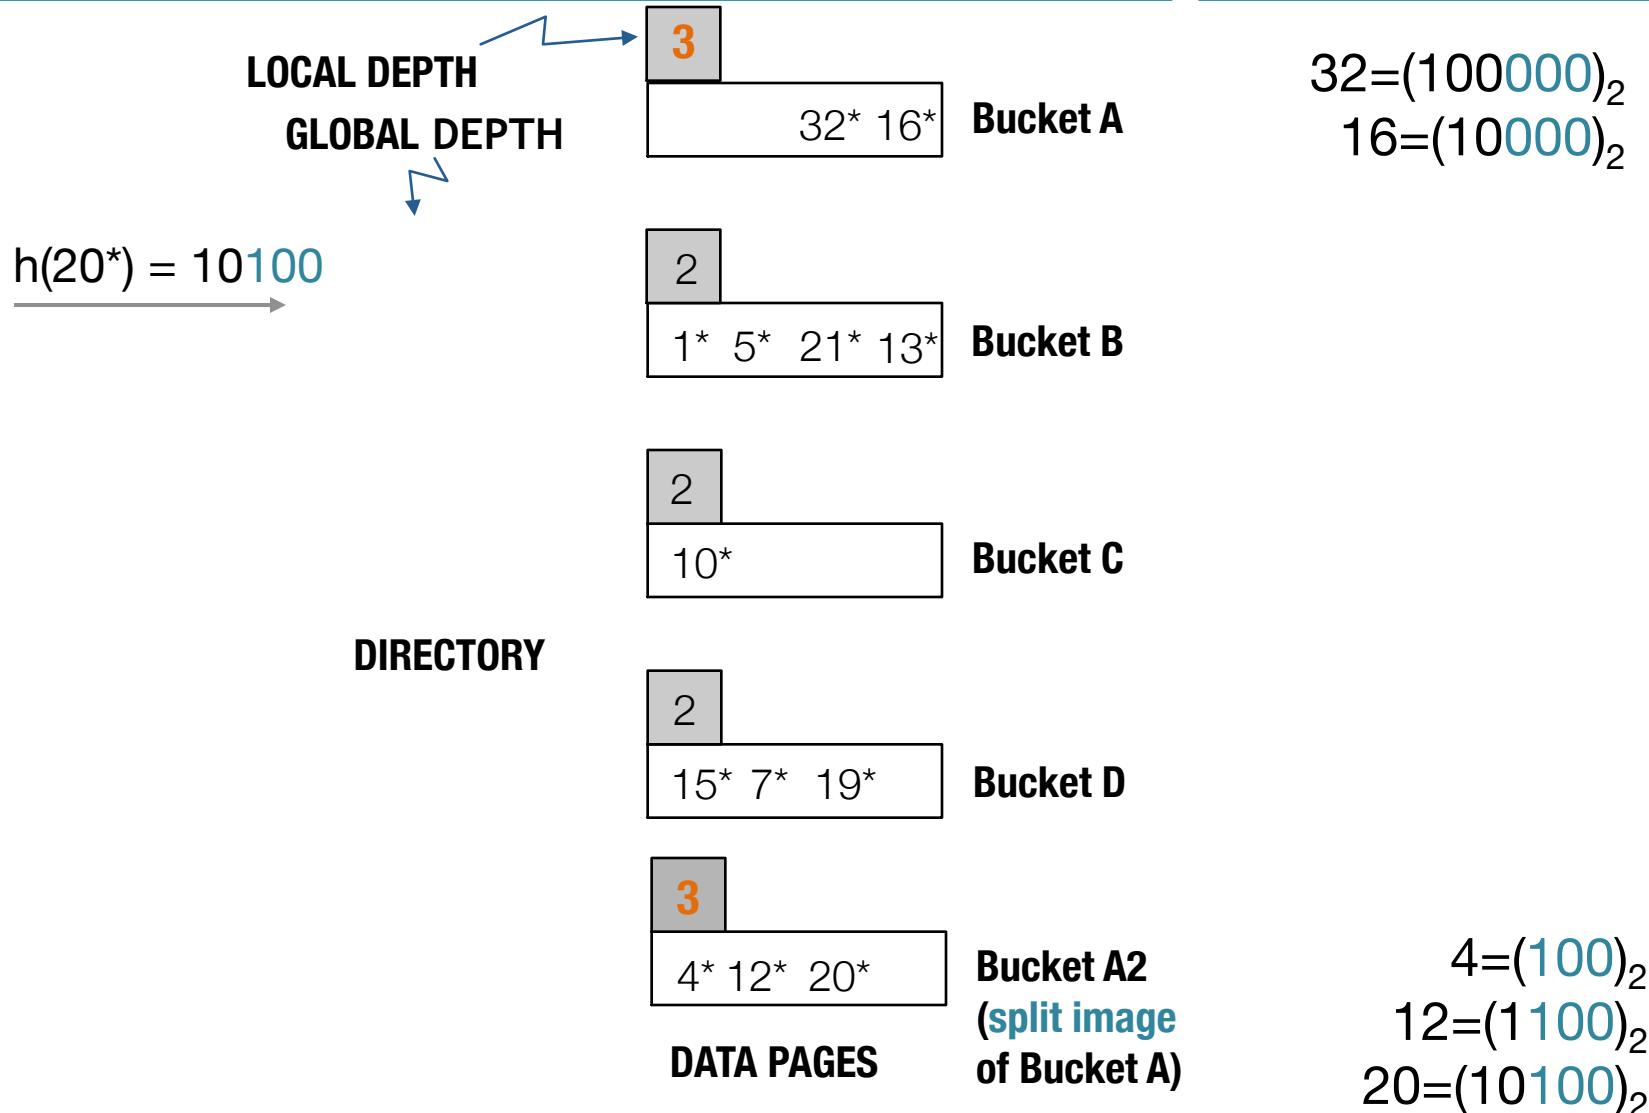
Example: Insertion into a full bucket

Insert $20 = (10100)_2$



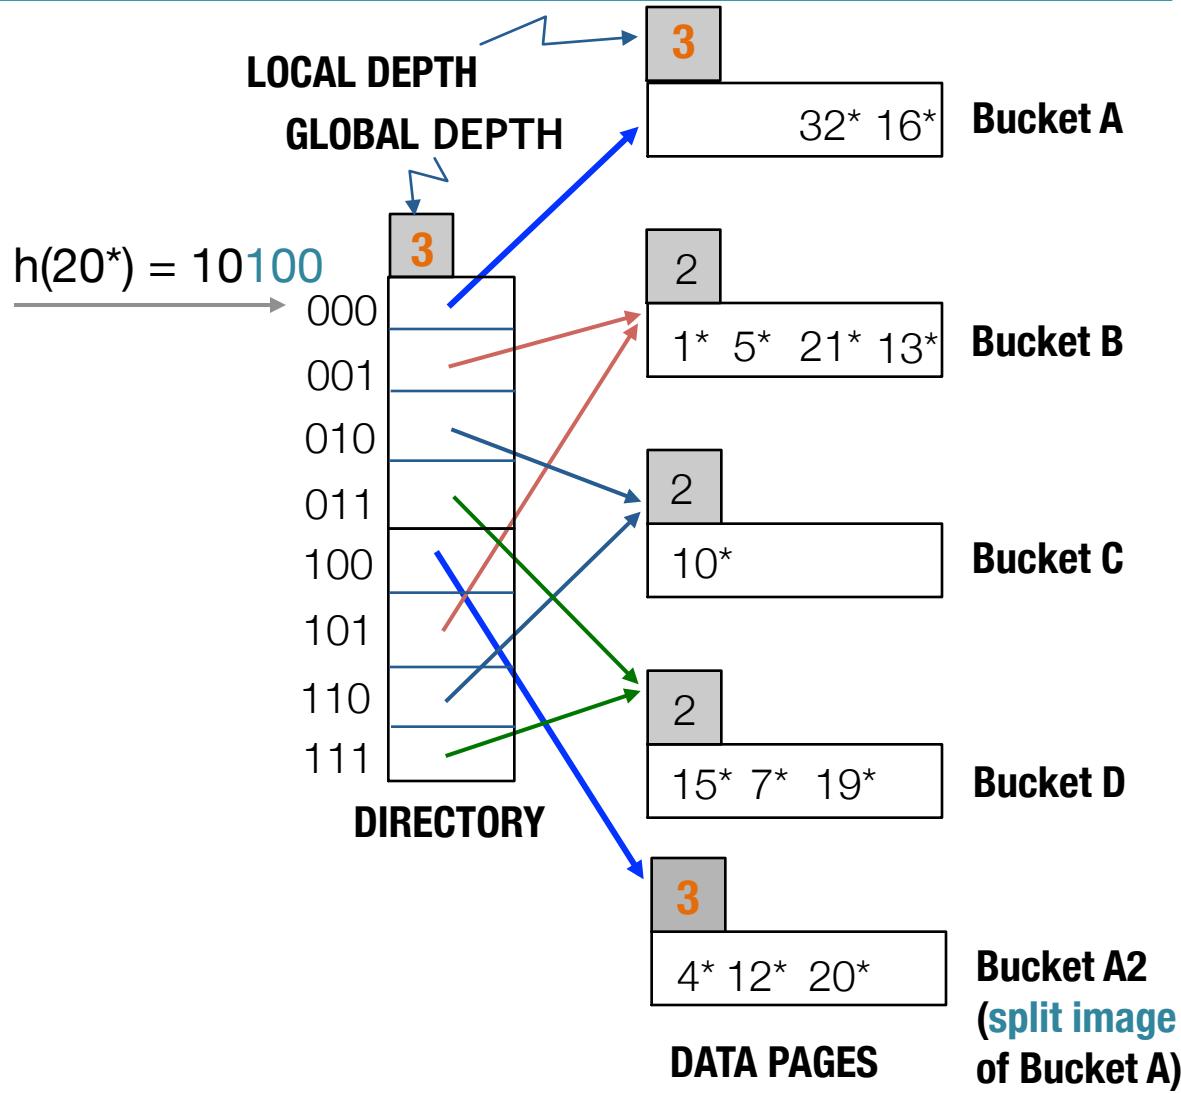
Example: Insertion into a full bucket

Insert $20 = (10100)_2$



Example: Insertion into a full bucket

Insert $20 = (10100)_2$



$$32 = (100000)_2$$
$$16 = (10000)_2$$

Notice that splitting a bucket only requires doubling the directory when $LD = GD$

Directory doubled by *copying it over* and fixing pointer to split image page

$$4 = (100)_2$$
$$12 = (1100)_2$$
$$20 = (10100)_2$$

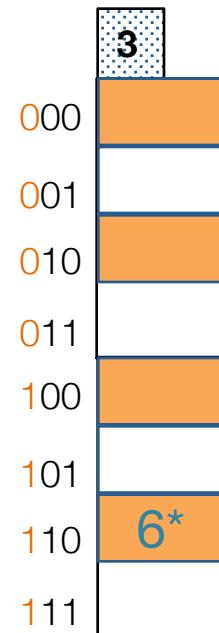
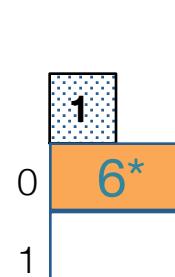
Directory Doubling



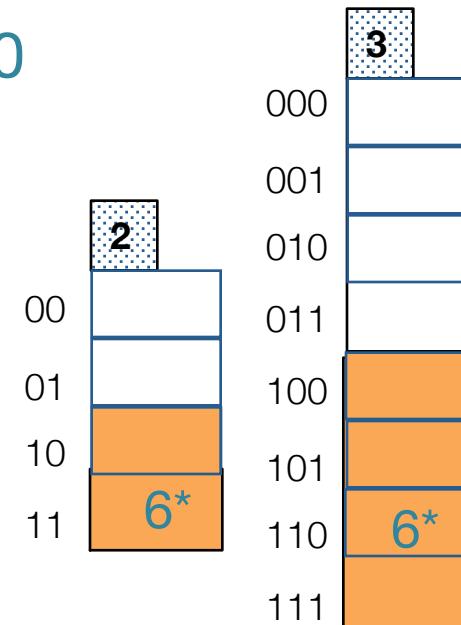
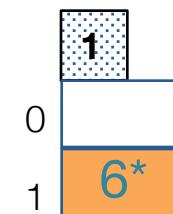
Why use least significant bits in directory?

Allows for doubling via copying!

$$6 = 110$$



$$6 = 110$$



Least Significant

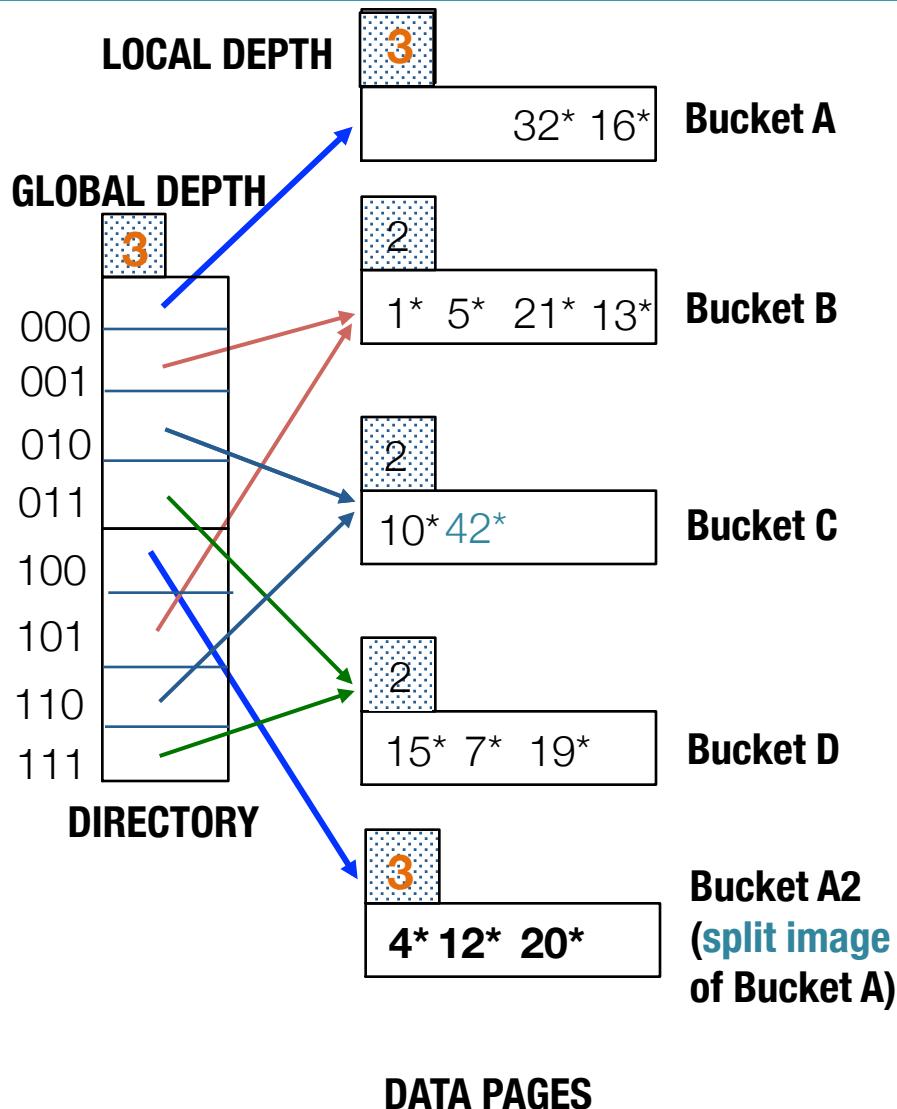
vs.

Most Significant

Comments on Extendible Hashing

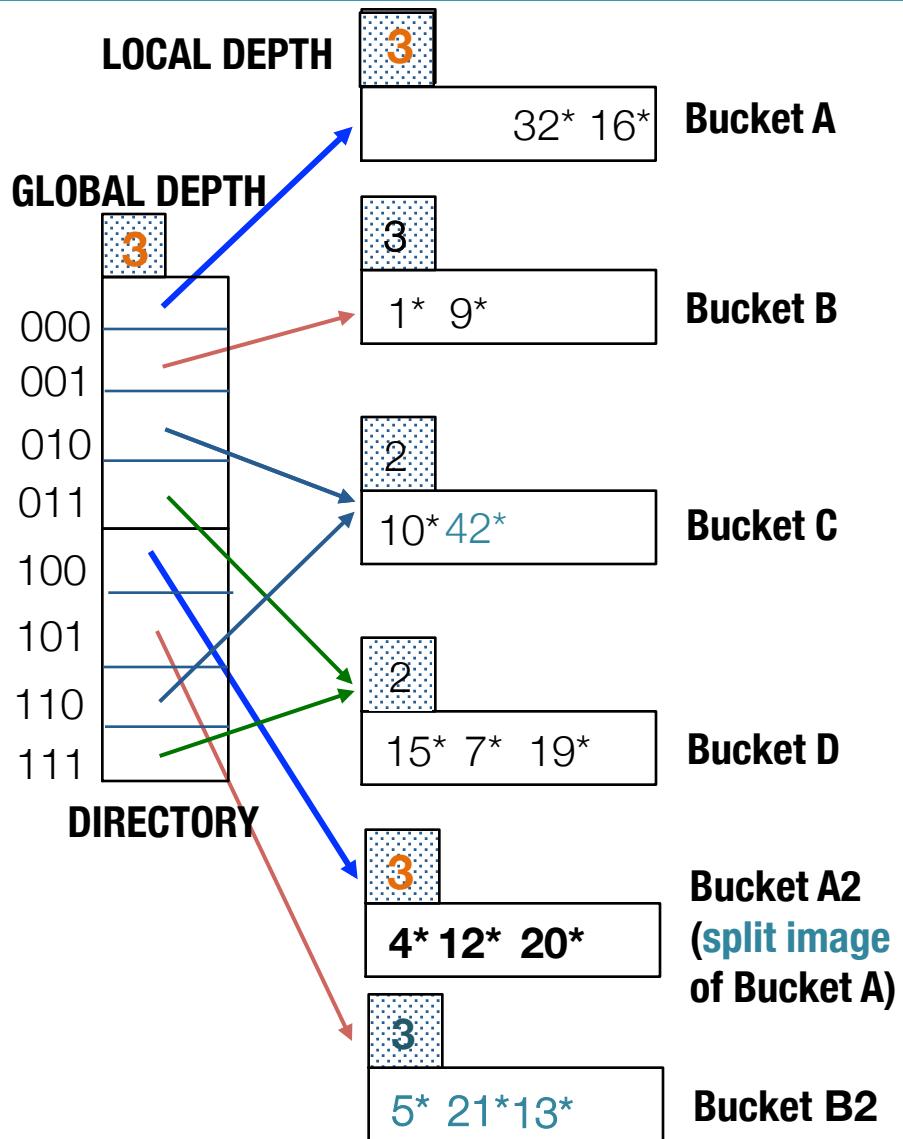
- How many disk accesses for equality search?
 - One if directory fits in memory, else two
- Directory grows in spurts, and, if the distribution of *hash values* is skewed, directory can grow large
- Do we ever need overflow pages?
 - Multiple entries with same hash value cause problems!
=> We need overflow pages
- **Delete:** Reverse of inserts – see textbook

Extendible Hashing Quiz: Insert 42, 9



$$42 = (101010)_2$$
$$9 = (1001)_2$$

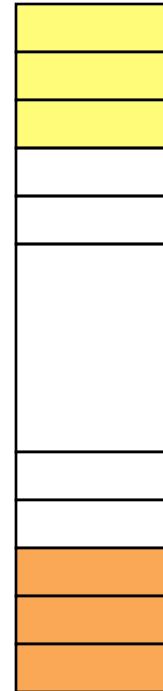
Answer: Insert 42, 9



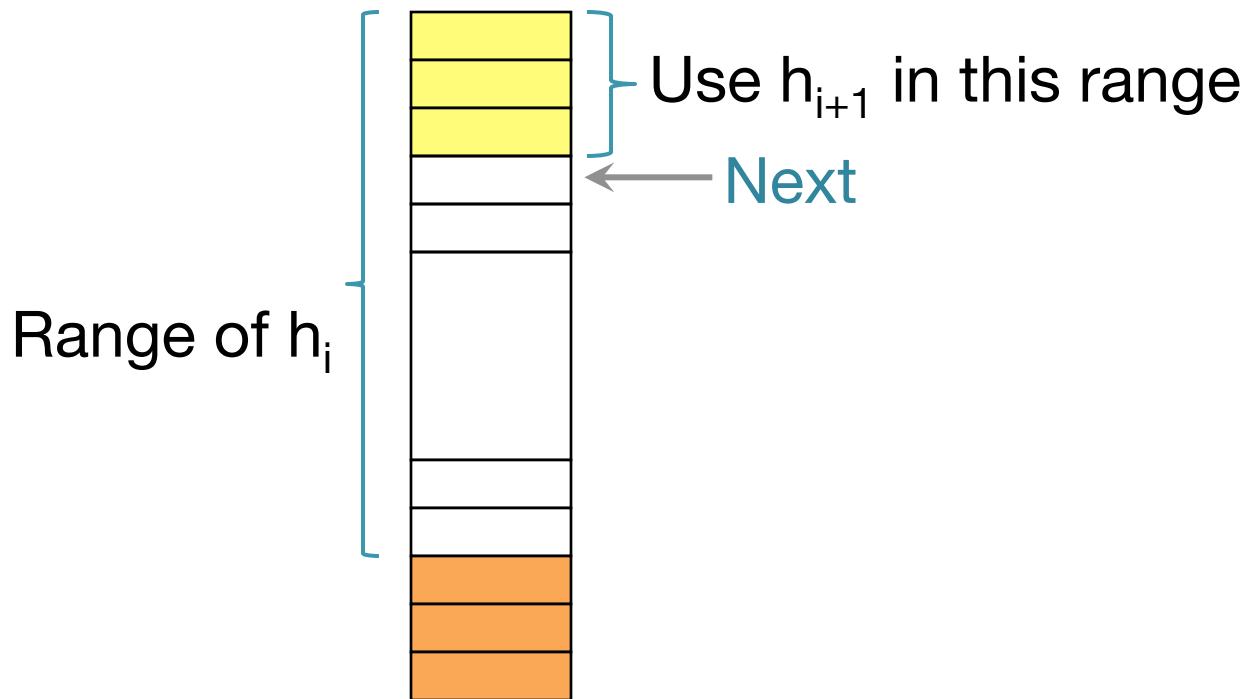
$$42 = (101010)_2$$
$$9 = (1001)_2$$

Linear Hashing

- Tries to overcome some of the drawbacks of Extendible Hashing
- Works in rounds
- Grows more gradually

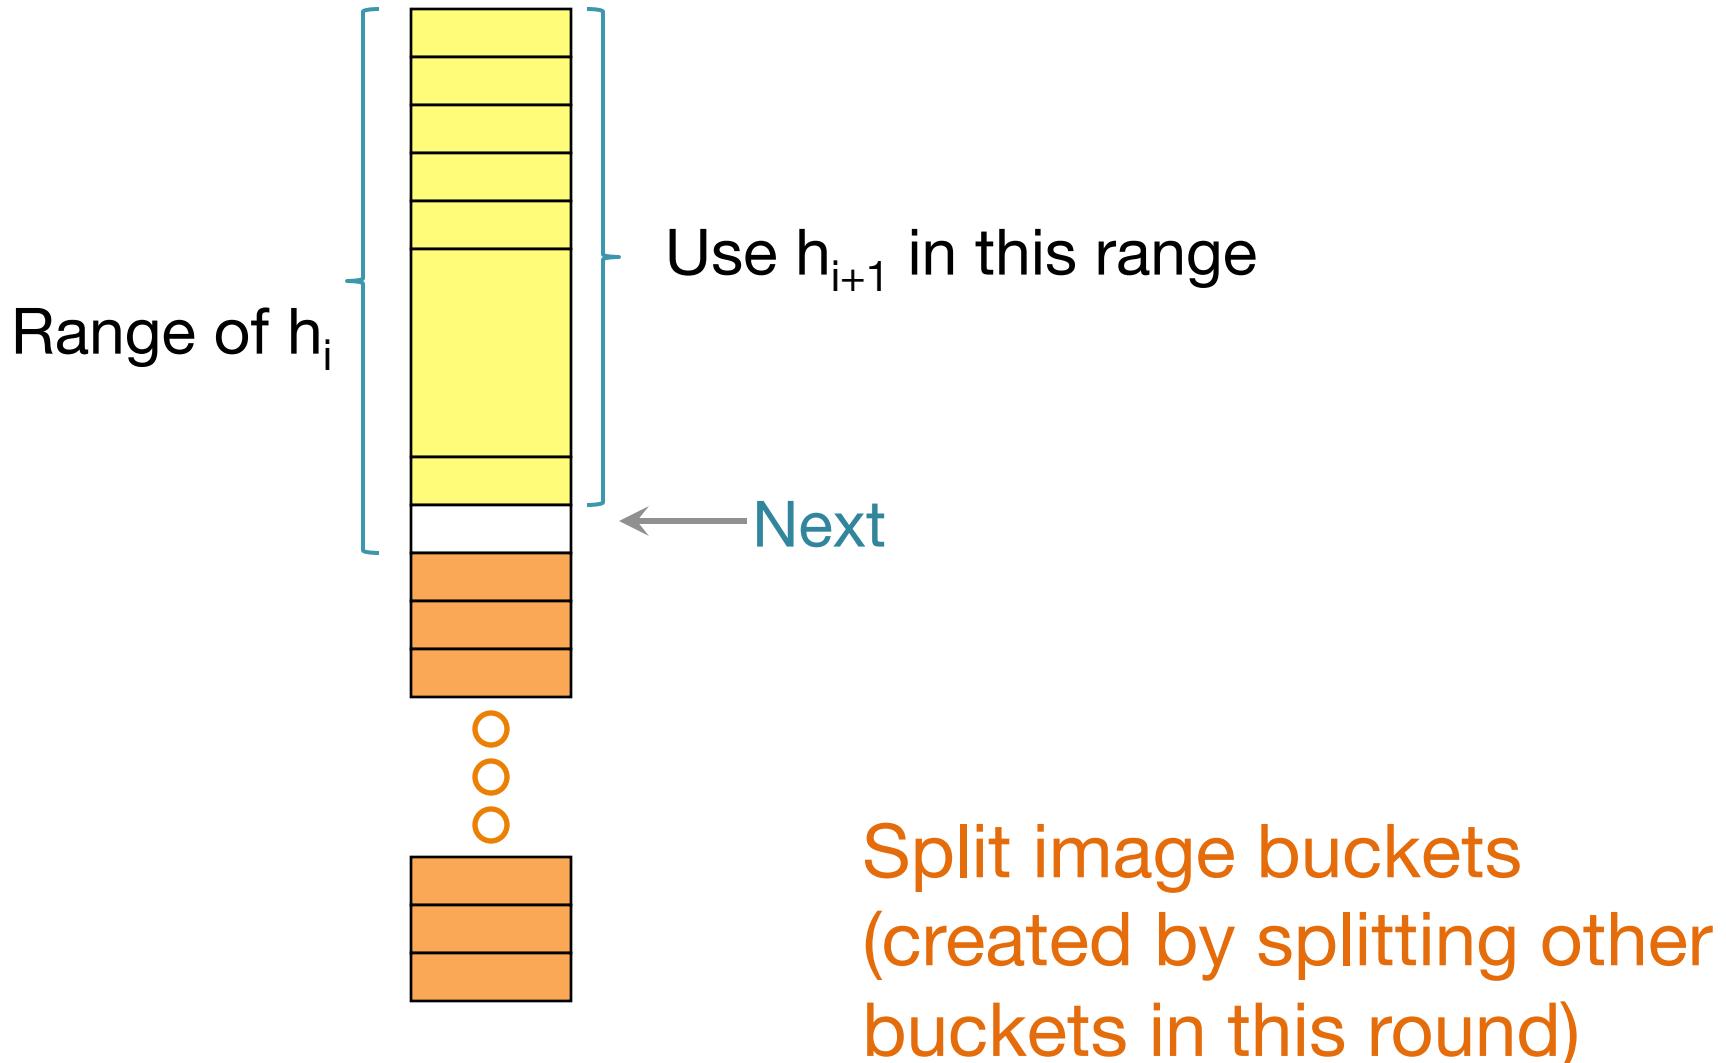


Buckets During Round i

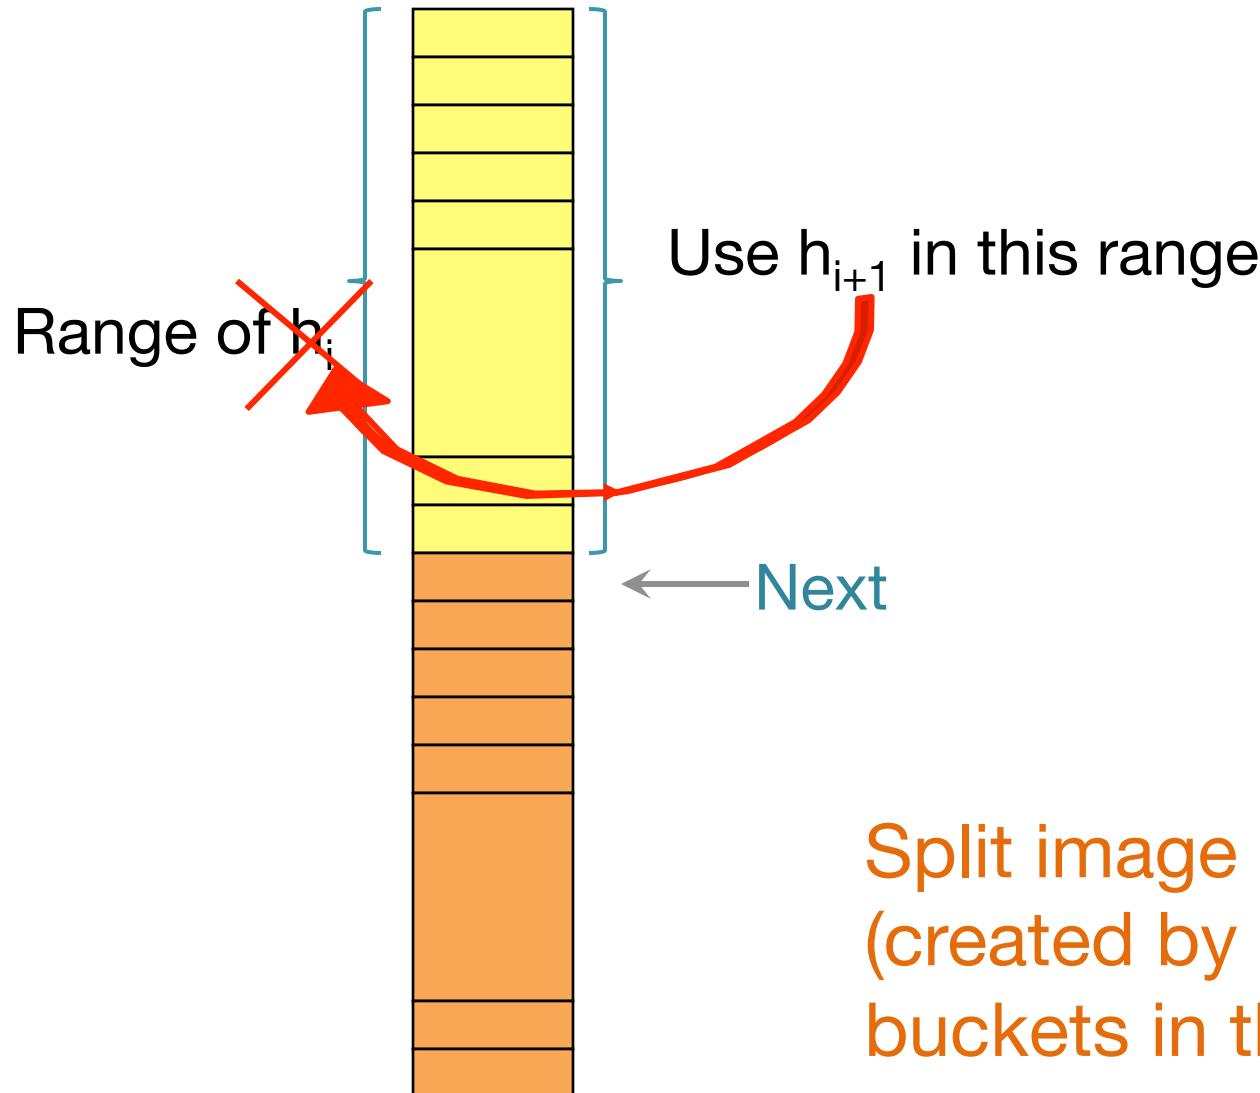


Split image buckets
(created by splitting other
buckets in this round)

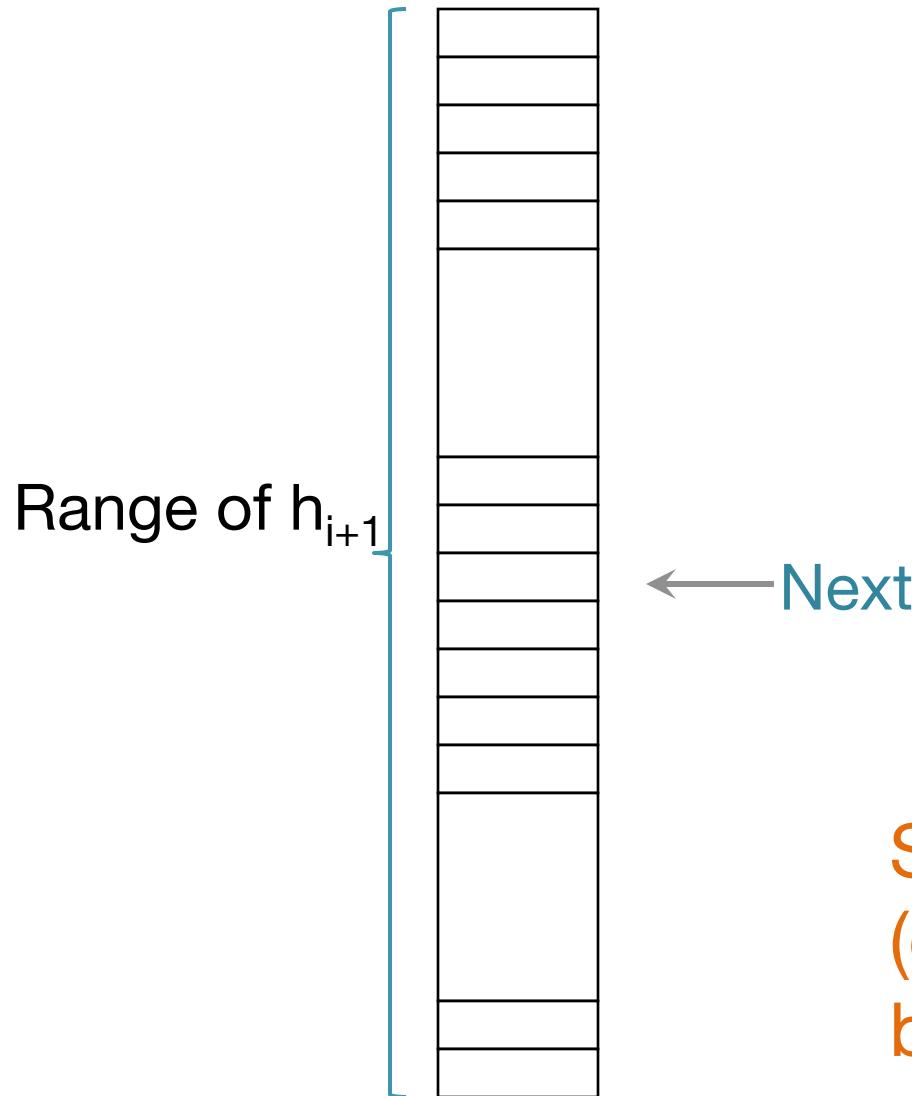
Buckets During Round i



End of Round i = Start of Round $i+1$

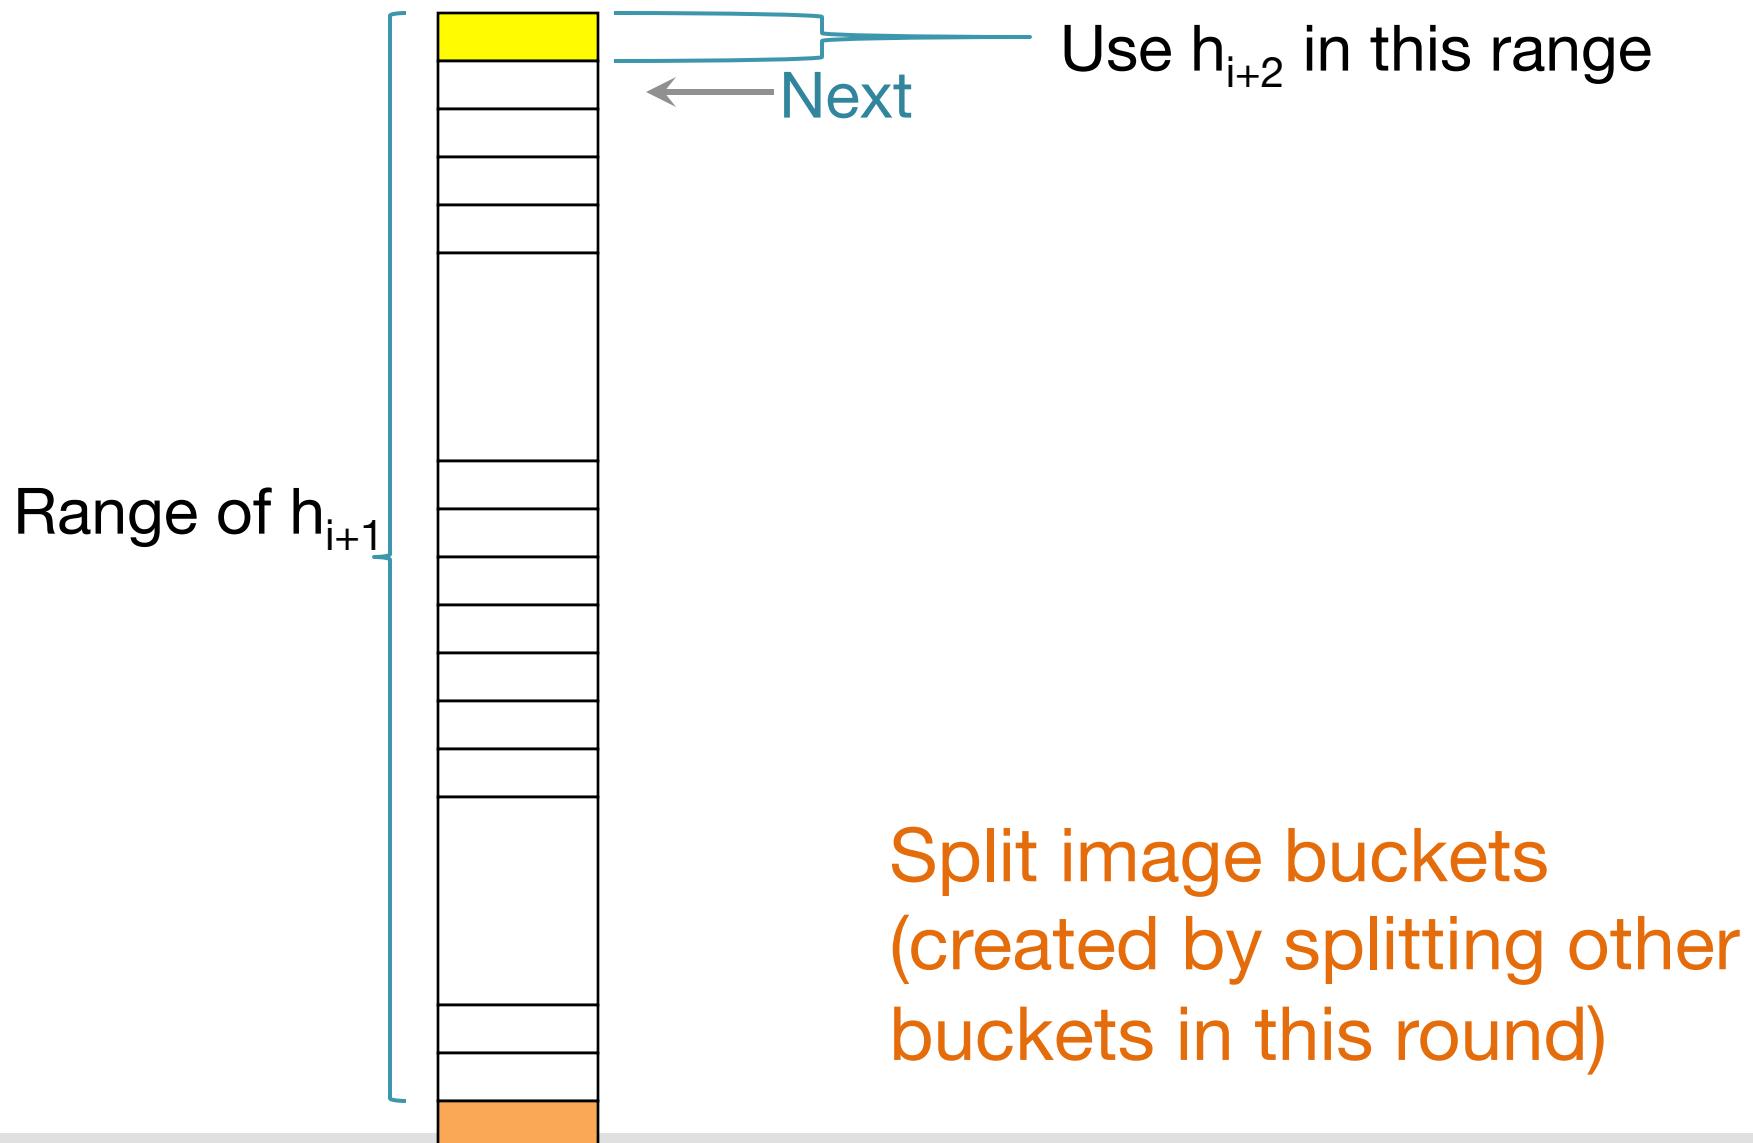


End of Round i = Start of Round $i+1$



Split image buckets
(created by splitting other
buckets in this round)

Round $i+1$ after the First Split



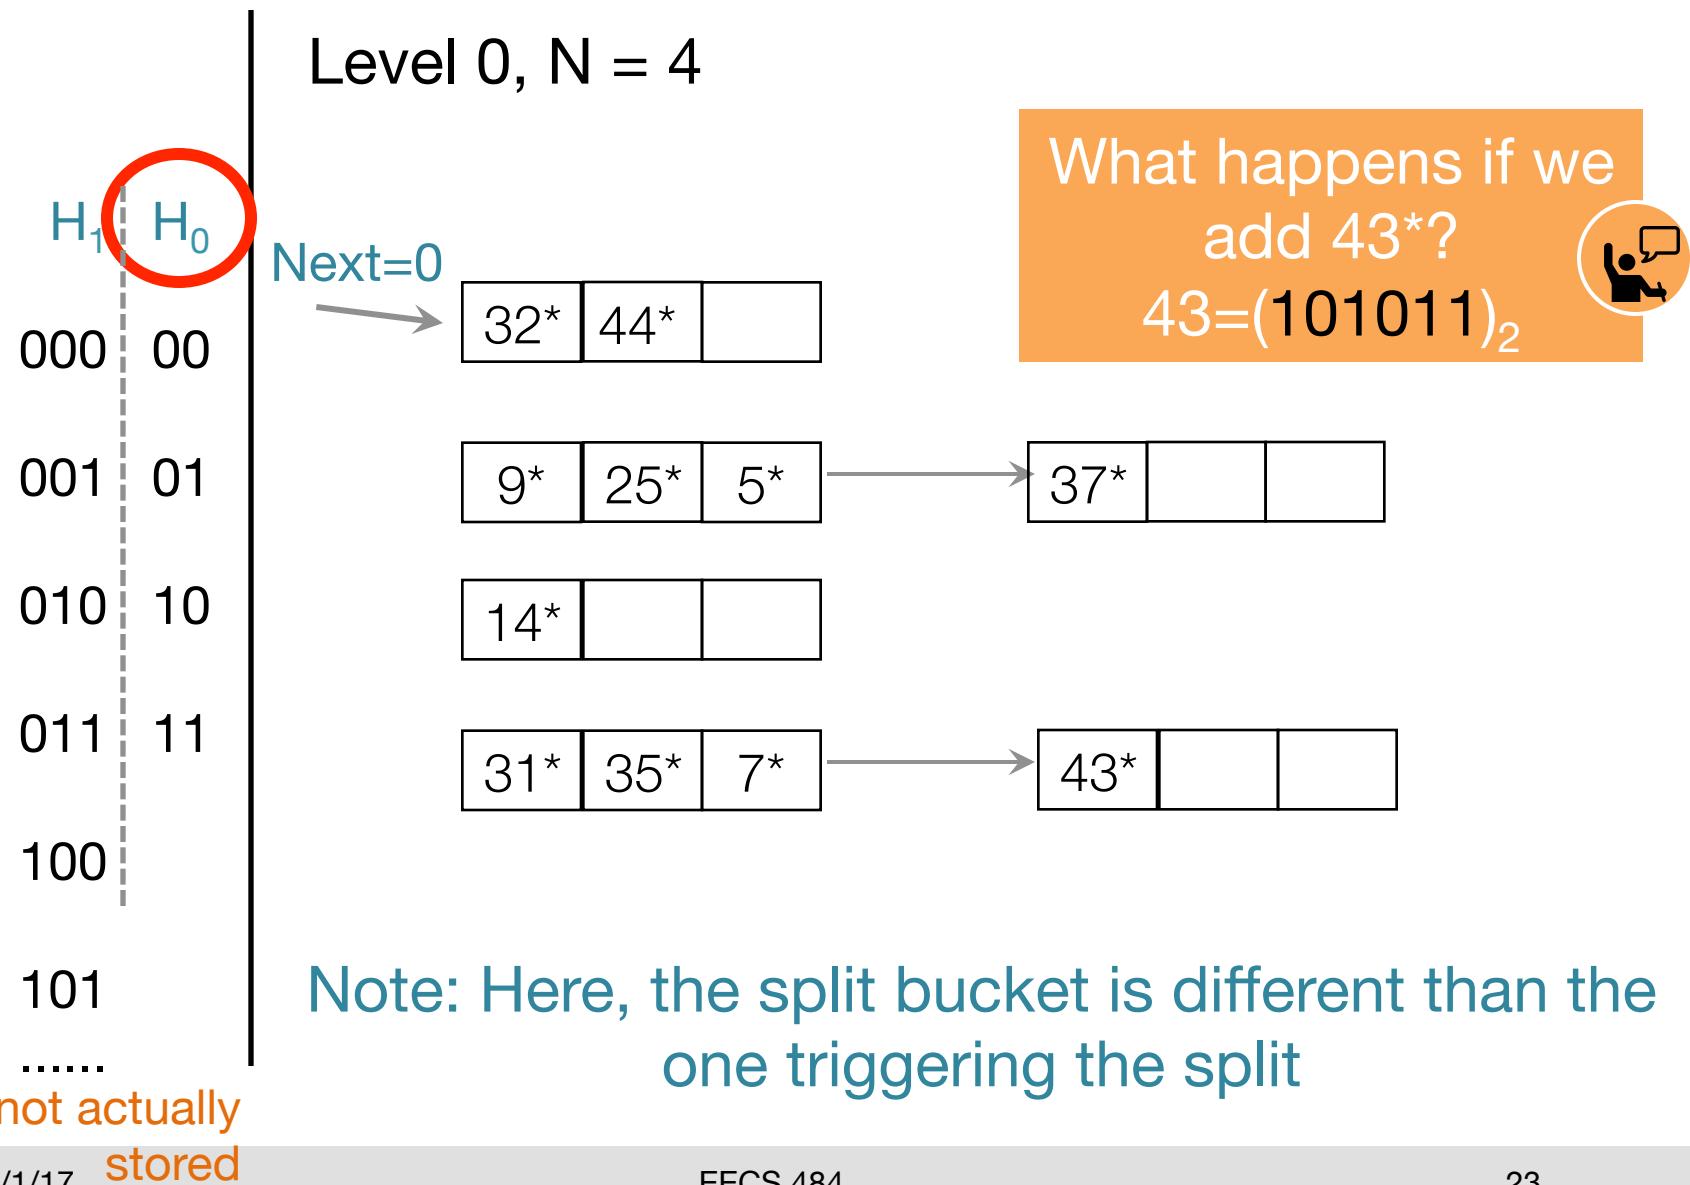
Linear Hashing: Main Idea

- Another dynamic hashing scheme
- Eliminates long overflow chains without using a directory
- **Idea:** Use a family of hash functions: h_0, h_1, h_2, \dots
 - h_{i+1} doubles the range of h_i (similar to directory doubling)
 - Hash family typically obtained by choosing hash function $h()$ and initial number of buckets N
 - Define $h_i(\text{value}) = h(\text{value}) \bmod (2^i N)$
 - If N is a power of 2, say $N=2^{d_0}$, apply hash function $h()$, and look at last d_i bits where $d_i = d_0 + i$

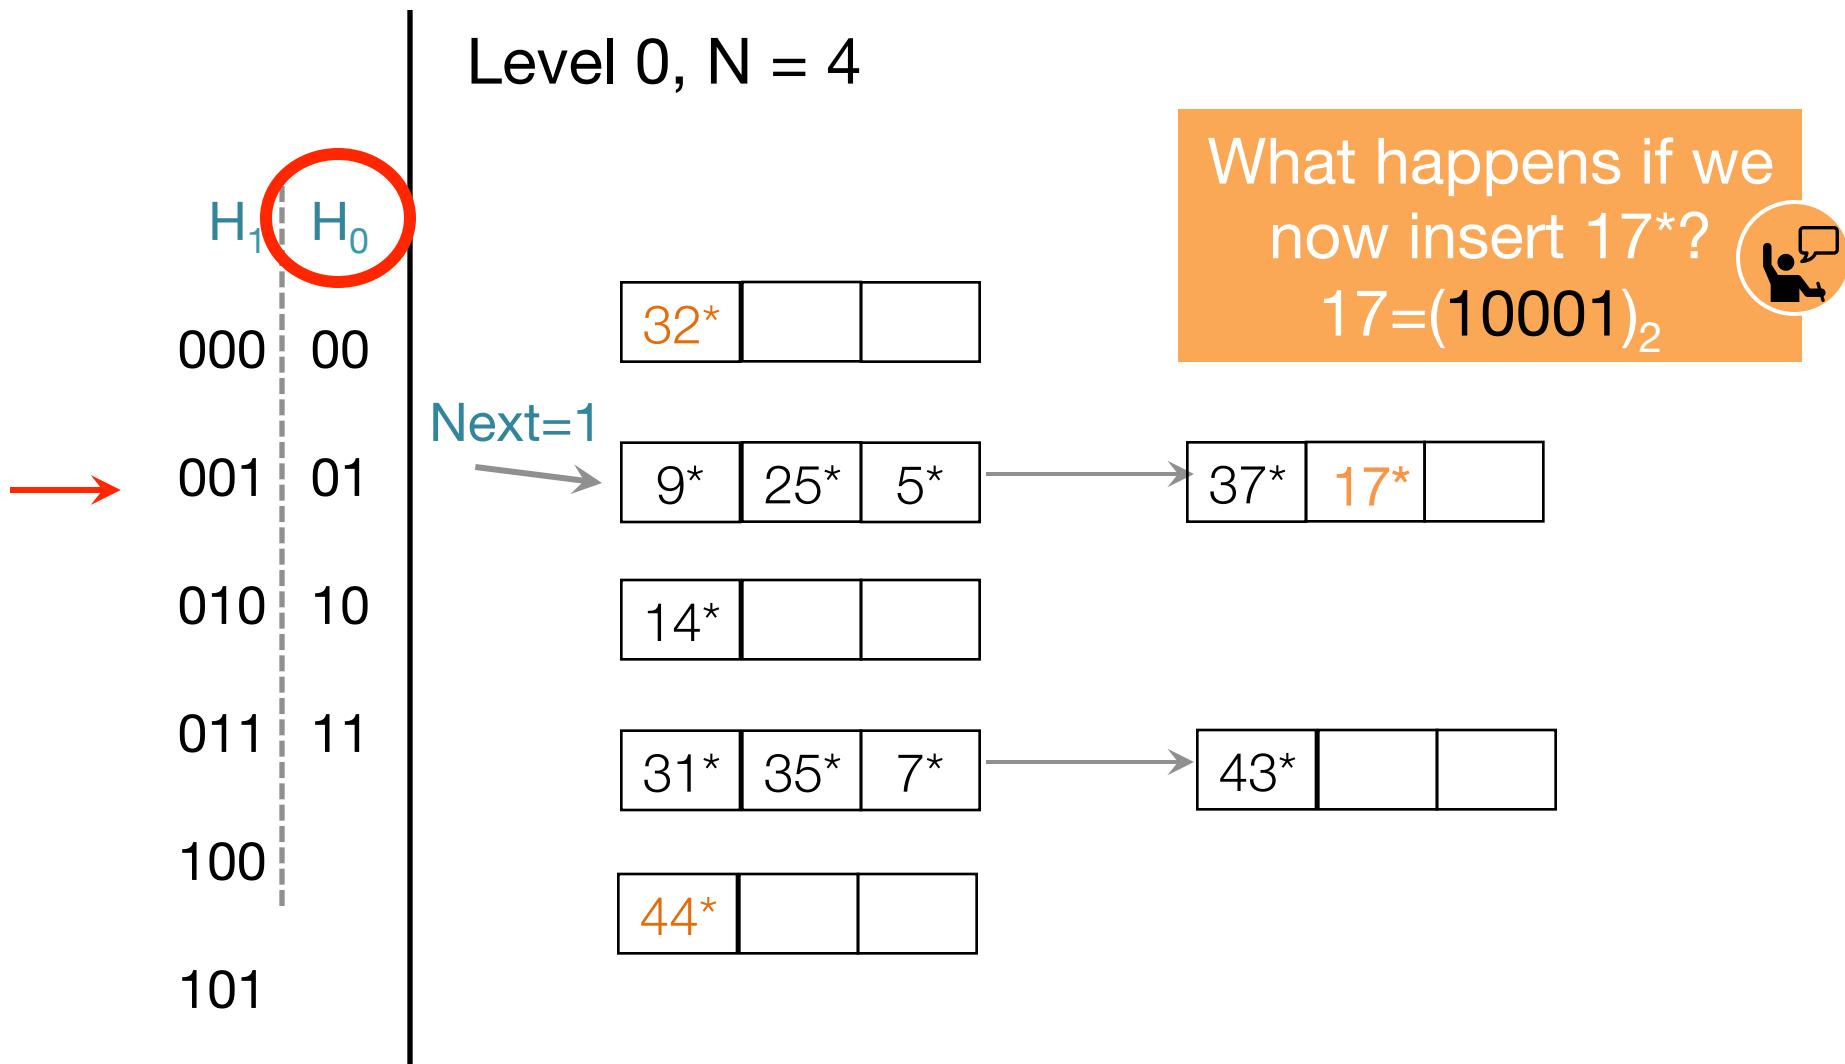
Linear Hashing: Details

- Splitting proceeds in **rounds**
 - During round *Level*, only h_{Level} and $h_{Level+1}$ are in use
- **Variables**
 - **Level**: Initialized to 0
 - **Next**: Pointer to the bucket being split
- At the beginning of round # **Level**,
 - # buckets in the file = $N_{Level} = N * 2^{Level}$
where N is initial number of buckets

Linear Hashing Example

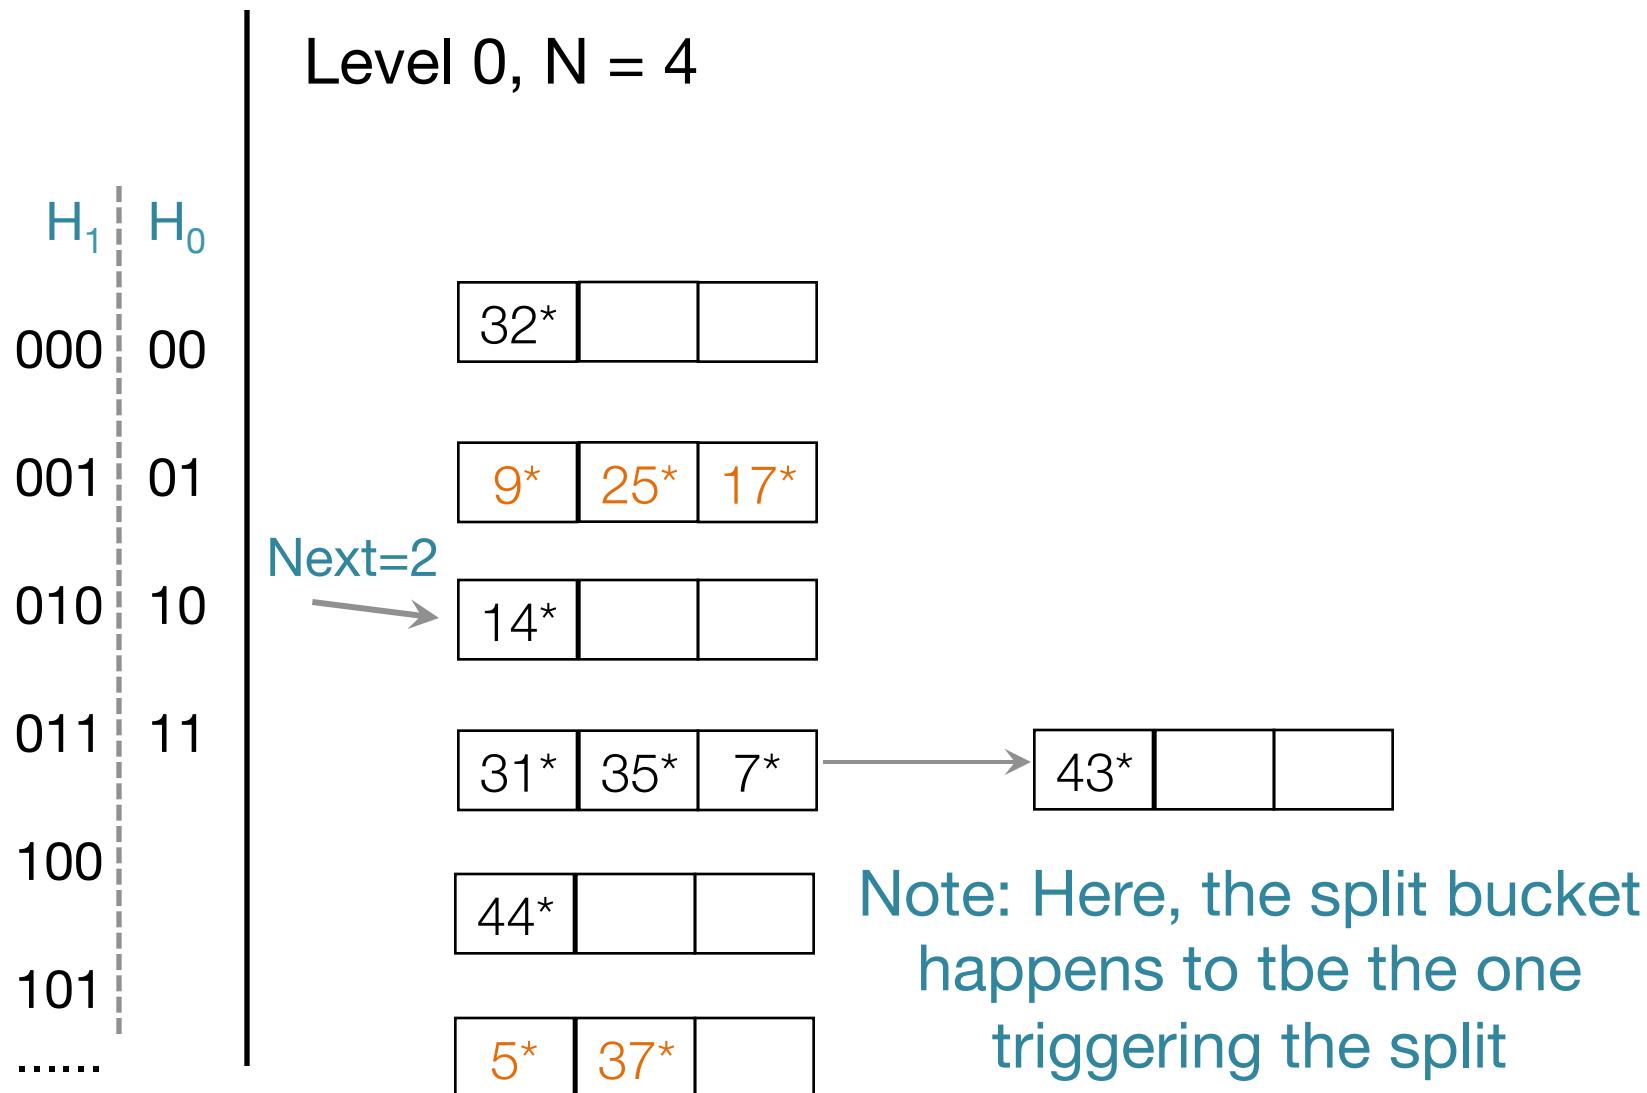


Linear Hashing Example



This is not actually stored

Linear Hashing Example



Linear Hashing Algorithm

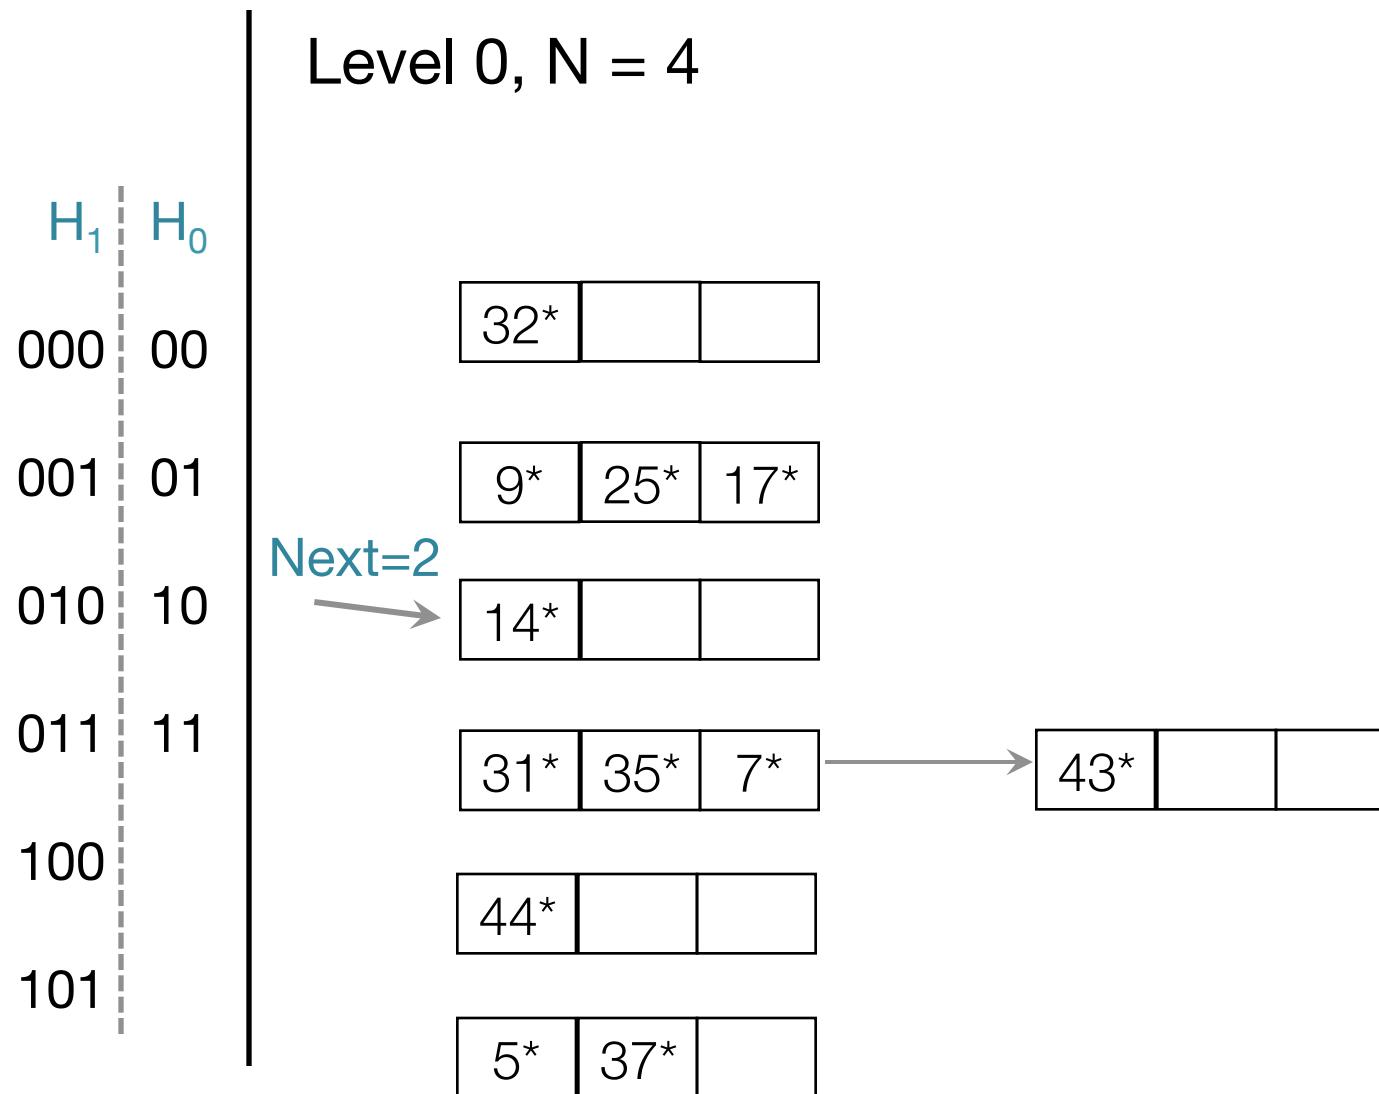
Insert:

- Find bucket by applying h_{Level}
 - Apply $h_{Level+1}$ if the value returned by $h_{Level} < Next$
- If bucket being inserted into is full:
 - Add overflow page and insert data entry
 - If bucket is $Next$, split first and then see if you still need overflow
 - Split $Next$ bucket and increment $Next$
 - If $Next = N_{Level} - 1$ and a split is triggered
 1. We split the bucket # $N_{Level}-1$ (i.e. the last bucket before the split image buckets)
 2. $Next = 0$ (reset the pointer)
 3. $Level = Level + 1$ (the size has doubled at this point)

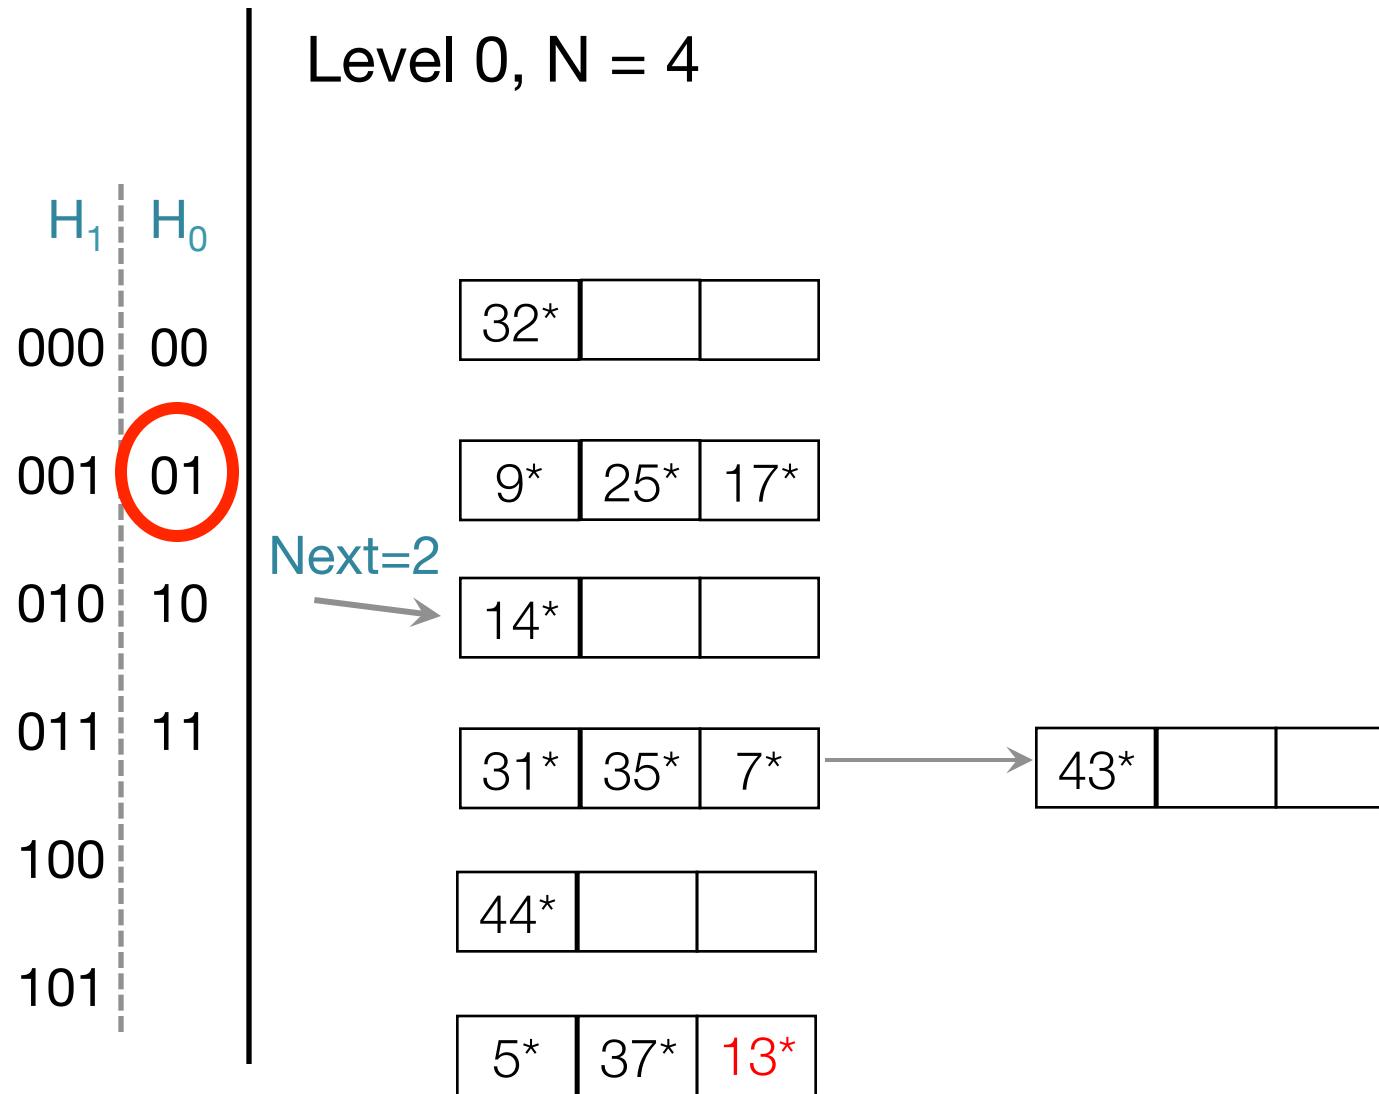
Linear Hashing (Contd.)

- Can choose any criterion to trigger split
 - e.g., split on an overflow (as in example)
 - e.g., space utilization on the page $> 90\%$
- Since buckets are split round-robin, long overflow chains don't develop!
- Deletes: see textbook

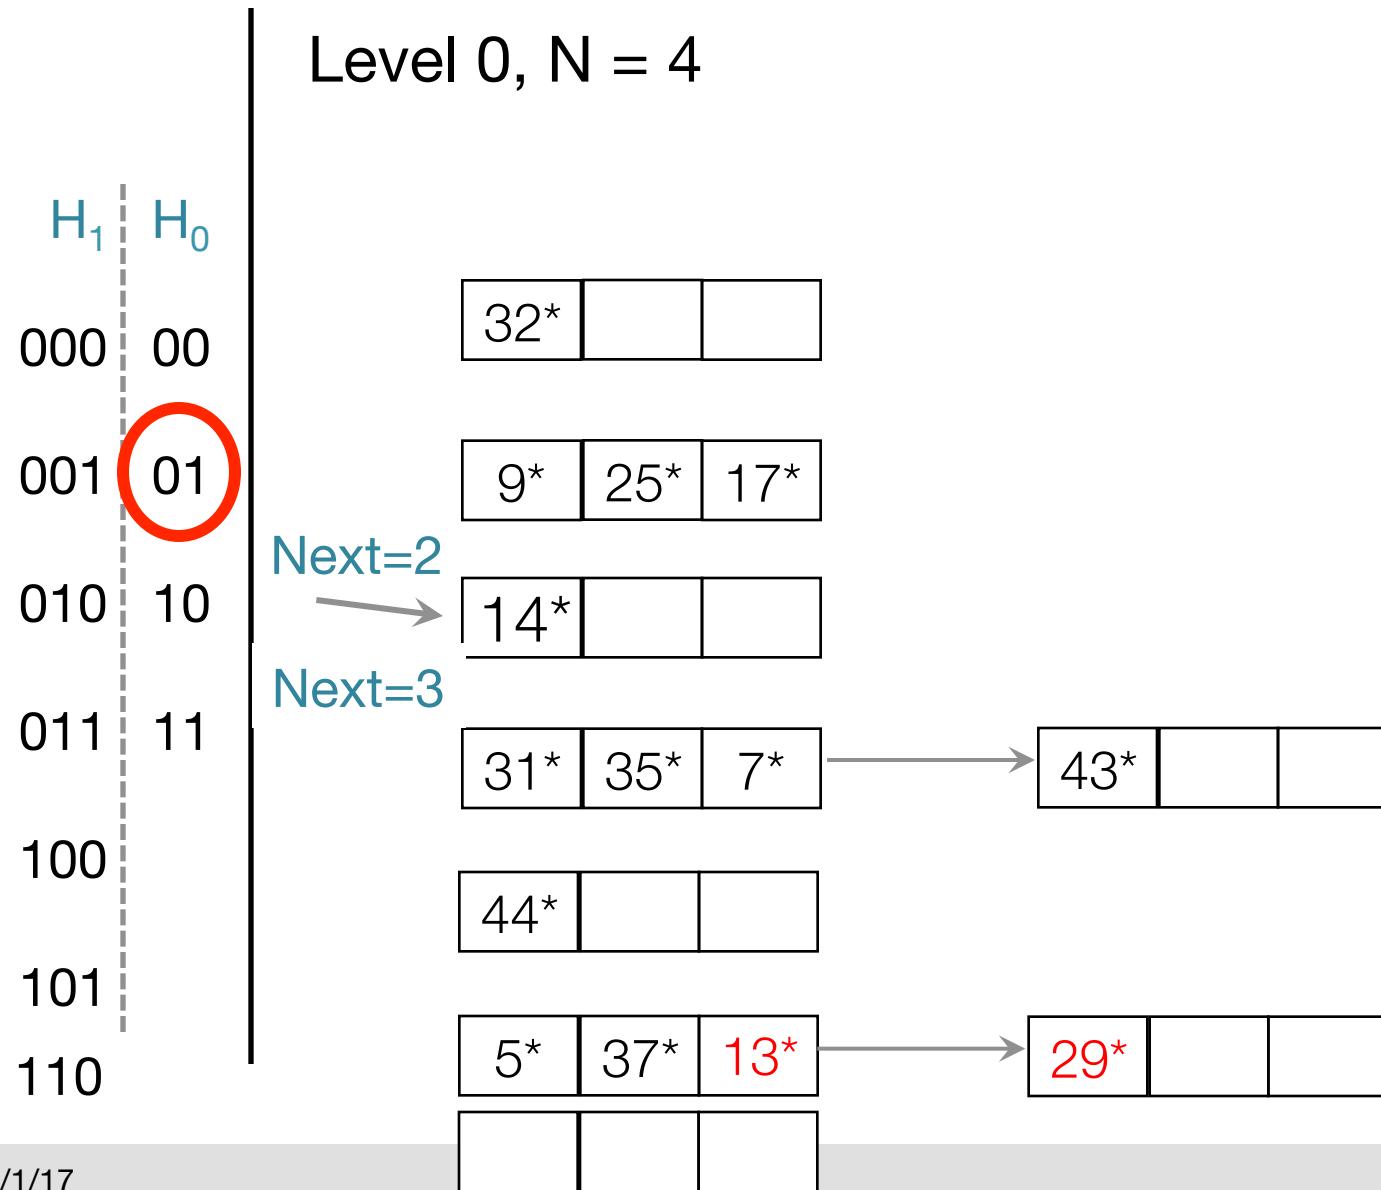
Quiz: Insert $13=(1101)_2$ & $29=(11101)_2$



Quiz: Insert $13=(1101)_2$ & $29=(11101)_2$



Quiz: Insert $13=(1101)_2$ & $29=(11101)_2$



Summary

- Discussed 3 kinds of hash-based indexes
- Static Hashing can lead to long overflow chains
- Extendible Hashing
 - Directory to keep track of buckets, doubles periodically
 - Always splits the “right” bucket
- Linear Hashing
 - Split buckets round-robin, and use overflow pages
 - Space utilization could be lower than Extensible Hashing

Optional Exercises

11.1, 11.3, 11.7, 11.9