

# Join Strategies

Chapter 12 and 14

# Operator Evaluation

- How to implement common operators?
  - ✓ • Selection
  - • Join
    - Basic strategies
    - Advanced strategies
  - Projection (optional DISTINCT)
  - Set Difference
  - Union
  - Aggregate operators (SUM, MIN, MAX, AVG)
  - GROUP BY

# Join Operator



```
SELECT  *
FROM    Reserves R, Sailors S
WHERE   R.sid = S.sid
```

- Commercial systems spend a lot of effort optimizing equality joins
  - Why?
  - What is the major source of performance cost when joining two (large) relations?
- **Cost Metric:** # of I/Os
  - We ignore final output cost since it's the same no matter which algorithm we use



**A SQL query walks up to two tables in a restaurant and asks: “Mind if I join you?”**

# Join Operator



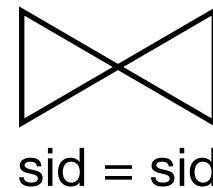
Many different ways of evaluating joins

Sailors

sid	name
1	Lucky
2	Rusty
3	Bob
4	Fred

Reserves

sid	name
1	100
1	200
3	300
4	200



How would you evaluate this join in memory?

# Simple Nested Loops Join

```
For each tuple r in R do
    for each tuple s in S do
        if r.sid == s.sid then add <r, s> to result
```

- Notation
    - $\|R\|$  = # tuples in R
    - $|R|$  = # pages in R
  - How many I/Os ?
    - $|R| + \|R\| * |S|$
- 
- Slightly different notation from textbook!

# Page-Oriented Nested Loops

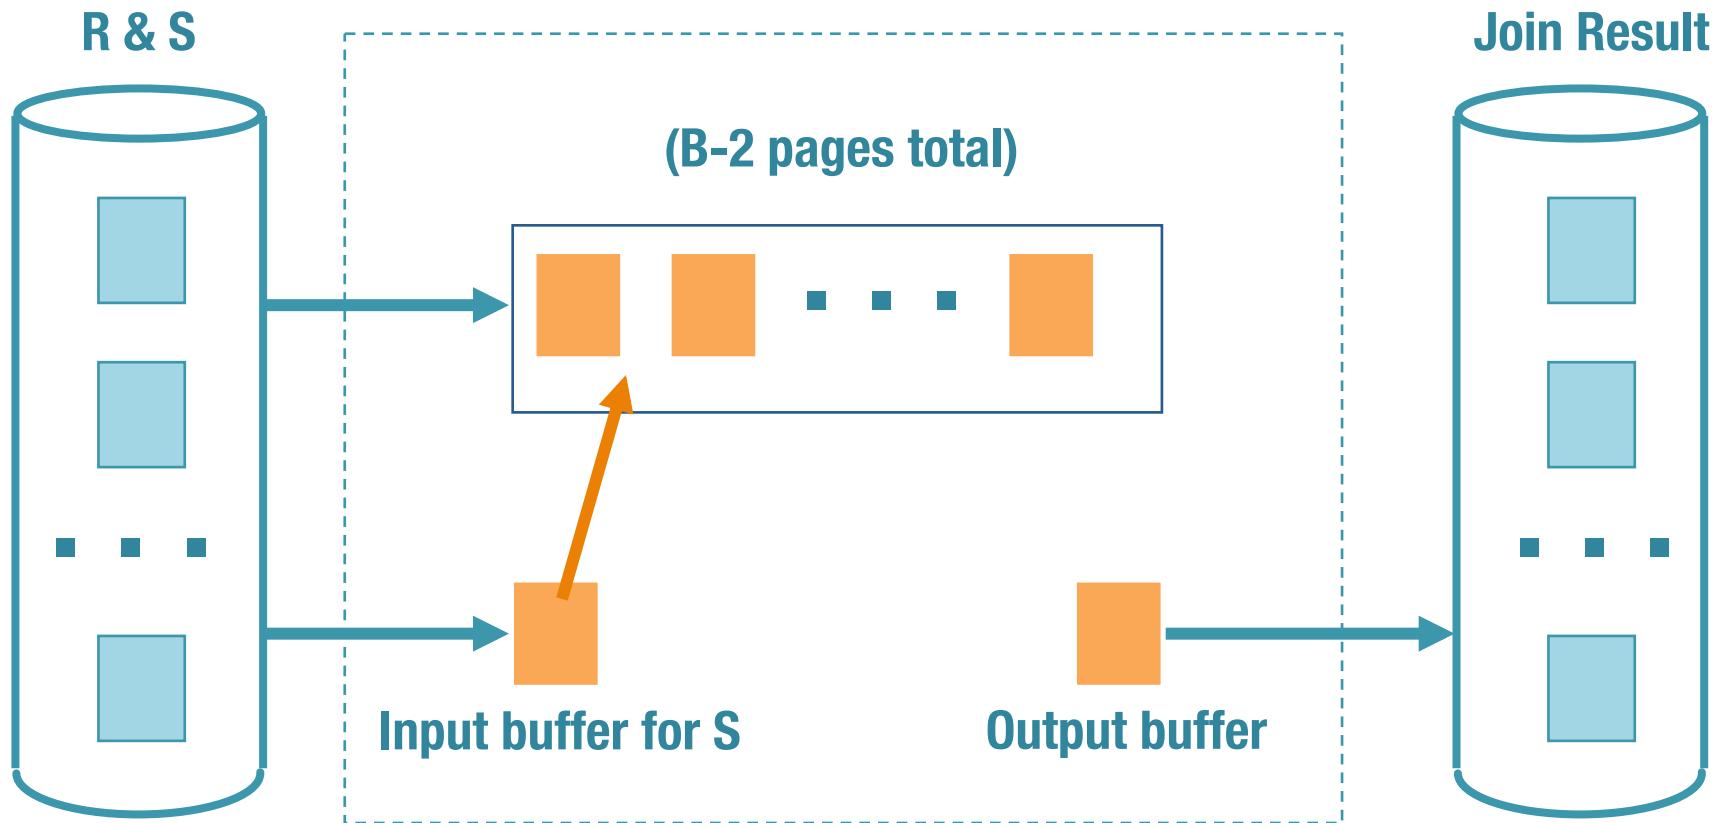
- Page-oriented Nested Loops join:
  - For each page of R, get each page of S, and join
  - How many I/Os? 
  - If S is the outer, then  $|S| + |R| * |S|$



How many buffer pages does this use?

# Block Nested Loops

Can we exploit available memory?  
Suppose we have B buffer pages available



# BNL's Cost Analysis

- Suppose we have B buffer pages available
- Use B-2 pages to hold a block of **outer R**

```
For each block of B-2 pages of R do
```

```
  for each page of S do {
```

```
    for each r in the B-2 R pages do
```

in memory

```
      for each tuple s in the S page do
```

```
        if r.sid == s.sid then add <r, s> to result
```

```
}
```

**Cost:**

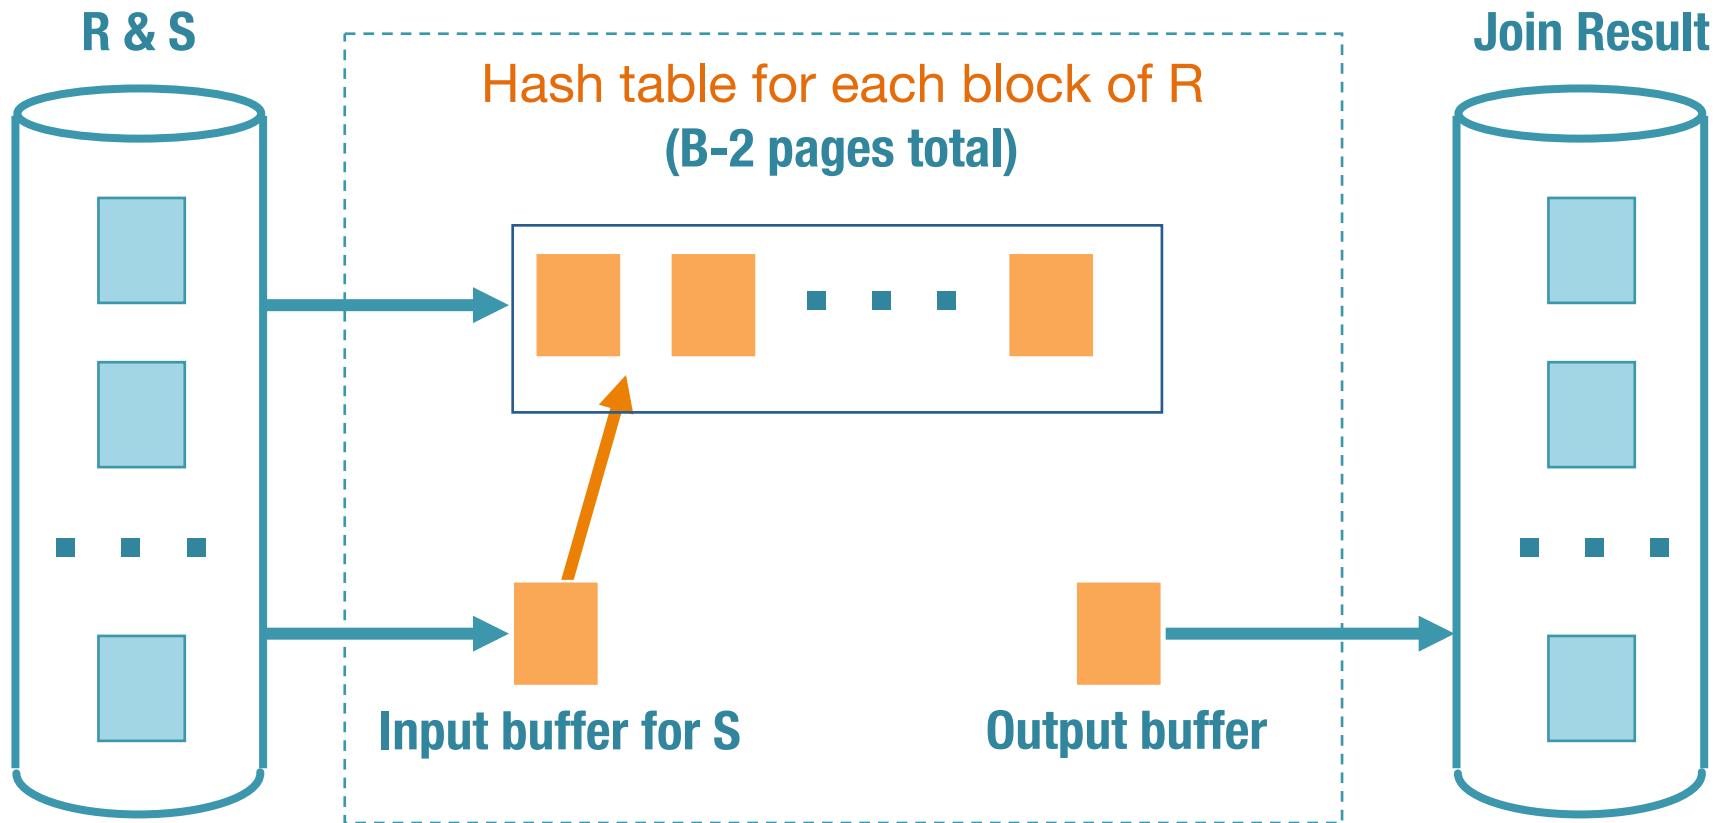
$$|R| + |S| * \left\lceil \frac{|R|}{B-2} \right\rceil$$



What is the cost  
if the smaller relation  
fits entirely in memory?

# Block Nested Loops

In-memory hash table: Small overhead for the hash table, but big savings in CPU costs (equality search)



# Quiz: PNL vs BNL

$|R| = 128$     $|S| = 64$     $B = 8$   
tuples/page for both S and R = 10

- **Page NL**

- Scan outer: 64
- Join:  $64 * 128 = 8192$
- TOTAL: 8256



In PNL, which  
rel. should be  
the outer?

- **Block NL**

- Scan outer: 64
- Join:  $\lceil 64/6 \rceil * 128 = 1408$
- TOTAL: 1472



What about  
simple nested  
loops?

# Operator Evaluation

- How to implement common operators?

✓ • Selection

• Join

✓ • Basic strategies

→ • Advanced strategies

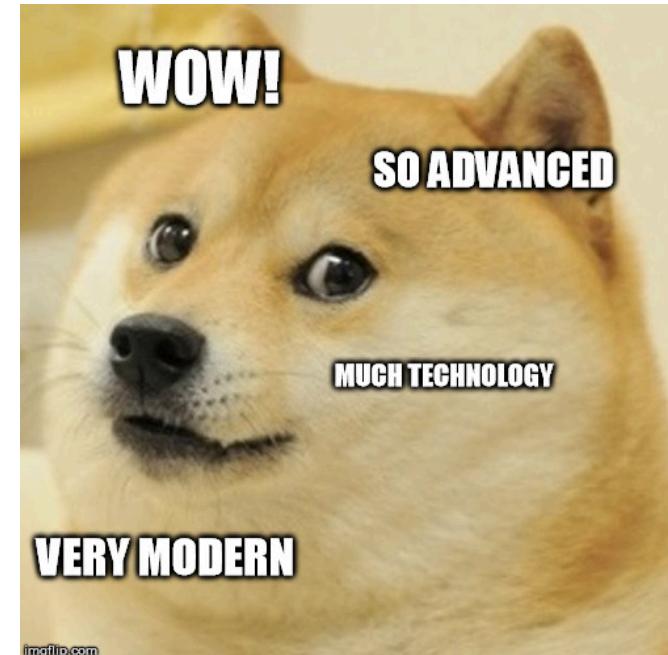
• Projection (optional DISTINCT)

• Set Difference

• Union

• Aggregate operators (SUM, MIN, MAX, AVG)

• GROUP BY



# AdvJ 1: Index Nested Loops

```
foreach tuple r in R do
    Probe Index on S.sid
        foreach matching tuple s do add <r, s> to result
```

- Let's say there is an index available on S
- Use index on join attribute of the inner relation
  - Cost:  $|R| + (|R| * \text{cost of finding matching S tuples})$
- Cost of finding matching S tuples per R tuple
  - Index cost: 1-2 for hash index
  - Index cost: 2-4 for B+ tree.
  - Record retrieval cost
    - Clustered index: one I/O (typical) for all S tuples per R tuple
    - Unclustered Index: Up to one I/O per matching S tuple

# AdvJ 2: Sort-Merge

- Sort R on the join attribute (if necessary)
- Sort S on the join attribute (if necessary)

- Merge

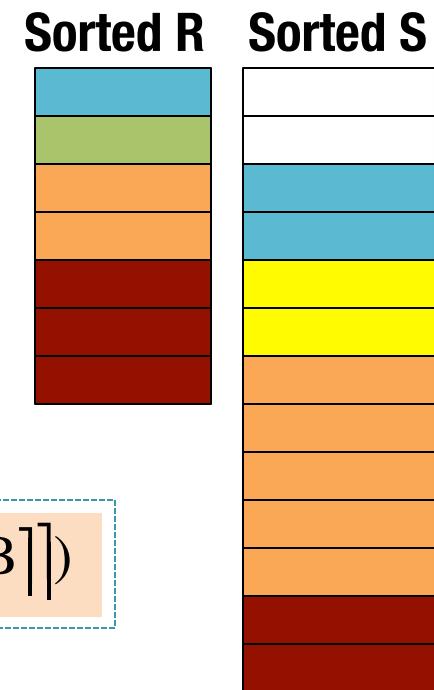


How many buffer pages  
are used in the merge?

- Sort:

$$2 \cdot |R| \cdot (1 + \lceil \log_{B-1} \lceil |R|/B \rceil \rceil) + 2 \cdot |S| \cdot (1 + \lceil \log_{B-1} \lceil |S|/B \rceil \rceil)$$

- Merge Cost:  $(|R|+|S|)$
- Merge Worst Case:  $(|R|^*|S|)$ 
  - When?
  - Backups (to disk) needed whenever # tuples in one of S's partitions exceeds buffer size



# Quiz: SM, Page NL, Block NL

- $|R| = 128 \quad |S| = 64 \quad B = 8$
- Find the cost of  $R \bowtie S$
- Sort-Merge (use two-way merge sort)
  - Sort R:  $2 * 128 (\log_2 128 + 1) = 2048$
  - Sort S:  $2 * 64 (\log_2 64 + 1) = 896$
  - Merge:  $128 + 64 = 192$
  - TOTAL: 3162
- Page NL
  - Scan outer: 64
  - Join:  $64 * 128 = 8192$
  - TOTAL: 8256
- Block NL
  - Scan outer: 64
  - Join:  $\lceil 64/6 \rceil * 128 = 1408$
  - TOTAL: 1472



In NL, which rel.  
is the outer?

# Quiz: SM, Page NL, Block NL

- $|R| = 128 \quad |S| = 64 \quad B = 8$
- Find the cost of  $R \bowtie S$
- Sort-Merge (use all the buffers for sort)

- Sort S:  $2 * 64 * (1 + \lceil \log_7 \lceil 64/8 \rceil \rceil) = 384$
- Sort R:  $2 * 128 * (1 + \lceil \log_7 \lceil 128/8 \rceil \rceil) = 768$
- Merge (input):  $64 + 128 = 192$
- TOTAL: 1344

$$\log_7 8 = 1.07$$
$$\log_7 16 = 1.42$$



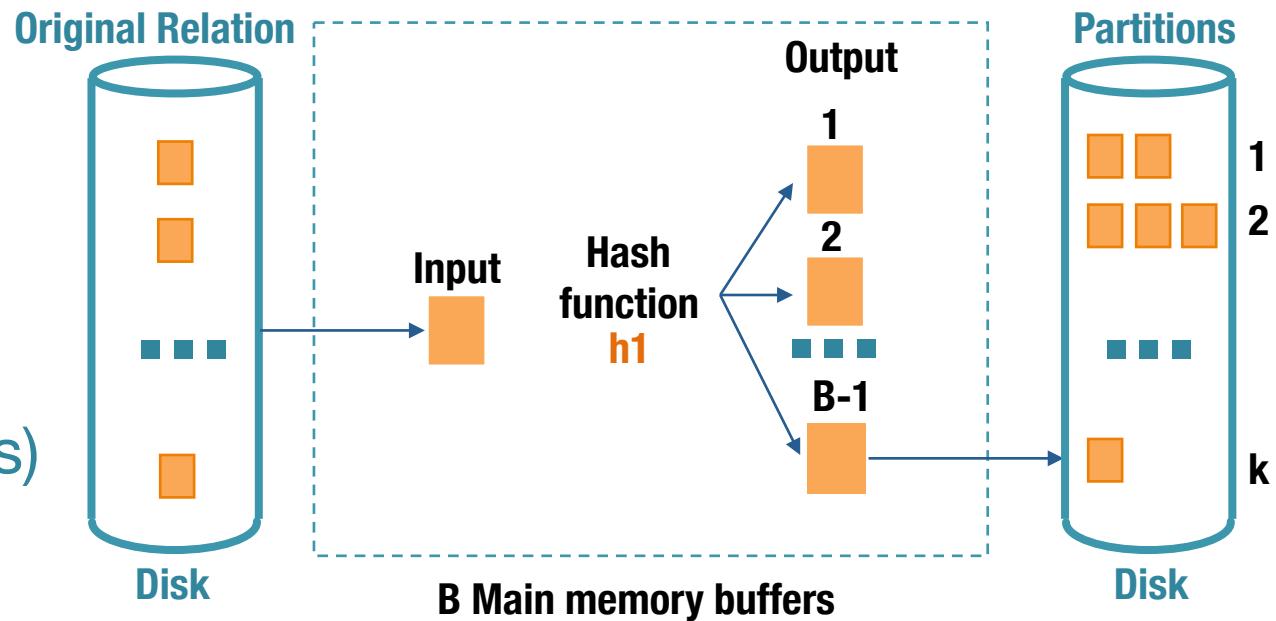
- 
- Page NL
    - Scan outer: 64
    - Join:  $64 * 128 = 8192$
    - TOTAL: 8256
  - Block NL
    - Scan outer: 64
    - Join:  $\lceil 64/6 \rceil * 128 = 1408$
    - TOTAL: 1472

# AdvJ 3: Grace Hash Join

- **Step 1:** (**Partition**) Hash both relations, R and S, on the join attribute, producing **k** disk-based partitions
  - Guarantees that R tuples can only join with S tuples in the same partition (e.g. R1-S1, R2-S2...)
- **Step 2:** (**Probe**) Read in complete partition of (smaller relation) R, and scan S partition for matches.
  - Use a different hash function  $h_2$  to reduce CPU cost of matching

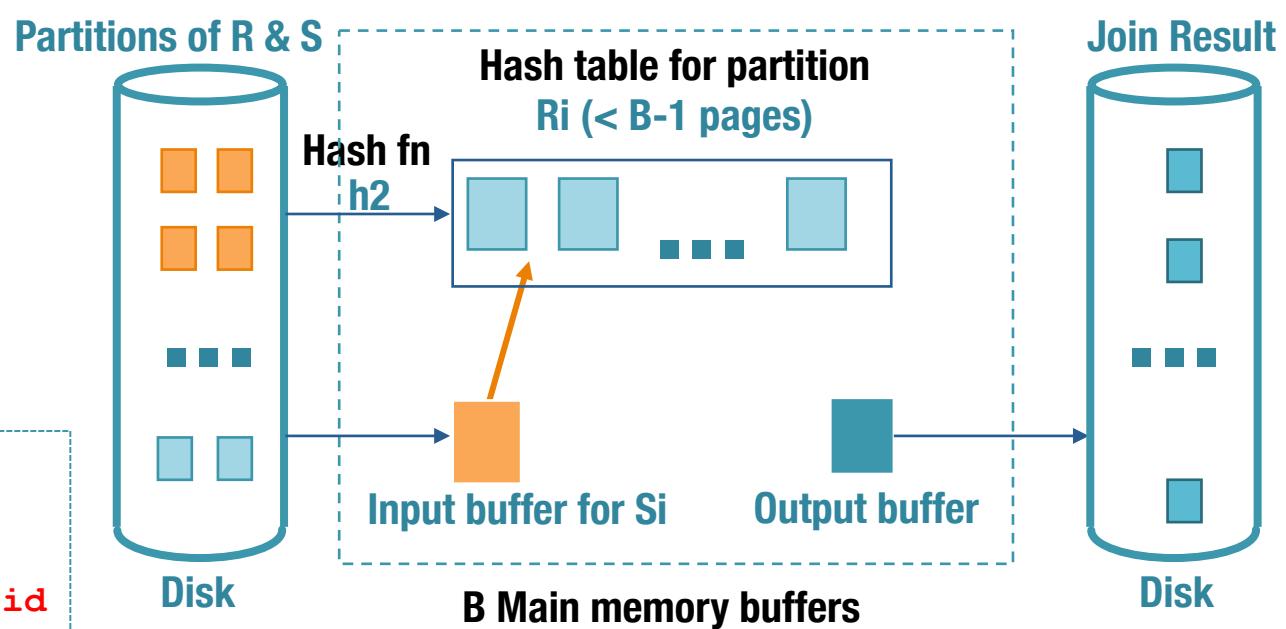
# Step 1: Partition

- $k = B-1$   
(# of buckets)
- $h1 \neq h2$



# Step 2: Probe

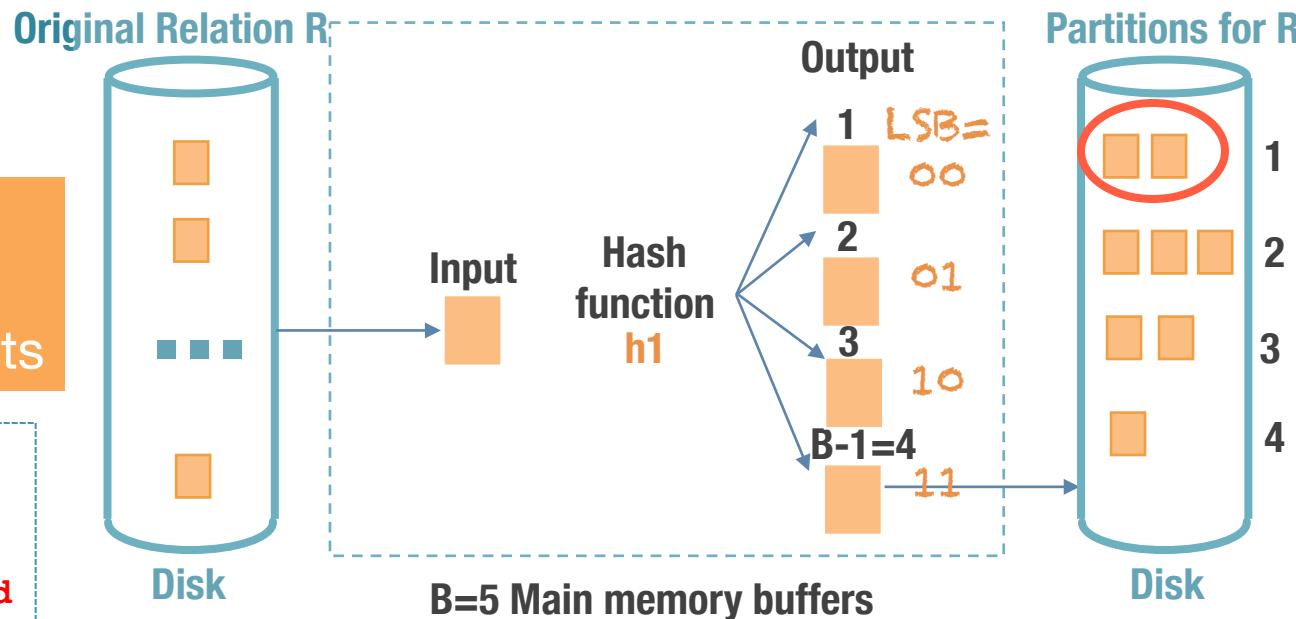
```
SELECT *
FROM   Reserves R,
       Sailors S
WHERE  R.sid = S.sid
```



# Step 1: Partition

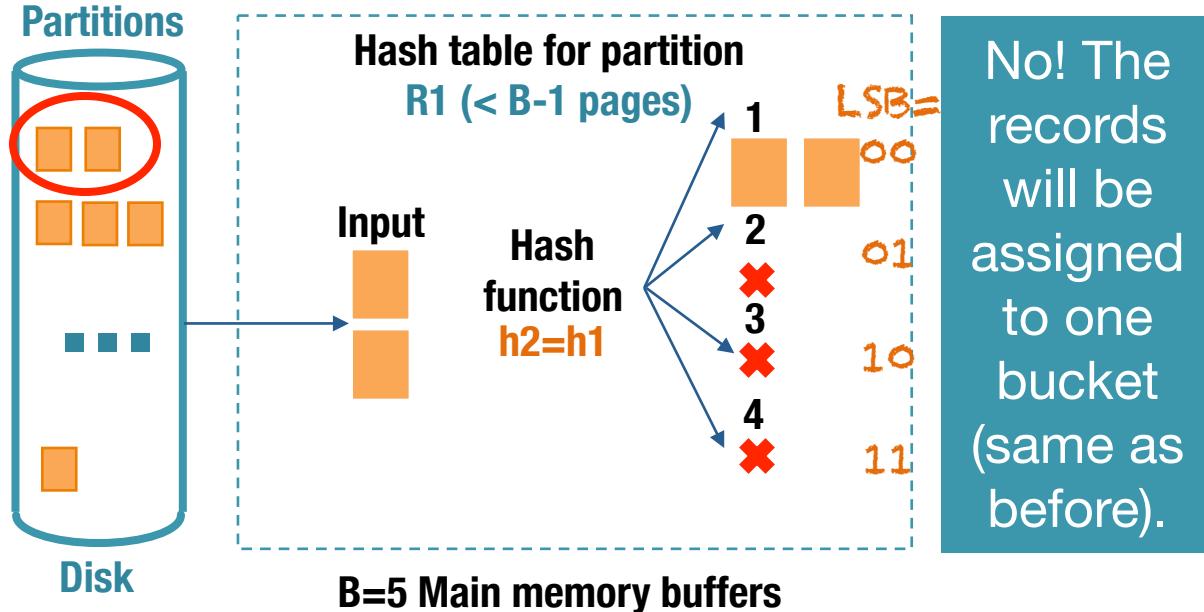
Can  $h1 = h2$ ?  
e.g.  $h1(sid) = m$   
Least Significant Bits

```
SELECT *  
FROM Reserves R,  
Sailors S  
WHERE R.sid = S.sid
```



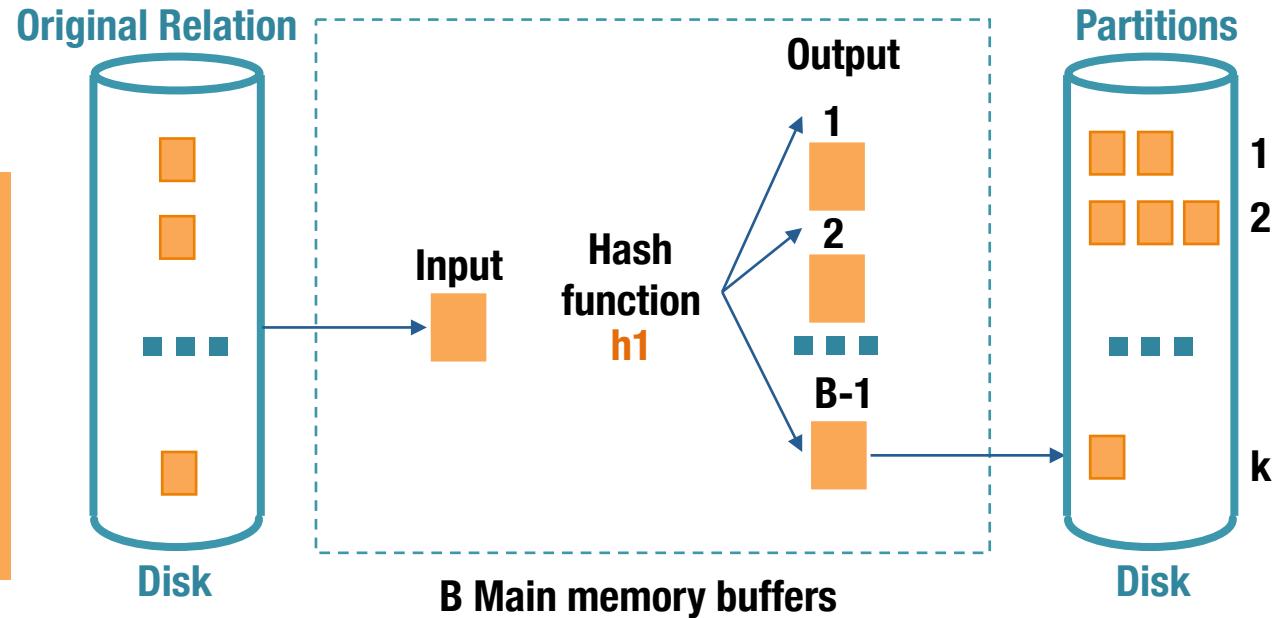
## Part of Step 2: h2 in-memory

Possible sids  
in partition 00:  
 $4 = (100)_2$   
 $12 = (1100)_2$   
 $20 = (10100)_2$   
 $32 = (100000)_2$   
...



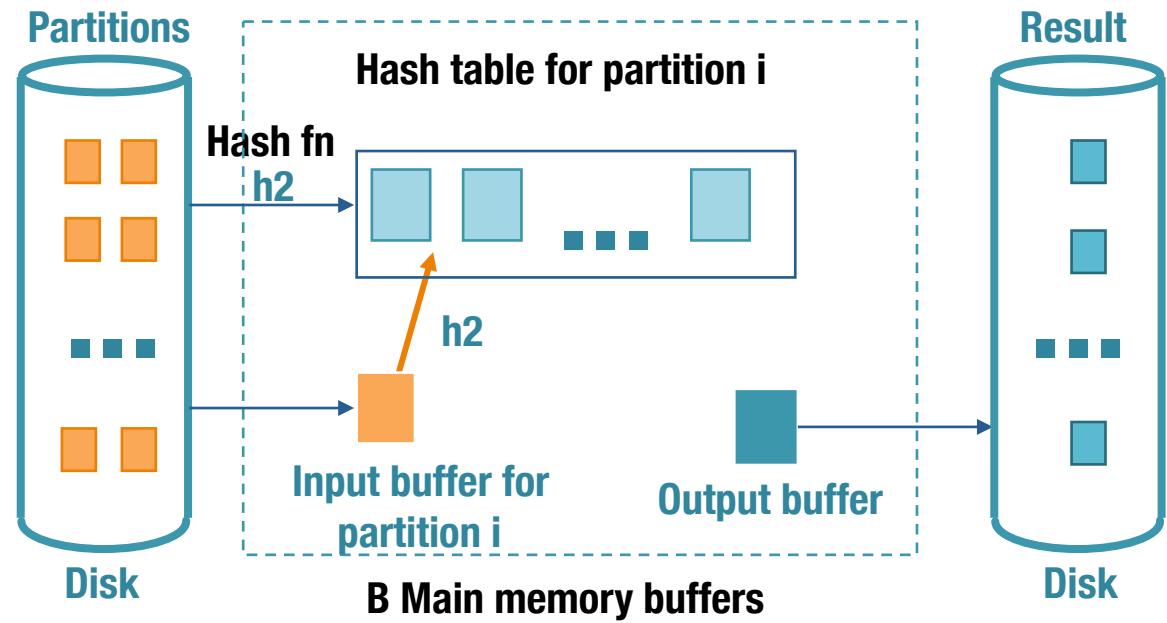
# Step 1: Partition

What if the hash table for a partition overflows, i.e. can't fit in memory?



# Step 2: Probe

Recursively apply hash-based projection technique to handle partition overflow problem



# Cost Analysis: Grace Hash Join

- Partition phase, read+write both relations;  $2(|R|+|S|)$
- In probe phase,
  - Read each partition of relation R once:
    - For that read the corresponding partition of S once
  - Total cost for all partitions:  $|R|+|S|$  I/Os.
  - Assumes each R partition fits in memory in probe phase

Grace hash join:  $3(|R| + |S|)$  I/Os

- The purpose of h2: reducing CPU costs
- Other variants of hash join exist

# How to Sort-Merge in $3(|R|+|S|)$ I/Os

- **Before:**
  - External sorting:  $2.|R|.(1+\lceil \log_{B-1} \lceil |R|/B \rceil \rceil) + 2.|S|.(1+\lceil \log_{B-1} \lceil |S|/B \rceil \rceil)$
  - Merge cost (no backups) =  $|R|+|S|$
  - Merge cost (backups) =  $|R| * |S|$
- **Can be done in  $3(|R|+|S|)$ :**
  - When larger relation  $|S| \leq B$ , i.e. each relation fits in memory
- **Rationale:**
  - R fits in memory, sort in  $2*|R|$  I/Os
  - S fits in memory, sort in  $2*|S|$  I/Os
  - Assuming no backups are required, merge would take only  $|R|+|S|$

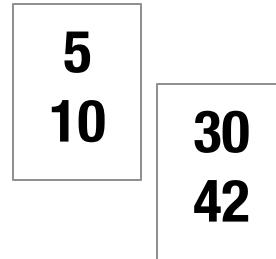
# How to Sort-Merge in $3(|R|+|S|)$ I/Os

- Before:
  - External sorting:  $2|R|(1 + \lceil \log_{B-1} \lceil |R|/B \rceil \rceil) + 2|S|(1 + \lceil \log_{B-1} \lceil |S|/B \rceil \rceil)$
  - Merge cost (no backups) =  $|R|+|S|$
  - Merge cost (backups) =  $|R| * |S|$
- **Can be done in  $3(|R|+|S|)$ :** *square root of  $|S|$  fits in memory*
  - When larger relation  $|S| \leq B$ , i.e. ~~each relation fits in memory~~
- **Example:** Assume  $|R| = 400$ ,  $|S|=10,000$  and  $B = 100$ 
  - Use replacement sort to produce runs of  $2^*M=2^*(B-2)=196$  pages long  $\approx 200=2^*B$  for each relation.
  - Then we have at most  $|R|/(2B) \approx 2$  runs of R and  $|R|/(2B) \approx 50$  runs of S.
  - Note that since  $|R| \leq |S| \leq B^2$  we have  $|R|/(2B) \leq B/2$  and  $|R|/(2B) \leq B/2$
  - Since total # of runs  $\leq B$ , we can read in memory one page per run (here 52 pages), merge, and apply join condition on-the-fly + one output page

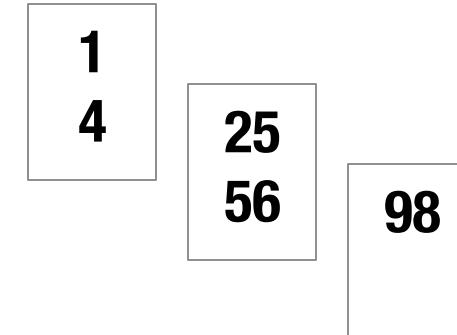
# Explanation: Why $|S| \leq B^2$ for SM join?

- **Assumption:** S is the largest relation
- The SM join algorithm that gives us the  $3(|R| + |S|)$  IO cost does the following:
  - Replacement sort for R to produce runs of  $\sim 2M = 2(B-2) \approx 2.B$  pages
  - Replacement sort of S to produce runs of  $\sim 2M = 2(B-2) \approx 2.B$  pages
  - Now each run is sorted, but each relation is NOT sorted globally. For example 2 of the runs of relation R might look like:

**Run 1:**



**Run 2:**



# Explanation: Why $|S| \leq B^2$ for SM join?

- **Assumption:** S is the largest relation
- The SM join algorithm that gives us the  $3(|R|+|S|)$  IO cost does the following:
  - Replacement sort for R to produce runs of  $\sim 2.M = 2.(B-2) \approx 2.B$  pages
  - Replacement sort of S to produce runs of  $\sim 2.M = 2.(B-2) \approx 2.B$  pages
  - Now each run is sorted, but each relation is NOT sorted globally.
  - To do the global sorting + join, we will need to bring in memory at least the first page of each run of R and S. How many runs does S (the largest relation) have? Let's call the quantity  $r_S$ .

$$r_S = (\# \text{ of pages of } S) / (\text{avg run length}) = |S| / (2.B) \quad (1)$$

- Relation R is smaller, so it will have fewer runs. As we said, we want the first page of each run to fit in memory. This means that we want:

$$r_R + r_S \leq B.$$

# Explanation: Why $|S| \leq B^2$ for SM join?

- As we said, we want the first page of each run to fit in memory. This means that we want:

$$r_R + r_S \leq B \quad (2)$$

- In the worst case R and S are of the same size, so we would have the same number of runs for both relations. In this case, relation (2) becomes:

$$2.r_S \leq B$$

- By substituting Eq. (1), we obtain:

$$\begin{aligned} 2.|S| / (2.B) &\leq B \Rightarrow \\ |S| &\leq B^2 \end{aligned}$$

So, if  $\sqrt{|S|}$  fits in memory, we can run the optimized SM join algorithm.  $\square$

# Reminder: Replacement Sort (for Pass 0)

- Start by reading a page from file into the input buffer
- Copy records from input buffer to current set
- Repeatedly pick smallest value from current set that is greater than largest value in output buffer
  - Write to output buffer (run). If buffer full, output
- Start a new run when no value in current set is larger than all values in output

On average, produces runs of size  $2M$ , i.e.  $2^*(B-2)$  pages.

# Join Algorithms: A Comparison

- Hash-Join vs. Sort-Merge Cost:
  - Sort-merge can be optimized to  $3(|R|+|S|)$  I/Os
    - When **larger** relation  $|S| \leq B^2$
    - **Memory requirements:** **larger** relation  $|S| \leq B^2$
  - Hash join costs  $3(|R|+|S|)$  I/Os
    - If each partition of the **smaller** relation fits in memory
    - **Memory requirements:** **Smaller** relation  $|R| \leq B^2$
- Hash Join superior if relation sizes differ greatly
- Hash Join is highly parallelizable
- Hash join worse if partitioning is skewed
- Sort-Merge results already sorted (if matters)

# General Join Conditions



- Equalities over several attributes is doable with all Joins!  
e.g. `R.sid=S.sid AND R.rname=S.sname`
  - Block Nested Loop
    - always works
  - Index Nested Loop:
    - index on `<sid, sname>` or `sid` or `sname`.
    - Usually more I/Os than BNL
  - Sort-merge join:
    - sort Reserves on `<sid, rname>` and Sailors on `<sid, sname>`
  - Hash join:
    - hash Reserves on `<sid, rname>` and Sailors on `<sid, sname>`

# General Join Conditions

- Inequality conditions (e.g.  $R.rname < S.sname$ ):
  - Block Nested Loop: still works
  - For Index Nested Loop, needs B+ tree index
    - Usually worse than Block Nested Loops
  - Sort-Merge and Hash-Join are not applicable



# Summary: Join Strategies

- Basic Join Strategies:
  - Simple nested loops
  - Page nested loops
  - Block nested loops
- Advanced Join Strategies:
  - Index Nested Loops
  - Sort-Merge
  - Hash-Join

# Optional Exercises

- 12.1 (1-4), 12.3, 12.5
- 13.1, 13.3
- 14.1 (2, 3, 4, 6, 7, 8, 9, 10), 14