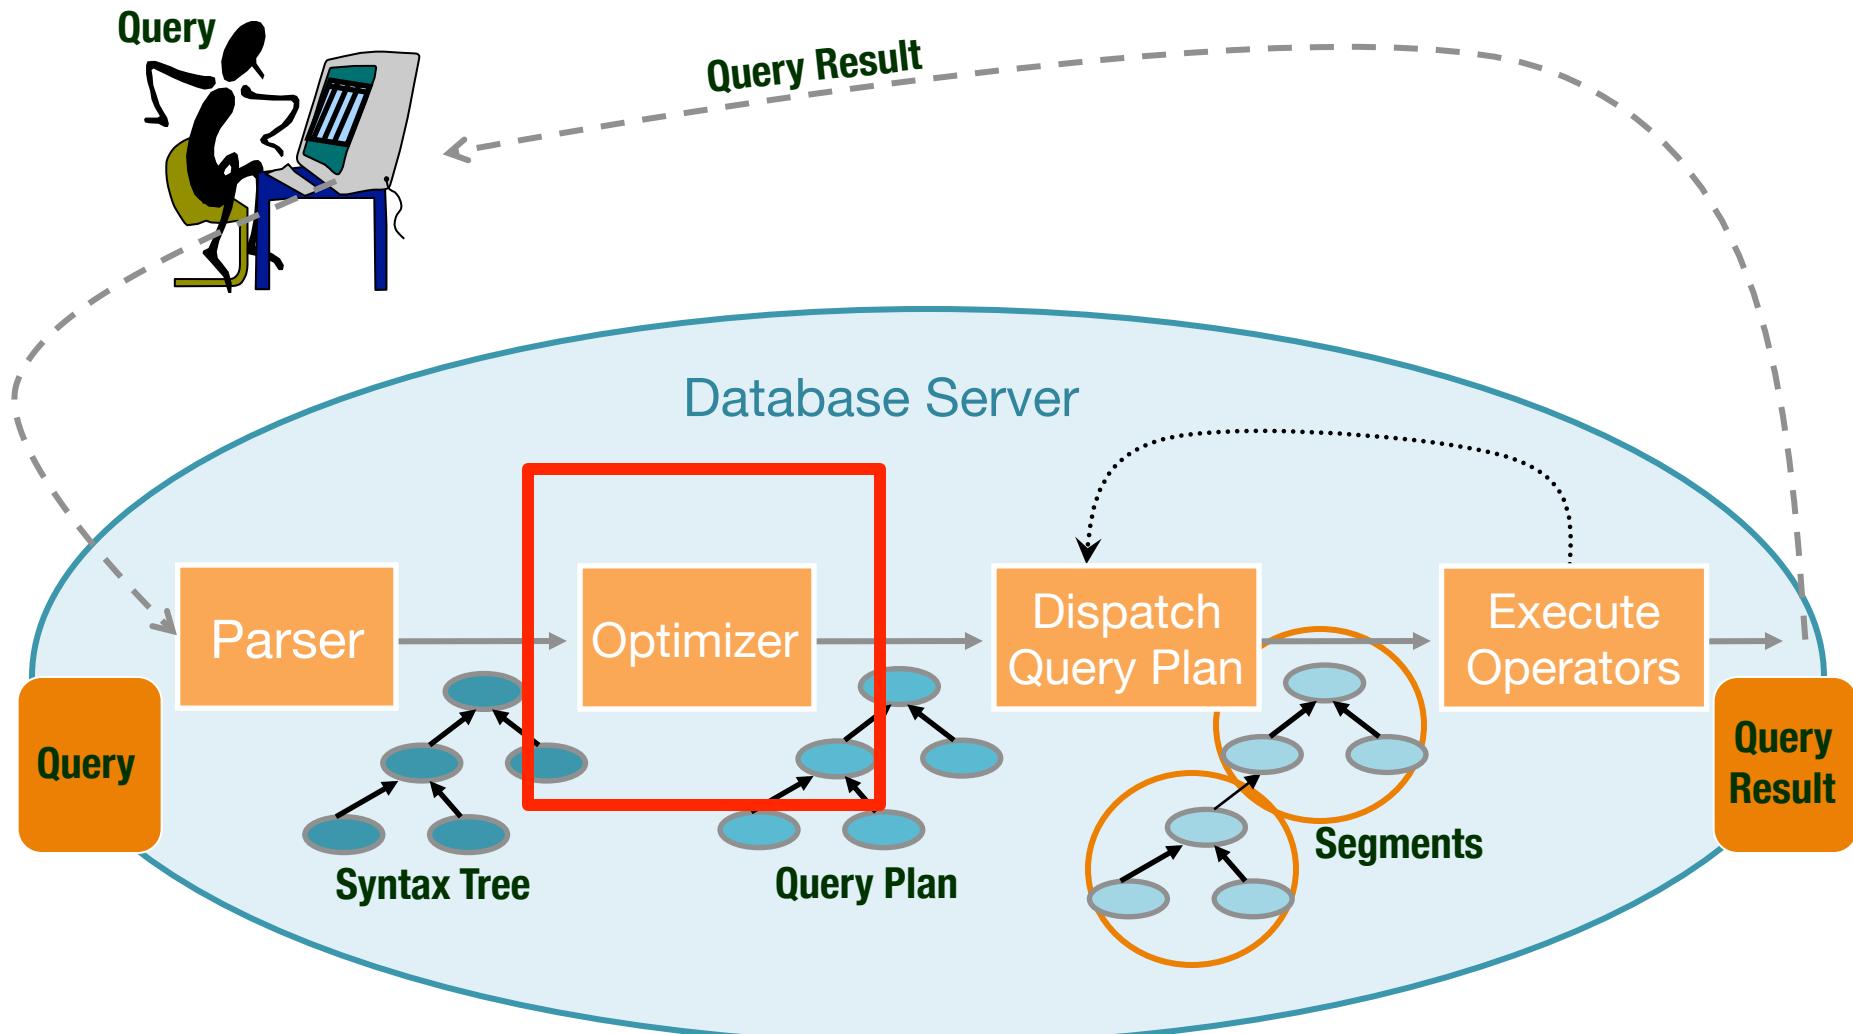


Query Optimization

Chapter 15

Query Execution Life-Cycle



Query Optimization

Query optimizer selects an evaluation plan with the least cost in two steps:

1. Plan Enumeration

Today

- Different query plans
 - Different implementations (i.e., evaluation algorithms) for each operator

2. Cost Estimation

- Cost of each operator
- Overall cost of the plan

Previous
lectures

Annotated RA Expressions

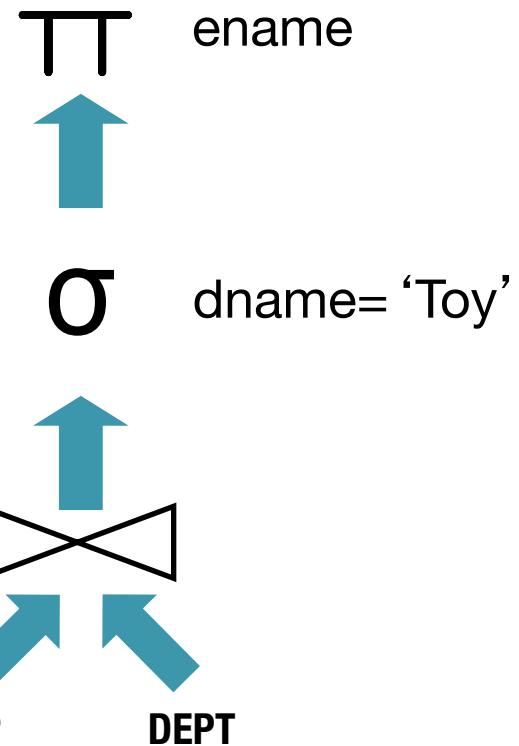
- Which algorithm to use for each operator
- Where **Intermediate Results** are:
 - **Pipelined:** Tuples resulting from one operator fed directly into the next
 - **Materialized:** Create a temporary table to store intermediate results

Example: RA Tree

```
EMP (ssn, ename, addr, sal, did)
```

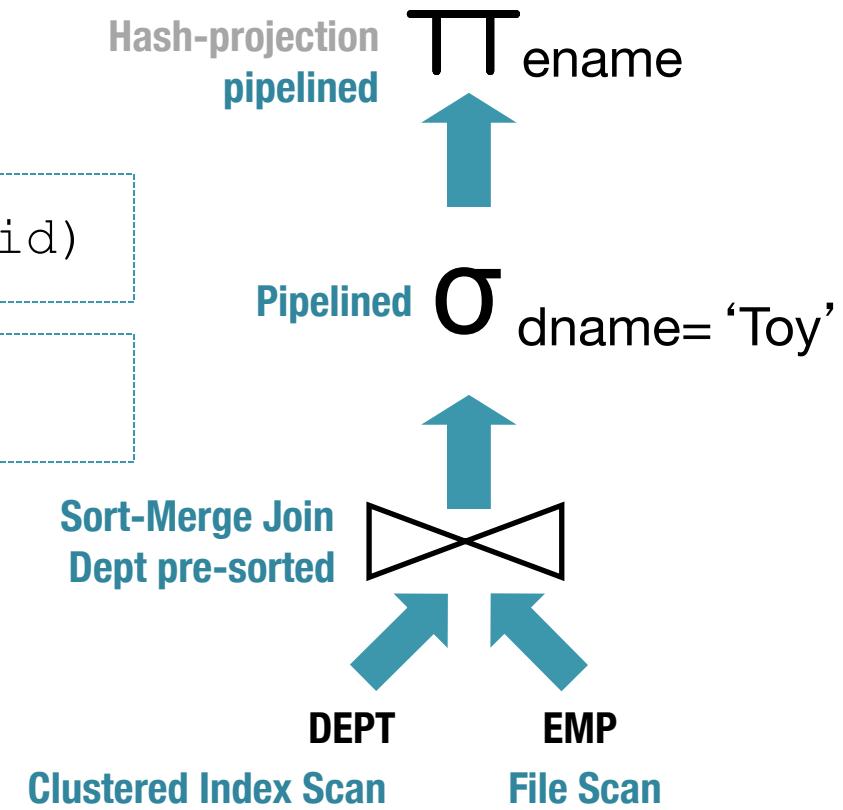
```
DEPT (did, dname, floor, mgr)
```

```
SELECT E.ename  
FROM Emp E, Dept D  
WHERE D.dname = 'Toy'  
AND D.did = E.did
```



Example: Annotated RA Tree

```
EMP (ssn, ename, addr, sal, did)  
  
DEPT (did, dname, floor, mgr)  
  
SELECT E.ename  
FROM Emp E, Dept D  
WHERE D.dname = 'Toy'  
AND D.did = E.did
```



Annotated RA Tree

Extended RA

```
HAVING MAX(SALARY) > 2 ( ... )  
GROUP BY D.did ( ... )
```

```
Select  
From  
Where  
Group By  
Having
```

```
E.did, Max (E.Salary)  
Emp E  
addr = 'Palo Alto'  
E.did  
count(*) > 10
```

```
 $\Pi_{did, \text{Max}(\text{salary})}$   
Havingcount(*) > 10 (Group Bydid ( $\sigma_{addr = 'Palo Alto'} \text{EMP}$ ))
```

Simplification: Only optimize the σ , Π , \times

- Project Group By/Having attributes
- Choose from different aggregate algorithms

RA Equivalence - Selections

$$\sigma_{P_1} (\sigma_{P_2} (R)) \equiv \sigma_{P_2} (\sigma_{P_1} (R)) \quad (\sigma \text{ commutativity})$$

$$\sigma_{P_1 \wedge P_2 \dots \wedge P_n} (R) \equiv \sigma_{P_1} (\sigma_{P_2} (\dots \sigma_{P_n} (R))) \quad (\text{cascading } \sigma)$$

Selection operation is commutative and multiple selections on the same relation can be combined into a single selection

RA Equivalence – Projections

$$\Pi_{a_1}(R) \equiv \Pi_{a_1}(\Pi_{a_2}(\dots \Pi_{a_k}(R) \dots)), \quad a_i \sqsubseteq a_{i+1} \text{ (cascading } \Pi\text{)}$$

- In the above, each a_i is a set of columns.
- For example: Let's say R has columns c_1, c_2, c_3, c_4

$$\Pi_{c_1, c_3}(R) \equiv (\Pi_{c_1, c_3}(\Pi_{c_1, c_2, c_3}(R)))$$

- Basically, we are eliminating one column at a time
- The above is useful when optimizing expressions that involve both projections and joins

RA Equivalence: Cross-Products & Joins

$$R \bowtie S \equiv S \bowtie R \text{ (commutativity)}$$
$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T \text{ (associativity)}$$
$$R \times S \equiv S \times R \text{ (commutativity)}$$
$$R \times (S \times T) \equiv (R \times S) \times T \text{ (associativity)}$$

- => joins and cross products can be performed in any order
- Same rows will result. Column order not important semantically
 - (SQL can reorder the columns at final presentation time to the user, if column order is relevant for UI purposes)

RA Equivalence – Multiple Ops

$\Pi_A(\sigma_c(R)) \equiv \sigma_c(\Pi_A(R))$ (if selection only involves attributes in the set A)

$\sigma_p(R \times S) \equiv (R \bowtie_p S)$

$\sigma_p(R \times S) \equiv \sigma_p(R) \times S$ (if p is only on R)

$\sigma_p(R \bowtie S) \equiv \sigma_p(R) \bowtie S$ (if p is only on R)

Key ideas: When possible, consider doing:

- Joins rather than cross-products
- Selections before joins
- Projections before selections

RA Equivalence – Multiple Ops

$$\sigma_P(R \ X \ S) \equiv \sigma_{P4}(\sigma_{P1}(R) \bowtie_{p3} \sigma_{P2}(S))$$

Note: $P=p1 \wedge p2 \wedge p3 \wedge p4$

Idea: Replace P by a cascade of conditions p1, p2, p3, and p4 such that:

P1: conditions only on R

P2: conditions only on S

P3: join conditions with equality on R and S

P4: other conditions involving both R and S columns

Example

$$\sigma_P (R \times S) \equiv \sigma_{P4} (\sigma_{P1}(R) \bowtie_{p3} \sigma_{P2}(S))$$

Note: $P=p1 \wedge p2 \wedge p3 \wedge p4$

$$\begin{aligned}\sigma_{\text{major=CSE} \wedge R.\text{umid}=S.\text{umid} \wedge \text{grade}>C} (R \times S) &\equiv \\ \sigma_{\text{true}} (\sigma_{\text{major=cse}}(R) \bowtie_{\text{umid}} \sigma_{\text{grade}>C}(S))\end{aligned}$$

umid	sname	major
23	Alice	CSE
71	Bob	ECE
11	Mary	CSE
13	John	CSE

umid	course	semes	grade
13	484	W17	A+
23	484	W17	A+
71	484	W17	A+
23	482	W17	C
71	482	W17	D
11	482	W17	D-

RA Equivalence – Multiple Ops

$$\Pi_P(R \times S) \equiv \Pi_{P1}(R) \times \Pi_{P2}(S)$$

Columns p in the cross-product consist of columns p1 from R and columns p2 from S



Why is this useful?

$$\Pi_P(R \bowtie_c S) \equiv \Pi_P(\Pi_{P1}(R) \bowtie_c \Pi_{P2}(S))$$

p1 are attrs of R that appear in p or c

p2 are attrs of S that appear in p or c

Example

$$\Pi_P (R \bowtie_c S) \equiv \Pi_P (\Pi_{P1}(R) \bowtie_c \Pi_{P2}(S))$$

$\Pi_{\text{surname, course}} (R \bowtie_{\text{umid}} S)$

$\equiv \Pi_{\text{surname, course}} (\Pi_{\text{surname, umid}}(R) \bowtie_{\text{umid}} \Pi_{\text{course, umid}}(S))$

umid	sname	major
23	Alice	CSE
71	Bob	ECE
11	Mary	CSE
13	John	CSE

umid	course	semes	grade
13	484	W17	A+
23	484	W17	A+
71	484	W17	A+
23	482	W17	C
71	482	W17	D
11	482	W17	D-

RA Equivalence – More Rules

$$\prod_{A_1, A_2, \dots, A_n} (\sigma_P (R)) \equiv \prod_{A_1, A_2, \dots, A_n} (\sigma_P (\prod_{A_1, \dots, A_n, B_1, \dots, B_M} R))$$

B1 ... BM attributes in P

$$\prod_{umid} (\sigma_{\text{major}=ECE} (R)) \equiv \prod_{umid} (\sigma_{\text{major}=ECE} (\prod_{umid, \text{major}} R))$$

umid	sname	major
23	Alice	CSE
71	Bob	ECE
11	Mary	CSE
13	John	CSE

RA Equivalence

- Additional equivalences possible when `union`, `intersection`, `set difference`, etc. are considered
- We do not discuss those further

Types of Optimization

- Rule-Based
 - If we know that plan A and plan B are equivalent, and that plan A is likely to be cheaper, then we should replace plan B with plan A.
 - E.g. Push selections in.
 - E.g. Push projections in.
- Cost-Based
 - If we can get (good) estimates of the cost of equivalent plans A and B, we can prefer the one that has lower cost.
 - Can use this even if there is no universal best choice.
 - E.g. choosing between join methods.
 - Major innovation in database systems

Query Optimization – Main Issues

Q1: Which (equivalent) plans do we consider?

Q2: How do we estimate the cost of a plan?

- **Ideal Goal:** Find the best plan
- **Pragmatic Goal:** Avoid worst plans!
- Typical optimizations
 - Re-ordering joins
 - “Pushing” selections and projections
 - e.g., avoid cartesian products

Cost Based Optimization

- Most popular currently; works well for < 10 joins
- **Cost estimation:** Approximate at best
 - Use statistics from systems catalogs
 - Combination of CPU and I/O costs
 - But this class focuses on I/O costs only
 - We also use **System R** approach
 - very inexact but ok in practice (better techniques are known now)

Estimating the Cost of a Plan

1. Estimate *cost* of each operation in plan tree

- Depends on input cardinalities (# of rows)
- Algorithm cost (see previous lectures)

2. Estimate *size of result*

- Use information about the input relations
- For selections and joins, assume independence of predicates



Why care about
result sizes?

Collecting & Maintaining Statistics

- Statistics stored in the catalogs (pg_stats & pg_class in Postgres)
 - Relation
 - Cardinality (# rows)
 - Size in pages
 - Index
 - Cardinality (# distinct keys)
 - Size in pages
 - Height
 - Range
- Catalogs update periodically
 - Can be slightly inconsistent
- Commercial systems use histograms
 - More accurate estimates

Estimating Output Size

```
SELECT attribute list  
FROM relation list  
WHERE term1 AND ... AND termk
```

Question: What is the cardinality of the result set?

- Max # tuples: product of input relation cardinalities
- Each term “filters” out some tuples: Reduction factor
- Result cardinality = Max # tuples * product of all RF’s
- Assumption: terms are independent!
- Term $col=value$ RF: $1/N\text{Keys}(I)$, given index I on col
- Term $col1=col2$ RF: $1/\text{MAX}(N\text{Keys}(I1), N\text{Keys}(I2))$
- Term $col>value$ RF: $(\text{High}(I)-value)/(\text{High}(I)-\text{Low}(I))$
If no index, Nkeys = # distinct values

Plan Enumeration

- Two main cases:
 - Single-relation plans
 - Multiple-relation plans
- Single-relation plan (no joins). Access Plans:
 - file scan
 - index scan(s): Clustered, Unclustered 
 - More than one index may “match” predicates
 - e.g. Clustered index I matching one or more selects:
Cost: $(N\text{Pages}(I) + N\text{Pages}(R)) * \text{product of RF's of matching selects}$
 - Choose the one with the least estimated cost.
 - Merge/pipeline selection and projection (and aggregation)
 - RID intersection techniques
 - Index aggregate evaluation (e.g. min or max age)

Quiz: Cost of different plans



EMP (ssn, ename, addr, sal, did)

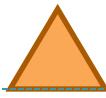
Estimate the cost of different access methods:

- Index on did:
 - Clustered index: ?
 - Unclustered index: ?
- Index on sal:
 - Clustered index: ?
 - Unclustered index: ...
- File scan: ?

```
SELECT E.ename  
FROM   Emp E  
WHERE  E.did=8  
AND    E.sal > 40K
```

1,000 data pages, 10K tuples
100 pages in B+-tree
depts: 10
Salary Range: 10K – 200K

Quiz: Cost of different plans



EMP (ssn, ename, addr, sal, did)

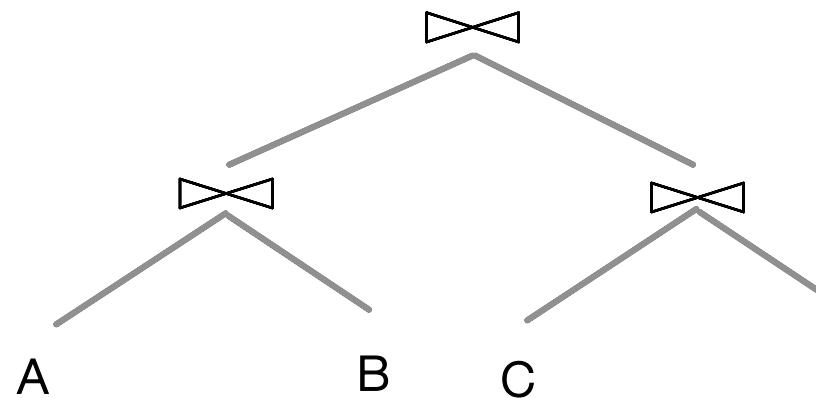
- Index on did:
 - Tuples Retrieved: $(1/10) * 10,000$
 - Clustered index: $(1/10) * (100+1,000)$ pages
 - Unclustered index: $(1/10) * (100+10,000)$ pages
- Index on sal:
 - Clustered index: $(200-40)/(200-10) * (100+1,000)$ pages
 - Unclustered index: $(200-40)/(200-10) * (100+10,000)$ pages
- File scan: 1,000 pages

```
SELECT E.ename  
FROM   Emp E  
WHERE  E.did=8  
AND    E.sal > 40K
```

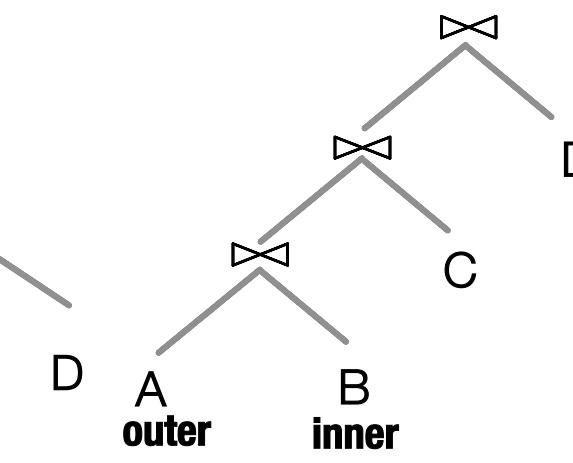
1,000 data pages, 10K tuples
100 pages in B+-tree
depts: 10
Salary Range: 10K – 200K

Multiple-Relation Plans

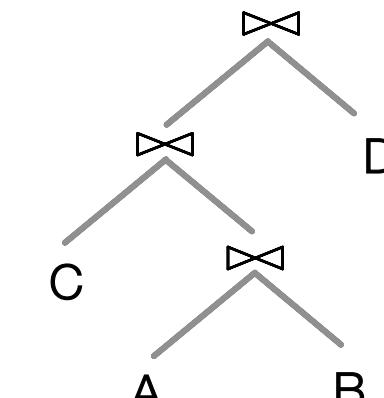
- System R: Only consider **left-deep** join trees
 - Used to restrict the search space
 - Left-deep plans can be **fully pipelined** (usually)
 - Intermediate results not written to temporary files
 - Not all left-deep trees are fully pipelined (e.g., SM join)



Not left-deep join tree



Left-deep join tree



Not left-deep join tree

Enumeration of Left-Deep Plans

- Decide:
 - Join order
 - Join method for each join
- Enumerated using N passes (if N relations joined):
 - **Pass 1:** Find best 1-relation plan for each relation (apply selections and projections first and consider using indexing)
 - For each relation, retain only:
 - Cheapest plan overall (e.g. File scan), plus
 - Cheapest plans that produce **ordered** tuples (e.g., using a B+-tree index). Order may be useful for a later sort-merge join, even though the plan may be more expensive at this point

Enumeration of Left-Deep Plans

Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation **(All 2-relation plans)**

Pass N: Find best way to join result of a (N-1)-relation plan (as outer) to the Nth relation **(All N-relation plans)**

- For each subset of relations, retain only:
 - Cheapest plan overall, plus
 - Cheapest plan for each **interesting order** of the tuples
 - Also, apply selections first, where possible
 - Apply projections first as well, where possible
- Assume pipelining of results to avoid additional I/O

Enumeration of Plans (Contd.)

- ORDER BY, GROUP BY handled as a final step,
- Only “join” relations if there is a connecting join condition
i.e. avoid Cartesian products if possible

Summary

- Query optimization critical to the DBMS performance
 - Helps understand performance impact of database design
- Two parts to optimizing a query:
 - Enumerate alternative plans. (Typically only consider left-deep plans)
 - Estimate cost of each plan: size of result and cost of algorithm
 - Key issues: Statistics, indexes, operator implementations
- Single-relation queries: Pick cheapest access plan + interesting order
- Multiple-relation queries:
 - All single-relation plans are first enumerated. Selections/projections considered as early as possible
 - For each 1-relation plan, consider all ways of joining another relation (as inner)
 - Keep adding 1-relation plan until done
 - At each level, retain cheapest plan, and best plan for each interesting order

Optional Exercises

12.1 (all parts), 15.1, 15.5, 15.7, 15.9