

Transaction Management

Chapter 16 (16.1-16.5),
Chapter 17 (17.1-17.4)


Transactions

- Foundation for concurrent execution and recovery in DBMS
- Transaction is an **atomic** unit of work
 - e.g. Transfer \$500 from chk to saving acct
- Transaction consists of multiple actions
- For performance, DBMS can **interleave** actions from different transaction
- Must guarantee same result as executing transactions **serially**



Example 1


```
Read (A);  
Check (A > $20);  
Pay ($20);  
A = A - 20;  
Write (A);
```



- Interleaving actions of different transactions can cause inconsistency
- DBMS should provide users an illusion of a single-user system
- Could insist on admitting only one transaction at a time
 - Lower utilization: CPU / IO overlap
 - Long running queries starve other queries, reduce overall response time

Example 2

```
Read (A);  
Check (A > $20);  
Pay ($20);  
A = A - 20;  
  
Write (A);
```



DBMS must also guarantee that changes made by partially completed transactions are not seen by other transactions

The ACID Properties

- Atomicity: All actions in the Xact happen, or none happen
 - Don't worry about incomplete Xacts
- Consistency: Consistent DB + consistent Xact \Rightarrow consistent DB
- Isolation: Execution of one Xact is isolated from that of other Xacts
 - Don't worry about other concurrent Xacts
- Durability: If a Xact commits, its effects persist

Schedules

- Transaction seen as a list of actions
 - Read(X), Write(X), commit, abort
- Schedule: An interleaving of the actions from a set of transactions
 - **Complete Schedule:** Each transaction ends in commit or abort
 - **Serial Schedule:** No interleaving actions among transactions
- Initial State + Schedule = Final State



T1	T2
Begin	
R(A)	
W(A)	
	Begin
	R(B)
	W(B)
R(C)	
W(C)	
	Commit
Abort	

Acceptable Schedules

- Serial schedule is “isolated” and “consistent”
- **Serializable schedule:**
 - Final state is what **some complete** serial schedule of **committed** transactions would have produced.
 - Can different serial schedules have different final states?
 - Yes, and we assume **all** are OK!
 - Aborted Xacts?
 - Should have no effect whatsoever
 - Other external actions (besides R/W to DB)
 - e.g. print a computed value, fire a missile, ...
 - Assume (for this class) these values are written to the DB, and can be undone

Example

- Initial State: A=1000, B=100
- Final State: A=990, B=220
- This is serializable because:
 - Running T1 to completion, then running T2 would lead to the same final state (990, 220)
- Note that running T2, then T1 would lead to (1000, 210) which is a different final state. But the schedule is still serializable. Why?

T1: Transfer \$100 from A to B	T2: Add 10% interest to A and B
Begin	Begin
R(A) /A -=100	
W(A)	
R(B) /B +=100	
W(B)	
	R(A) /A *= 1.1
	W(A)
	R(B) /B *= 1.1
	W(B)
	Commit
Commit	

Anomalies Due to Interleaving

- Two actions on the same data “object” conflict if at least one is a write()
- Three anomalous situations for transactions T1 and T2
 - Write-read (WR) conflict anomaly
 - Read-write (RW) conflict anomaly
 - Write-write (WW) conflict anomaly

WR Conflict Anomaly

- Dirty read: reading an uncommitted change
- Could lead to a non-serializable execution
- Example:
 - $\text{@Start } (A, B) = (1000, 100)$
 - End $(990, 210)$
 - $T1 \rightarrow T2:$
 - $(900, 200) \rightarrow (990, 220)$
 - $T2 \rightarrow T1:$
 - $(1100, 110) \rightarrow (1000, 210)$

T1: Transfer \$100 from A to B	T2: Add 10% interest to A and B
Begin	
	Begin
R(A) /A -=100	
W(A)	
	R(A) /A *= 1.1
	W(A)
	R(B) /B *= 1.1
	W(B)
	Commit
R(B) /B +=100	
W(B)	
Commit	

Inconsistent Database

RW Conflict Anomaly

- “Unrepeatable read”
- T1 reads the value of object A, then T2 updates A (before T1 has committed)

T1	T2
$R(A) = 3$	
$W(B) : B=A$	
	$W(A) = 5$
	Commit
$R(A) = 5$	
$W(A) : A=A+1$	
Commit	



Is this
schedule serializable?
Why?

WW Conflict Anomaly

- Overwriting uncommitted changes
- T2 overwrites what T1 wrote
- Usually occurs in conjunction with other anomalies
 - Unless you have “blind writes”

Students in same group (A and B) should get the same project grade

T1 (GSI)	T2 (Prof)
W(A) = 80	
	W(A) = 95
	W(B) = 95
W(B) = 80	
Commit	
	Commit



Is this schedule
Serializable? Why?

What About Aborts?

- Serializable schedule:
Effect equivalent to a
serial schedule of
committed transactions
- As if aborted
transactions **never**
happened

T1 (GSI)	T2 (Prof)
	$W(A) = 2$
	$W(B) = 2$
$W(A) = 1$	
$W(B) = 1$	
Abort	
	Commit

What About Aborts?

- How does one undo the effects of a transaction T1?
 - Covered in logging and recovery lecture
- What if another transaction T2 sees these effects??
 - Must abort and undo T2 too!
 - (Called **cascading abort**)

Can we always undo effects of aborted transactions?

No! Not all schedules
are recoverable

Example:

T1 reads & decrements
account balance by \$100

T2 adds 5% cashback

T1	T2
R(balance) = 1000	
W(balance) = 1000-100	
	R(balance) = 900
	W(balance) = 900 * 1.05
	Commit
Abort	


Desirable Goals

- Different goals:
 - Serializable schedule – (Discussed earlier)
 - Recoverable schedule - Transactions commit only after all transactions whose changes they have read commit
 - Note: A recoverable schedule is **not necessarily** serializable and vice versa
 - Avoid cascading aborts (ACA) - Transactions read only the changes of committed transactions
 - ACA ensures Recoverability but Rec. doesn't ensure ACA

Precedence Graph

- Precedence (or serializability) graph for schedule L
 - A node for each committed transaction in L
 - An arc from T_i to T_j if some action in T_i precedes and conflicts with some action in T_j
 - (Remember that a conflict may or may not be an anomaly)

T1	T2	T3
R(A)		
	W(A)	
	Commit	
W(A)		
Commit		
		W(A)
		Commit



Terminology


- Two schedules are **conflict equivalent** if
 - involve the same transactions
 - order each pair of conflicting transactions in the same way
- A schedule is **conflict serializable** if it is conflict equivalent to a serial schedule
- All **conflict serializable schedules** are also **serializable** (opposite is not true!)

Example



Is this schedule serializable? Is it conflict-serializable too?


T1	T2	T3
R(A)		
	W(A)	
	Commit	
W(A)		
Commit		
		W(A)
		Commit




Conflict Serializability & Graphs


- **Theorem:** A schedule is conflict serializable if and only if its precedence graph is acyclic
 - Equivalent serial schedule is given by any topological sort over graph

T1	T2	T3
	W(A)	
	Commit	
R(A)		
W(A)		
Commit		
		W(A)
		Commit



Quiz: Are these schedules conflict-serializable?

T1: Transfer \$100 from A to B	T2: Add 10% interest to A and B
Begin	
	Begin
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
	W(B)
	Commit
R(B)	
W(B)	
Commit	

T1: Transfer \$100 from A to B	T2: Add 10% interest to A and B
Begin	
	Begin
R(A)	
W(A)	
R(B)	
W(B)	
Commit	
	
	R(A)
	W(A)
	R(B)
	W(B)
	Commit



Intro to Locking

- What can a DBMS do to guarantee an acceptable schedule?
- Lock info maintained by a “lock manager”:
 - Stores (XID, RID, Mode) triples.
 - This is a simplistic view; suffices for now.
 - Mode $\in \{S, X\}$
 - Lock compatibility table:
- If a transaction can’t get a lock
 - Suspended on a wait queue

	--	S	X
--	✓	✓	✓
S	✓	✓	
X	✓		

Two-Phase Locking (2PL)

- 2PL:
 - If T wants to read (modify) an object, first obtains a shared (exclusive) lock
 - If T releases any lock, it can acquire **no new locks!**
 - Guarantees conflict-serializability (and thus serializability)!
- Strict 2PL:
 - **Hold all locks until end of Xact**
 - Guarantees conflict-serializability (and thus serializability)
 - Guarantees ACA (and thus recoverability)



Example

- Two accounts, A and B
- Two transactions, T1 and T2
- T1: Transfer \$100 from A to B
- T2: Add 10% interest to A and B

Example - Strict 2PL

T1 acquires S lock on A

T1 acquires X lock on A

T1 acquires S lock on B

T1 acquires X lock on B

T1 releases all locks

T2 acquires S lock on A

...

T1: Transfer \$100 from A to B	T2: Add 10% interest to A and B
Begin	Begin
R(A); /A -= 100	
W(A)	
R(B); B += 100	
W(B)	
Commit	
	R(A) /A * = 1.1
	W(A)
	R(B), /B * = 1.1
	W(B)
	Commit

2PL & Serializability

- 2PL guarantees acyclic precedence graph
 - Guarantees a conflict serializable schedule
 - Intuitively, equivalent serial schedule given by order in which transactions enter shrinking phase
- Why Strict 2PL?
 - Guarantees ACA (read only committed values)
 - How? Hold X locks until end of transaction;

Quiz

- Is this schedule allowed by 2PL?
Strict 2PL?
- Is it serializable?
 - If so, what is the corresponding serial schedule?
- Is the schedule recoverable? ACA?

T1: Transfer \$100 from A to B	T2: Add 10% interest to A, B & C
Begin	
	Begin
	R(C)
	W(C)
R(A)	
W(A)	
R(B)	
W(B)	
Commit	
	R(A)
	W(A)
	R(B)
	W(B)
	Commit

Scheduling Transactions

- A transaction can be in a state of:
 - Running
 - On a runnable queue, waiting for CPU
 - On a lock queue, waiting for a lock
 - Suspended, waiting for I/O device
- Transaction moves between the running state or one of the queues

Lock Manager Implementation

- **Lock Table:** A hash table of Lock Entries. Each entry:
 - OID: object ID
 - Number of trxs holding a lock on this object
 - Mode: shared (read) or exclusive (write) lock
 - **List:** A wait queue of other trxs requesting a lock on this object

On lock requests for object OID:

- If S lock requested:
 - If its queue is empty and currently not in X mode:
 - Grant the lock, set mode to S and increment the count
- If X lock requested
 - If no one is currently holding the lock (\Rightarrow queue will be empty too)
 - Grant the lock, set mode to X, set count to 1
- Otherwise, add this trx to the queue and suspend it

Lock Manager Implementation

- On lock release (OID, XID): happens upon commit/abort
 - Update the lock entry
 - Examine lock's wait queue, grant lock to the head of the queue (or to multiple of them if S mode)
- **Note:** Lock request handled atomically! (via mutex in OS)


What About Deadlocks?

T1 gets S, then X lock on A

T2 gets S, then X lock on B

T1 waits for lock on B

T2 waits for lock on A



“Waits-for” Graph

T1: Transfer \$100 from A to B	T2: Add 10% interest to A and B
Begin	
	Begin
R(A)	
W(A)	
	R(B)
	W(B)
R(B)	
W(B)	
Commit	
	R(A)
	W(A)
	Commit

Deadlock Detection

- Observation:
 - Deadlocks are rare
 - Often involve only a few transactions
 - Detect rather than prevent
- Approach #1: Lock Mgr maintains waits-for graph
 - Periodically check graph for cycles
 - Abort some Xact to break the cycle
- Approach #2 (Simpler hack): use *time-outs*
 - Abort if no progress made for a while

Deadlock Prevention

- Assign priorities to transactions
 - Use timestamp to assign priorities
- T_i requests a lock, held by T_j (in a conflicting mode)
 - Approach #1 (Wait-Die): If T_i has higher priority, it waits; else T_i aborts
 - Approach #2 (Wound-Wait): If T_i has higher priority, abort T_j ; else T_i waits
 - After abort, restart with original timestamp
 - Guarantees the transaction eventually runs!

Non-Locking CC Protocols

- Locking most common technique for guaranteeing transaction serializability
 - Used in most commercial systems
- Other approaches exist (beyond scope of class):
 - Optimistic CC
 - 2PL locking protocols avoid conflict by blocking (waiting)
 - Optimistic CC instead undoes transactions if a conflict occurs
 - Anticipates that conflicts rarely occur
 - Multiversion CC
 - Make sure transactions never have to wait to read an object
 - Maintain multiple versions of each object, each with a timestamp
 - Used by Oracle

Transaction and Constraints

Create Table A (akey, bref, ...)

Create Table B (bkey, aref, ...)



Q: How to insert
the first tuple,
either in A or B?

- Solution:
 - Insert tuples in the same transaction
 - Defer the constraint checking
- SQL constraint modes
 - **BEGIN DEFERRED**: Check at commit time.
 - **BEGIN IMMEDIATE**: Check immediately

Summary

- Locking commonly used by DBMS for concurrency control
- 2PL guarantees a serializable schedule
- Strict 2PL guarantees a serializable, recoverable, ACA schedule
- Lock manager handles lock requests from transactions
- Deadlock rare, but must be detected or prevented

Optional Exercises

- From the textbook: 16.1, 16.3, 17.5
- **Additional Review Exercise:** Recall the three kinds of conflicts from last class (RW, WR, WW). For each, think of an example where the conflict occurs. How does 2PL prevent these conflicts?